

Ron S. Kenett, Shelemyahu Zacks, Peter
Gedeck

Industrial Statistics: A Computer Based Approach with Python

Volume 2 - Solutions

September 10, 2022

Springer Nature

Contents

1	Introduction to industrial statistics	5
2	Basic Tools and Principles of Process Control	9
3	Advanced Methods of Statistical Process Control	27
4	Multivariate Statistical Process Control	53
5	Classical Design and Analysis of Experiments	63
6	Quality by Design	79
7	Computer Experiments	85
8	Cybermanufacturing and digital twins	95
9	Reliability Analysis	99
10	Bayesian Reliability Estimation and Prediction	115
11	Sampling Plans for Batch and Sequential Inspection	119

Chapter 1

Introduction to industrial statistics

Exercise 1.1 Describe three work environments where quality is assured by 100% inspection of outputs (as opposed to process control).

Solution 1.1 Examples of 100% inspection:

1. Airplane pilots rely heavily on checklists before takeoff. These lists provide a systematic check on safety and functional items that can be verified by relatively simple pilot inspection. The lists are used in every flight.
2. The education system relies on tests to evaluate what students learn in various courses and classes. Tests are administered to all students attending a class. The grades received on these tests reflect the student's ability, the course syllabus and training material, the instructor's ability, the classroom environment, etc.
3. Production lines of electronic products typically include automatic testers that screen products as PASS or FAIL. Failed products are directed to rework departments where the problems are diagnosed and fixed before the products are retested.

Exercise 1.2 Search periodicals, such as Business Week, Fortune, Time and Newsweek and newspapers such as the New York Times and Wall Street Journal for information on quality initiatives in service, healthcare, governmental and industrial organizations. Summarize three such initiatives indicating what was done and what were the concrete outcomes of these initiatives.

Solution 1.2 Possible articles to evaluate are:

- <https://www.industryweek.com/operations/quality/article/21213860/quality-initiatives-deserve-better-from-industry-40>
- <https://www.nytimes.com/2017/03/13/upshot/lousy-customer-service-a-better-way-in-health-care.html>
- <https://www.nytimes.com/2007/02/08/business/08scene.html>
- <https://www.wsj.com/articles/BL-CIOB-6959>
- <https://www.newsweek.com/trends-opportunities-threats-manufacturing-operations-management-software-market-1696962>

Exercise 1.3 Provide examples of the three types of production systems.

- a) Continuous flow production.
- b) Discrete mass production
- c) Industry 4 production

Solution 1.3 Examples of production systems:

- a) *Continuous flow production*: Polymerization is used in the petroleum industry and in the manufacture of synthetic rubber. In polymerization, a reaction occurs in which two or more molecules of the same substance combined to form a compound from which the original substance may or may not be regenerated. Typical parameters affecting the process are feed rate, polymerizer temperature, and sludge levels in the separator kettles. Typical measured responses are yield, concentration, color, and texture.
- b) *Discrete mass production*: A college or university is, in fact, a discrete mass production system where students are acquiring knowledge and experience through various combinations of classes, laboratories, projects, and homework assignments.
- c) *Job shop*: Modern print houses are typical job shop operations. Customers provide existing material, computer files, or just an idea. The print house is responsible for producing hard copies in various sizes, colours, and quantities. The process involves several steps combining human labor with machine processing.

Exercise 1.4 What management approach cannot work with continuous flow production?

Solution 1.4 Sample inspection for acceptance sampling procedures determining the quality of batches of units.

Exercise 1.5 What management approach characterizes

- a) a school system?
- b) a group of scouts?
- c) a football team?

Solution 1.5 a) *A school system?*: In most cases, school systems rely on individual learning combined with 100% inspection. Many schools also encourage cooperative learning efforts such as team projects.
b) *A military unit?*: Traditionally, this is a classical organization operating with strict regulations and 100% inspection. Improvements are achieved through training and exercises with comprehensive performance reviews.
c) *A football team?*: Serious sports teams invest huge efforts in putting together a winning team combination. This includes screening and selection, team building exercises, on-going feedback during training and games, and detailed analysis and review of self and other teams performance.

Exercise 1.6 Provide examples of how you, personally, apply inspection, process control or quality by design approaches.

- a) as a student.
- b) in your parents' house.
- c) with your friends

Solution 1.6 Examples of how you personally apply the four management approaches:

- a) As a student: With an old car that you maintain with a "don't fix what ain't broke" strategy you are fire fighting. Reviewing your bank account statements is typically done on a 100% basis. Serious learning involves on going process control to assure success in the final exam. Proper preparation, ahead of time, through self study and learning from the experience of others is a form of ensuring quality by design.
- b) In your parent's home: This is a personal exercise for self thinking.
- c) With your friends: This is also a personal exercise for self thinking.

Exercise 1.7 Evaluate the information quality of a case study provided by your instructor.

Solution 1.7 Three case studies are provided in slides 20, 21 and 22 in https://ceeds.unimi.it/wp-content/uploads/2020/02/Kenett_Analytics_2020-1.pdf. Their information quality assessment is provided in slides 45-59.

Chapter 2

Basic Tools and Principles of Process Control

Import required modules and define required functions

```
import numpy as np
import pandas as pd
import pingouin as pg
from scipy import stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
import statsmodels.stats as sms
from statsmodels.graphics.mosaicplot import mosaic
import seaborn as sns
import matplotlib.pyplot as plt
import pwlf

import mistat
```

Exercise 2.1 Use Python and file **OELECT.csv** to chart the individual electrical outputs of the 99 circuits. Do you observe any trend or non-random pattern in the data? [Use under SPC the option of Individual chart. For Mu and Sigma use “historical” values, which are \bar{X} and S .]

Solution 2.1 The following Python code creates the chart shown in Fig. 2.1.

{fig:xbarone-oelect}

```
oelect = mistat.load_data('OELECT')
qcc = mistat.QualityControlChart(oelect, qcc_type='xbarone',
                                std_dev='SD')
qcc.plot()
plt.show()
```

It seems that the data are randomly distributed around a mean value of 219.25.

Exercise 2.2 Chart the individual variability of the length of steel rods, in **STEEL-ROD.csv** file. Is there any perceived assignable cause of non-randomness?

Solution 2.2 The following Python code creates the chart shown in Fig. 2.2.

{fig:xbarone-steelrod}

```
steelrod = mistat.load_data('STEELROD')
qcc = mistat.QualityControlChart(steelrod, qcc_type='xbarone',
```




Fig. 2.1 X-barone chart of **OELECT.csv** dataset

{fig:xbarone-oelect}

```
qcc.plot()
plt.show()
```

It seems that the data are randomly distributed around a mean value of 19.89.

Exercise 2.3 Examine the chart of the previous Exercise for possible patterns of non-randomness.

Solution 2.3 No patterns of non-randomness are apparent.


Exercise 2.4 Test the data in file **OTURB2.csv**¹ for lack of randomness. In this file we have three columns. In the first we have the sample size. In the second and third we have the sample means and standard deviation. Chart the individual means. For the historical mean use the mean of column *c2*. For historical standard deviation use $(\hat{\sigma}^2/5)^{1/2}$, where $\hat{\sigma}^2$ is the pooled sample **variance**.

{fig:xbarone-oturb2}

Solution 2.4 The following Python code creates the chart shown in Fig. 2.3.

```
oturb2 = mistat.load_data('OTURB2')
# print(oturb2)
sd = np.sqrt(oturb2['xbar'].var() / 5)
center = oturb2['xbar'].mean()
print(sd, center)
qcc = mistat.QualityControlChart(oturb2['xbar'], qcc_type='xbarone',
```

¹ **TODO:** The dataset has 10 entries while the figure in the old solutions pdf has 20. With only 10 entries, we see a downward trend in the data.



{fig:xbarone-steelrod}

Fig. 2.2 X-bar chart of STEELROD.csv dataset

```
center=center, std_dev=sd)
qcc.plot()
plt.show()
```

| 0.07199111056234651 0.6526

No pattern of nonrandomness is apparent.

Exercise 2.5 A sudden change in a process lowers the process mean by one standard deviation. It has been determined that the quality characteristic being measured is approximately normally distributed and that the change had no effect on the process variance.

- What percentage of points are expected to fall outside the control limits on the \bar{X} chart if the subgroup size is 4?
- Answer the same question for subgroups of size 6.
- Answer the same question for subgroups of size 9.

Solution 2.5 If $\bar{X}_n \sim N(\mu - \sigma, \frac{\sigma}{\sqrt{n}})$ the proportion of points expected to fall outside the control limits is

$$\pi_n \equiv \Pr \left\{ \bar{X}_n < \mu - \frac{3\sigma}{\sqrt{n}} \right\} + \Pr \left\{ \bar{X}_n > \mu + \frac{3\sigma}{\sqrt{n}} \right\}.$$

- (a) For $n = 4$, $\pi_n = 0.159$; (b) For $n = 6$, $\pi_n = 0.291$; (c) For $n = 9$, $\pi_n = 0.500$.

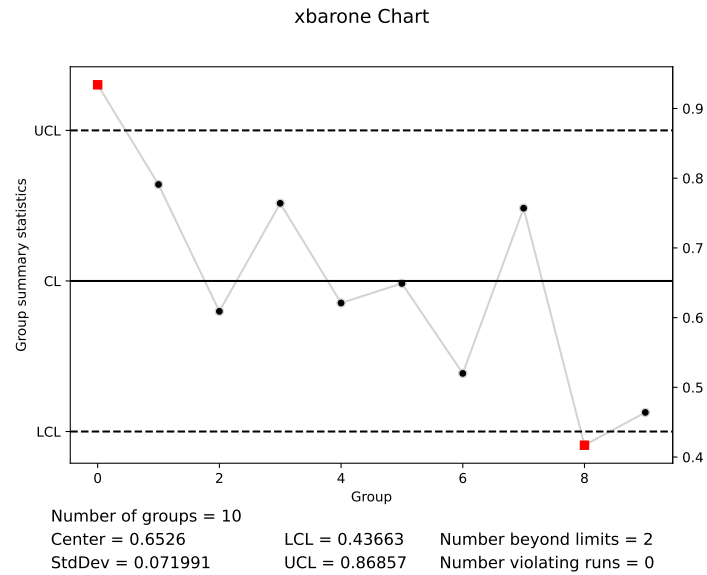


Fig. 2.3 X-barone chart of **OTURB2.csv** dataset

(fig:xbarone-oturb2)

Exercise 2.6 Make capability analysis of the electric output (volts) of 99 circuits in data file **OELECT.csv**, with target value of $\mu_0 = 220$ and $LSL = 210$, $USL = 230$.

Solution 2.6 In Python:

```
oelect = mistat.load_data('OELECT')

qcc = mistat.QualityControlChart(oelect, qcc_type='xbarone',
                                std_dev='SD')
pc = mistat.ProcessCapability(qcc, spec_limits = [210, 230])
pc.plot()
plt.show()
pc.summary()
```


```
Process Capability Analysis

Number of obs = 99          Target = 220.00
Center = 219.25           LSL = 210.00
StdDev = 4.014219         USL = 230.00

Capability indices:

      Value      2.5%      97.5%
Cp      0.8304      0.7142      0.9463
Cp_l     0.7679      0.6622      0.8737
Cp_u     0.8928      0.7743      1.0113
Cp_k     0.7679      0.6420      0.8939
Cpm     0.8162      0.7007      0.9315

Exp<LSL    1%    Obs<LSL    0%
Exp>USL    0%    Obs>USL    0%
```



{fig:process-capability} **Fig. 2.4** Process capability analysis of **OELECT.csv** dataset

{fig:process-capability} Performing the capability analysis (see Fig. 2.4) we found that $C_p = 0.83$. Although $C_p = 0.83$, there is room for improvement by designing the process with reduced variability. See Chapter 5.

{ch:design-of-exp}

Exercise 2.7 Estimate the capability index C_{pk} for the output of the electronic circuits, based on data file **OELECT.csv** when $LSL = 210$ and $USL = 230$. Determine the point estimate as well as its confidence interval, with confidence level 0.95.

Solution 2.7 Using the **mistat** package, we get:

```
oelect = mistat.load_data('OELECT')
qcc = mistat.QualityControlChart(oelect, qcc_type='xbarone',
                                std_dev=4.004, center=219.25)
pc = mistat.ProcessCapability(qcc, spec_limits = [210, 230],
                             confidence_level=0.975)

pc.summary()
pc = mistat.ProcessCapability(qcc, spec_limits = [210, 230],
                             confidence_level=0.95)

pc.summary()
```

```
Process Capability Analysis

Number of obs = 99      Target = 220.00
Center = 219.25        LSL = 210.00
StdDev = 4.004000      USL = 230.00

Capability indices:

      Value      2.5%      97.5%
Cp      0.8325      0.7002      0.9663
```

```

Cp_l    0.7701  0.6438  0.8963
Cp_u    0.8949  0.7535  1.0364
Cp_k    0.7701  0.6257  0.9144
Cpm     0.8183  0.6868  0.9513

Exp<LSL  1%   Obs<LSL  0%
Exp>USL  0%   Obs>USL  0%
Process Capability Analysis

Number of obs = 99          Target = 220.00
Center = 219.25           LSL = 210.00
StdDev = 4.004000         USL = 230.00

Capability indices:

      Value      2.5%   97.5%
Cp      0.8325   0.7161  0.9488
Cp_l    0.7701   0.6641  0.8760
Cp_u    0.8949   0.7762  1.0136
Cp_k    0.7701   0.6438  0.8963
Cpm     0.8183   0.7025  0.9339

Exp<LSL  1%   Obs<LSL  0%
Exp>USL  0%   Obs>USL  0%

```

With $\bar{X} = 219.25$, $S = 4.004$, $n = 99$, $\xi_U = 230$ and $\xi_L = 210$ we obtained $\hat{C}_{pl} = 0.77$, $\hat{C}_{pu} = 0.895$ and $\hat{C}_{pk} = 0.77$.

The `miostat` implementation uses the normal distribution to calculate confidence intervals. Using the F-distribution with $F_{.975}[1, 98] = 5.1982$, we get:

```

def confidenceLimitSL(Cp, n, alpha):
    F = stats.f(1, n-1).ppf(1-(1-alpha)/2)
    a = np.sqrt(F/n) * np.sqrt(Cp**2/2 + (1 - F/(2*n))/9)
    b = 1 - F/(2*n)
    return (Cp - a) / b, (Cp + a) / b

n = len(oelect)
rho_1L, rho_1U = confidenceLimitSL(pc.Cp_l, n, 0.95)
rho_2L, rho_2U = confidenceLimitSL(pc.Cp_u, n, 0.95)
print(rho_1L, rho_1U)
print(rho_2L, rho_2U)

```

```

0.6413037077789147 0.9402124117777945
0.751435033838329 1.0865431596145048

```

The confidence intervals for C_{pl} and C_{pu} , at 0.95 confidence level, are:

	Lower Limit	Upper Limit
C_{pl}	0.6413	0.9402
C_{pu}	0.7514	1.0865
C_{pk}	0.6413	0.9402

The 0.95-Confidence interval for C_{pk} is (0.6413, 0.9402).

Exercise 2.8 Estimate the capability index for the steel rods, given in data file **STEELROD.csv**, when the length specifications are $\xi_L = 19$ and $\xi_U = 21$ [cm] and the level of confidence is $1 - \alpha = 0.95$.

Solution 2.8 In Python:

```

steelrod = mistat.load_data('STEELROD')
qcc = mistat.QualityControlChart(steelrod, qcc_type='xbarone',
                                std_dev='SD')
pc = mistat.ProcessCapability(qcc, spec_limits = [19, 21],
                              confidence_level=0.95)
pc.summary()

```

```

Process Capability Analysis

Number of obs = 50          Target = 20.00
      Center = 19.89          LSL = 19.00
      StdDev = 1.129109       USL = 21.00

Capability indices:

      Value      2.5%      97.5%
Cp      0.2952    0.2369    0.3534
Cp_l     0.2628    0.1738    0.3518
Cp_u     0.3276    0.2329    0.4224
Cp_k     0.2628    0.1568    0.3688
Cpm     0.2938    0.2361    0.3514

Exp<LSL  22%    Obs<LSL  22%
Exp>USL  16%    Obs>USL  20%

```

Using the STEELROD.csv data with $\xi_L = 19$ and $\xi_U = 21$ we obtained $\hat{C}_p = 0.295$, $\hat{C}_{pu} = 0.328$, $\hat{C}_{pl} = 0.263$ and $\hat{C}_{pk} = 0.263$. We use the function from the previous exercise to calculate the confidence interval.

```

n = len(steelrod)
rho_1L, rho_1U = confidenceLimitSL(pc.Cp_l, n, 0.95)
rho_2L, rho_2U = confidenceLimitSL(pc.Cp_u, n, 0.95)
print(rho_1L, pc.Cp_l, rho_1U)
print(rho_2L, pc.Cp_u, rho_2U)

```

```

0.14851207708654524 0.26280004915567645 0.40677992256113854
0.2084480020651775 0.32763581095436156 0.4838408510212421

```

A 0.95-confidence interval for C_{pk} is (0.1485, 0.4068).

Exercise 2.9 The specification limits of the piston cycle times are 0.05 ± 0.01 seconds. Generate 20 cycle times at the lower level of the 7 control parameters.

- Compute C_p and C_{pk} .
- Compute a 95% confidence interval for C_{pk} .

Generate 20 cycle times at the upper level of the 7 control factors.

- Recompute C_p and C_{pk} .
- Recompute a 95% confidence interval for C_{pk} .
- Is there a significant difference in process capability between lower and upper operating levels in the piston simulator?

Solution 2.9 ² Simulate 20 observations at the low level of the 7 control parameters.

² **TODO:** Adjusted the exercise to version 2 of the piston simulator - please check that it is reasonable.

```
settings = {'m': 30, 's': 0.005, 'k': 1000, 't': 290,
            'p0': 90_000, 'v0': 0.002, 't0': 340}

simulator = mistat.PistonSimulator(n_simulation=20, n_replicate=1, seed=1,
                                   **settings)

Ps = simulator.simulate()
cycleTime = mistat.qcc_groups(Ps['seconds'], Ps['group'])

qcc = mistat.QualityControlChart(cycleTime, std_dev= Ps['seconds'].std())
print(f'Mean      {qcc.center:.4f}')
print(f'Std.Dev   {qcc.std_dev:.4f}')
```

```
| Mean      0.0257
| Std.Dev   0.0078
```

If all the 7 control parameters are at their low level, the mean cycle time of 20 observations is $\bar{X}_{20} = 0.0257$. This value is outside the specification limits, 0.04 and 0.06. Thus the process, at the low parameter values is incapable of satisfying the specs. The standard deviation is $S_{20} = 0.0078$.

(a) Calculate the capability index C_p

```
pc = mistat.ProcessCapability(qcc, spec_limits = [0.04, 0.06],
                             confidence_level=0.95)
print(f'C_p {pc.Cp:.3f}')
```

```
| C_p 0.425
```

If the process mean can be moved to 0.05, the C_p value is 0.425.

(b)³

(c) Under the high values of the seven control parameters we get

```
settings = {'m': 60, 's': 0.02, 'k': 5_000, 't': 296,
            'p0': 110_000, 'v0': 0.01, 't0': 360}

simulator = mistat.PistonSimulator(n_simulation=20, n_replicate=1, seed=1, **settings)
Ps = simulator.simulate()
cycleTime = mistat.qcc_groups(Ps['seconds'], Ps['group'])
qcc = mistat.QualityControlChart(cycleTime, std_dev= Ps['seconds'].std())
print(f'Mean      {qcc.center:.4f}')
print(f'Std.Dev   {qcc.std_dev:.4f}')
```

```
| Mean      0.0568
| Std.Dev   0.0037
```

$\bar{X}_{20} = 0.0568$ and $S_{20} = 0.0037$. The mean is within the specification limit 0.05 ± 0.01 .

For these values, we get the following for the capability indices.

```
pc = mistat.ProcessCapability(qcc, spec_limits = [0.04, 0.06],
                             confidence_level=0.95)
pc.summary()
```

³ **TODO:** Is 0.425 low? The limits of C_{pk} are -0.3657 -0.8503. What does this mean?


```

Process Capability Analysis

Number of obs = 20          Target = 0.05
      Center = 0.06          LSL = 0.04
      StdDev = 0.003720      USL = 0.06

Capability indices:

      Value      2.5%      97.5%
Cp      0.8960      0.6135      1.1782
Cp_l     1.5090      1.0881      1.9298
Cp_u     0.2830      0.1390      0.4270
Cp_k     0.2830      0.1115      0.4546
Cpm      0.4281      0.2543      0.6020

Exp<LSL    0%    Obs<LSL    0%
Exp>USL    20%   Obs>USL    10%

```

The values are: $C_p = 0.890$ and $C_{pk} = 0.283$.

(d) Under the normal approximation the confidence interval for C_{pk} , at level 0.95, is (0.1115, 0.4546).

(e) As mentioned above, the gas turbine at low control levels is incapable of satisfying the specifications of 0.05 ± 0.01 seconds. At high levels its C_{pk} is not greater than 0.282. In Chapter 5 we study experimental methods for finding the combination of control levels, which maximizes the capability. (ch:design-of-exp)

Exercise 2.10 A fiber manufacturer has a large contract which stipulates that its fiber, among other properties, have tensile strength greater than 1.800 [grams/fiber] in 95% of the fiber used. The manufacturer states the standard deviation of the process is 0.015 grams.

- Assuming a process under statistical control, what is the smallest nominal value of the mean that will assure compliance with the contract?
- Given the nominal value in part a) what are the control limits of \bar{X} and S charts for subgroups of size 6?
- What is the process capability, if the process mean is $\mu = 1.82$?

Solution 2.10 ⁴ (a) 1.8247 (b) 1.8063 and 1.8430 (c) $C_{pk} = 0.44$.

Exercise 2.11 The output voltage of a power supply is specified as 350 ± 5 volts DC. Subgroups of four units are drawn from every batch and submitted to special quality control tests. The data from 30 subgroups on output voltage produced $\sum_{i=1}^{30} \bar{X} = 10,500.00$ ⁵ and $\sum_{i=1}^{30} R_i = 86.5$

- Compute the control limits for \bar{X} and R .
- Assuming statistical control and a normal distribution of output voltage, what properties of defective product is being made?
- If the power supplies are set to a nominal value of 350 volts, what is now the proportion of defective products?
- Compute the new control limits for \bar{X} and R .

⁴ **TODO:** Check how this is calculated.

⁵ **TODO:** should this be 10,950 as used in the solution?

- (e) If these new control limits are used but the adjustment to 350 volts is not carried out, what is the probability that this fact will not be detected on the first subgroup?
- (f) What is the process capability before and after the adjustment of the nominal value to 350 volts? Compute both C_p and C_{pk} .

Solution 2.11 ⁶ (a) $\bar{\bar{X}} = \frac{10950}{30} = 365$, $\bar{S} = 1.153$, $n = 4$, $\hat{\sigma} = \bar{S}/c(n) = 1.252$. The control limits for \bar{X}_4 are

$$UCL = 365 + 3 \times \frac{1.252}{\sqrt{4}} = 366.878 \quad \text{and} \quad LCL = 365 - 3 \times \frac{1.252}{\sqrt{4}} = 363.122.$$

(b) When $X \sim N(\mu = 365, \sigma = 1.252)$,

$$1 - \Pr\{365 < X < 375\} = 1 - \Phi\left(\frac{375 - 365}{1.252}\right) + \Phi\left(\frac{365 - 365}{1.252}\right) = 0.5.$$

(c) When $X \sim N(\mu = 370, \sigma = 1.252)$,

$$1 - \Pr\{365 < X < 375\} = 2 \times \left(1 - \Phi\left(\frac{5}{1.252}\right)\right) = 0.000065.$$

(d) $UCL = 370 + 3 \times \frac{1.252}{\sqrt{4}} = 371.878$ and $LCL = 370 - 3 \times \frac{1.252}{\sqrt{4}} = 368.122$.

(e) When $\bar{X}_4 \sim N(\mu = 365, \sigma = 1.252/\sqrt{4})$,

$$\Pr\{368.122 < \bar{X}_4 < 371.878\} =$$

$$\Phi\left(\frac{371.878 - 365}{1.252/\sqrt{4}}\right) - \Phi\left(\frac{368.122 - 365}{1.252/\sqrt{4}}\right) = 0.$$

(f) The process capability before and after the adjustment is

	Before	After
C_p	1.331	1.331
C_{pk}	0	1.331

{exc:pareto-stat-anal}

Exercise 2.12 The following data were collected in a circuit pack production plant during October

⁶ **TODO:** Check how this is calculated.

	Number of Nonconformities
Missing component	293
Wrong component	431
Too much solder	120
Insufficient solder	132
Failed component	183

An improvement team recommended several changes that were implemented in the first week of November. The following data were collected in the second week of November.

	Number of Nonconformities
Missing component	34
Wrong component	52
Too much solder	25
Insufficient solder	34
Failed component	18

- Construct Pareto charts of the nonconformities in October and the second week of November.
- Has the improvement team produced significant differences in the type of nonconformities?

Solution 2.12 (a) Fig. 2.5 is created using the following code

[fig:pareto-nonconformities]

```

october = pd.DataFrame([
    ['Missing component', 293],
    ['Wrong component', 431],
    ['Too much solder', 120],
    ['Insufficient solder', 132],
    ['Failed component', 183],
], columns=['Issue', 'Count'])
november = pd.DataFrame([
    ['Missing component', 34],
    ['Wrong component', 52],
    ['Too much solder', 25],
    ['Insufficient solder', 34],
    ['Failed component', 18],
], columns=['Issue', 'Count'])

def makeParetoChart(data, ax, title):
    paretoChart = mistat.ParetoChart(data['Count'], labels=data['Issue'])
    paretoChart.plot(rotation=30, ha='right', ax=ax)
    ax.set_title(title)

fig, axes = plt.subplots(ncols=2, figsize=(8,4))
makeParetoChart(october, axes[0], 'October')
makeParetoChart(november, axes[1], 'November (2nd week)')
fig.suptitle('')
plt.tight_layout()
plt.show()

```

(b) Using the test described in Section 2.6 we find that the improvement team has produced significant differences in the type of nonconformities. In particular,

[sec:stat-anal-pareto]




Fig. 2.5 Pareto chart noconformities in October and the second week of November

{fig:pareto-nonconformities}

the difference in the ‘Insufficient solder’ category is significant at the 1% level and the difference in the ‘Too much solder’ category is significant at the 10% level (see Table 2.5). The computations involved are summarized in Table 2.1.

{tbl:pareto-test-and-nonconform}

Table 2.1 Computations for Pareto significance analysis (Exercise 2.12)

Category	October	p	November	Expected	(Obs. - Exp.)	Z
Missing	293	0.253	34	41.239	-7.239	-1.304
Wrong	431	0.372	52	60.636	-8.636	-1.399
Too Much	120	0.103	25	16.789	8.211	2.116
Insufficient	132	0.114	34	18.582	15.418	3.800
Failed	183	0.158	18	25.754	-7.754	-1.665
Total	1159		163			

{tbl:pareto-calc-nonconform}

Exercise 2.13 Control charts for \bar{X} and R are maintained on total soluble solids produced at 20°C in parts per million (ppm). Samples are drawn from production containers every hour and tested in a special test device. The test results are organized into subgroups of $n = 5$ measurements, corresponding to 5 hours of production. After 125 hours of production we find that $\sum_{i=1}^{25} \bar{X}_i = 390.8$ and $\sum_{i=1}^{25} R_i = 84$. The specification on the process states that containers with more than 18 ppm of total soluble solids should be reprocessed.

- Compute an appropriate capability index.
- Assuming a normal distribution and statistical control, what proportion of the sample measurements are expected to be out of spec?
- Compute the control limits for \bar{X} and R .

Solution 2.13 From the data, $\bar{\bar{X}} = 15.632$ ppm and $\bar{R} = 3.36$. An estimate of the process standard deviation is $\hat{\sigma} = \bar{R}/d(5) = 1.4446$.

$$(a) C_{pu} = \frac{18 - 15.632}{3 * \hat{\sigma}} = 0.546.$$

(b) The proportion expected to be out of spec is 0.051.

(c) The control limits for \bar{X} are $LCL = 13.694$ and $UCL = 17.570$. The control limits for R_5 are $LCL = 0$ and $UCL = 7.106$.

{exc:piston-time-shift}

Exercise 2.14 Part I: Run the piston simulator at the lower levels of the 7 piston parameters and generate 100 cycle times. Add 0.02 to the last 50 cycle times.

(a) Compute control limits of \bar{X} and R by constructing subgroups of size 5, and analyze the control charts.

Part II: Assign a random number, R_i , from $U(0, 1)$, to each cycle time. Sort the 100 cycle times by $R_i, i = 1, \dots, 100$.⁷

(b) Recompute the control limits of \bar{X} and R and reanalyze the control charts.

(c) Explain the differences between (a) and (b).

Solution 2.14 Part I (a) Fig. 2.6 is created using the following code.

{fig:piston-time-shift}

```
settings = {'m': 30, 's': 0.005, 'v0': 0.002, 'k': 1000,
           'p0': 90_000, 't': 290, 't0': 340}

simulator = mistat.PistonSimulator(n_simulation=20, n_replicate=5, seed=1, **settings)
Ps = simulator.simulate()

# Add 0.02 seconds to last 50 simulation results
Ps.loc[50:,'seconds'] = Ps.loc[50:,'seconds'] + 0.02

def makeQCCplot(data, reference, qcc_type, title):
    # convert to groups
    data = mistat.qcc_groups(data['seconds'], data['group'])
    reference = mistat.qcc_groups(reference['seconds'], reference['group'])
    # calculate control limits based on reference data
    qcc_ref = mistat.QualityControlChart(reference, qcc_type=qcc_type)
    qcc = mistat.QualityControlChart(data, qcc_type=qcc_type,
                                     center=qcc_ref.center, limits=qcc_ref.limits)
    return qcc.plot(title=title)

# create xbar and R control charts
reference = Ps.iloc[:50, ]

makeQCCplot(Ps, reference, 'xbar', 'for cycleTime')
plt.show()
makeQCCplot(Ps, reference, 'R', 'for cycleTime')
plt.show()
```

See the control charts in Fig. 2.6. The increase in cycle times after group 9 is so large that the means of the subgroup are immediately out of the control limits. The shift has no influence on the ranges within subgroups.

{fig:piston-time-shift}

Part II (b) Fig. 2.7 is created using the following code.

{fig:piston-randomization}

```
# Create random numbers from U(0,1)
R = stats.uniform.rvs(size=100)
```

⁷ **TODO:** This is just randomly shuffling the calculated cycle times. Why not say it this way?

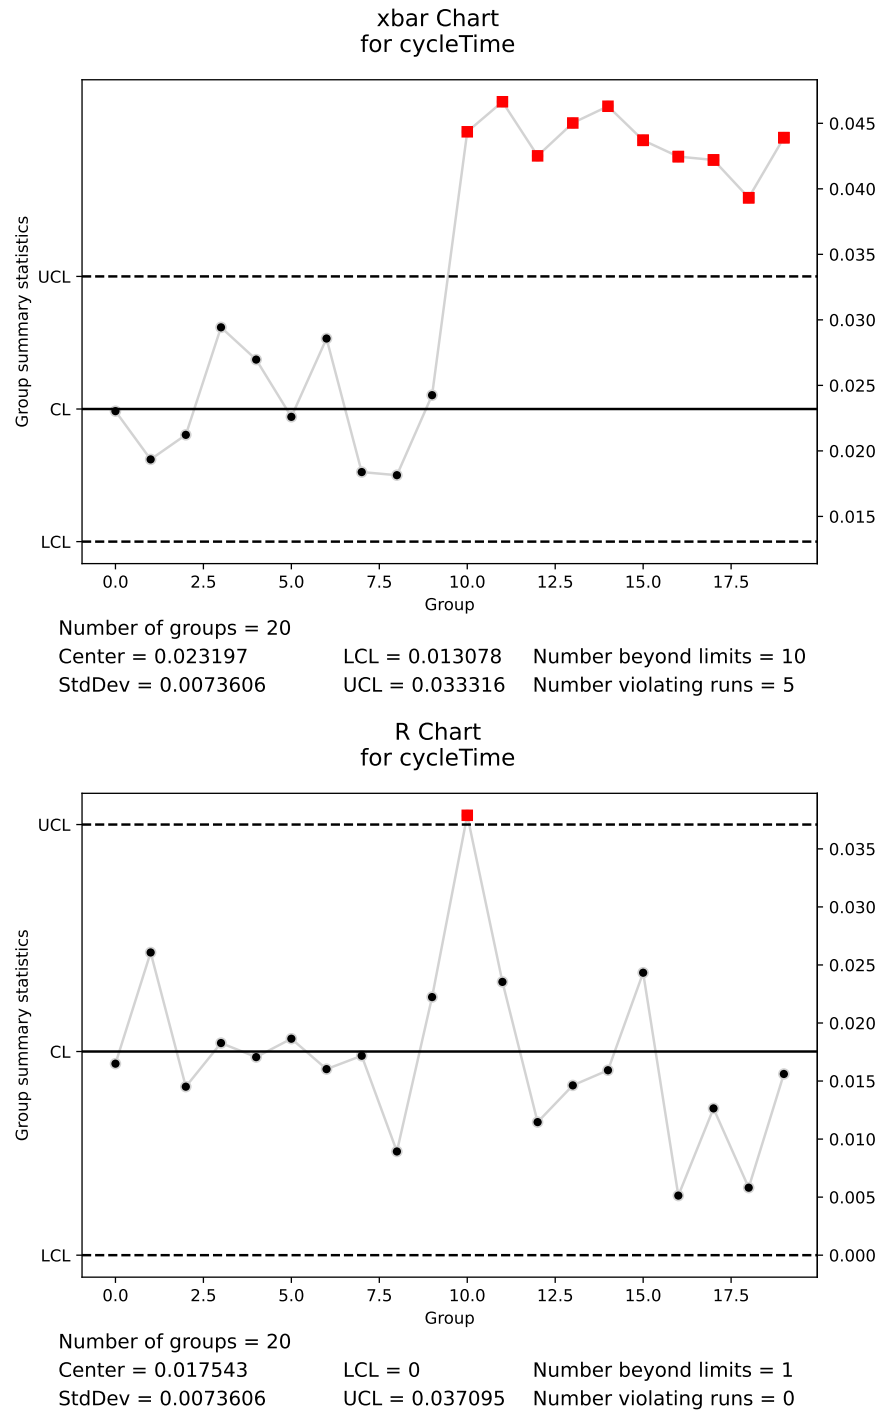


Fig. 2.6 Control charts for time shift Piston simulation (Exercise 2.14)

(fig:piston-time-shift)

```
# sort cycle times using by the order of the random numbers
# this randomizes the cycle times
Ps['seconds'] = Ps['seconds'][R.argsort()].values
# alternative versions
# Ps['seconds'] = Ps['seconds'].sample(frac=1).values
# Ps['seconds'] = [s for _, s in sorted(zip(R, Ps['seconds']))]

makeQCCplot(Ps, Ps, 'xbar', 'for cycleTime')
plt.show()
makeQCCplot(Ps, Ps, 'R', 'for cycleTime')
plt.show()
```

See the control charts in Fig. 2.7. Due to the randomization of the cycle times, the increased values are distributed over all groups and the mean cycle time increases by 0.01. The variability as shown in the R chart is increased too.

(fig:piston-randomization)

Exercise 2.15 Part I: Run the piston simulator by specifying the 7 piston parameters within their acceptable range. Record the 7 operating levels you used and generate 20 subgroups of size 5.

1. Compute the control limits for \bar{X} and S .

Part II: Rerun the piston simulator at the same operating conditions and generate 20 subgroups of size 10.

2. Recompute the control limits for \bar{X} and S .
3. Explain the differences between a) and b).

Solution 2.15 Part I (a): We run the simulation with all 7 control parameters at their lowest value. We obtained $\bar{\bar{X}} = 0.0234$, $\bar{S} = 0.0068$.

```
settings = {'m': 30, 's': 0.005, 'v0': 0.002, 'k': 1000,
            'p0': 90_000, 't': 290, 't0': 340}

simulator = mistat.PistonSimulator(n_simulation=20, n_replicate=5, seed=1, **settings)
Ps = simulator.simulate()
data = mistat.qcc_groups(Ps['seconds'], Ps['group'])
for qcc_type in ('xbar', 'S'):
    qcc = mistat.QualityControlChart(data, qcc_type=qcc_type)
    print(f'{qcc_type:4s} Center {qcc.center:.4f} Control limits ' +
          f'[{qcc.limits.LCL[0]:.5f}, {qcc.limits.UCL[0]:.4f}]')
```

```
xbar Center 0.0234 Control limits [0.01354, 0.0333]
S      Center 0.0068 Control limits [0.00000, 0.0141]
```

Part II (b): We obtained $\bar{\bar{X}} = 0.0253$, $\bar{S} = 0.0093$.

```
simulator = mistat.PistonSimulator(n_simulation=20, n_replicate=10, seed=1, **settings)
Ps = simulator.simulate()
data = mistat.qcc_groups(Ps['seconds'], Ps['group'])
for qcc_type in ('xbar', 'S'):
    qcc = mistat.QualityControlChart(data, qcc_type=qcc_type)
    print(f'{qcc_type:4s} Center {qcc.center:.4f} Control limits ' +
          f'[{qcc.limits.LCL[0]:.5f}, {qcc.limits.UCL[0]:.4f}]')
```

```
xbar Center 0.0253 Control limits [0.01623, 0.0344]
S      Center 0.0093 Control limits [0.00264, 0.0159]
```

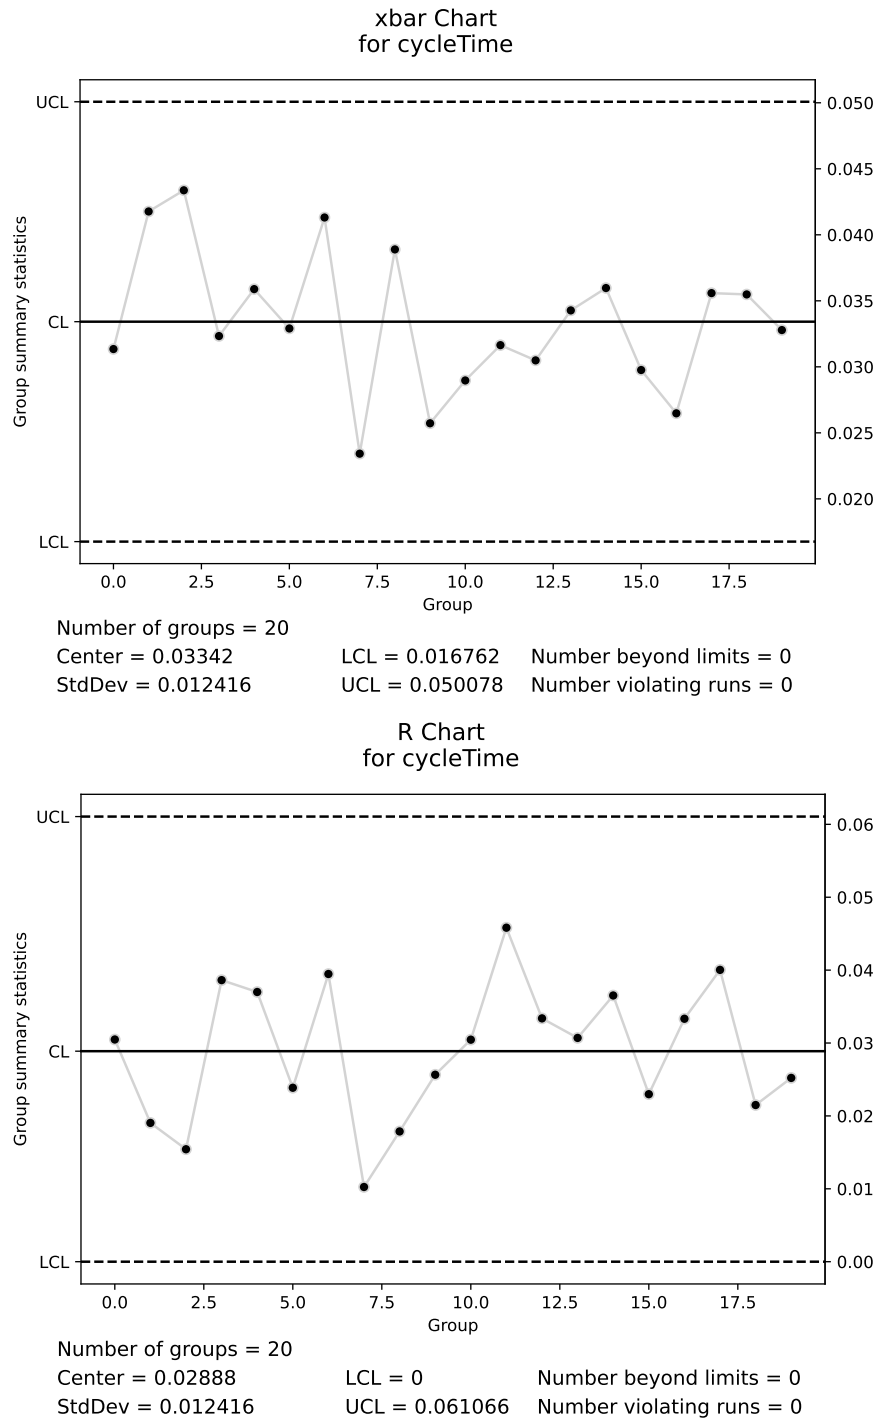


Fig. 2.7 Control charts for randomized cycle time in Piston simulation (Exercise 2.14)

(fig:piston-randomization)

(c) We see that the control limits of \bar{X}_{10} and of S_{10} are closer to the centerline than those of \bar{X}_5 and S_5 , respectively.

Exercise 2.16 Repeat the data segment analysis from Section 2.8.2 with a piecewise linear regression fit. This can be done using `pwlf.PiecewiseLinFit` initialized with the keyword argument `degree=1`.

{sec:segment-analysis-fda}

- (a) Prepare models with a variety of knots
- (b) Compare the results to the step function fits
- (c) When would you use a piecewise linear fit compared to a step function fit?

Solution 2.16 (a) Use code like the following to explore piecewise linear fits to the sensor data using different number of knots.

```
data = mistat.load_data('PROCESS_SEGMENT')

def sensorData(data, label):
    series = data[label]
    return pd.DataFrame({
        'Time': np.arange(len(series)),
        'values': series,
    })

sensorX = sensorData(data, 'X')
sensorZ = sensorData(data, 'Z')

def fitPiecewiseLinearFit(sensor, knots):
    model = pwlf.PiecewiseLinFit(sensor['Time'], sensor['values'], degree=1)
    model.fit(knots)
    return model

modelX = fitPiecewiseLinearFit(sensorX, 6)
modelZ = fitPiecewiseLinearFit(sensorZ, 3)

def plotPiecewiseLinearFit(sensor, model, ax, label):
    for bp in model.fit_breaks[1:-1]:
        ax.axvline(bp, color='lightgrey')
    ax.scatter(sensor['Time'], sensor['values'], color='grey', alpha=0.5)
    ax.plot(sensor['Time'], model.predict(sensor['Time']), color='black')
    ax.set_xlabel('Time')
    ax.set_ylabel(label)
    return ax

fig, axes = plt.subplots(ncols=2, figsize=(8,4))
plotPiecewiseLinearFit(sensorX, modelX, axes[0], 'Sensor X')
plotPiecewiseLinearFit(sensorZ, modelZ, axes[1], 'Sensor Z')
plt.tight_layout()
plt.show()
```

The figure shows results for sensor X and sensor Z data. By trying different number of segments, we can see that we require 6 segments for the sensor X data while 3 segments are sufficient for sensor Z. In contrast to the step function fit, the piecewise linear function doesn't enforce the linear segments to have a slope of 0. While visually, we can see that the segment that is characteristics for a process under control are essentially horizontal, it is useful to analyze the individual segments. The slope can be extracted from the model using the method `modelX.calc_slopes()`. We get:




Fig. 2.8 Piecewise linear fits of sensor X and sensor Z data.

{fig:sensor-pwlf-example}

Sensor X			
	Range	slope	prediction
Segment 1	Time < 894.3	35.30	26.3–4634.1
Segment 2	130.5 ≤ Time < 894.3	-0.02	4634.1–4619.8
Segment 3	894.3 ≤ Time < 910.7	32.89	4619.8–5158.9
Segment 4	910.7 ≤ Time < 1426.3	-0.43	5158.9–4936.1
Segment 5	1426.3 ≤ Time < 1502.6	-63.67	4936.1–77.9
Segment 6	1502.6 ≤ Time	-0.02	77.9–70.0

Sensor Z			
	Range	slope	prediction
Segment 1	Time < 831.1	38.97	125.2–4813.7
Segment 2	120.3 ≤ Time < 831.1	-0.06	4813.7–4770.3
Segment 3	831.1 ≤ Time	3.09	4770.3–5295.3

The slope of the longer second segment is almost zero for both sensor data. The segment analysis can therefore be used.

{sec:segment-analysis-fda}

(b) The step function fits of Section 2.8.2 required 9 segments for sensor X and 6 segments for sensor Z to describe the stable phase sufficiently. With the piecewise linear fit, we require fewer segments.

(c) The piecewise linear fit can model a linear ramp-up phase like we see at the beginning of both run charts better than a step function fit. It also allows us identifying and estimating drifts in the run charts.

Chapter 3

Advanced Methods of Statistical Process Control

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats
from scipy.special import factorial
import statsmodels.formula.api as smf
import mistat
import matplotlib.pyplot as plt
```

Exercise 3.1 Generate the distribution of the number of runs in a sample of size $n = 25$, if the number of elements above the sample mean is $m_2 = 10$.

- (i) What are Q_1 , Me and Q_3 of this distribution?
- (ii) Compute the expected value, μ_R , and the standard deviation σ_R .
- (iii) What is $\Pr\{10 \leq R \leq 16\}$?
- (iv) Determine the normal approximation to $\Pr\{10 \leq R \leq 16\}$.

Solution 3.1 First define a function to calculate the probability of runs:

```
from scipy.special import binom
def probRuns(m1, m2, R):
    n = m1 + m2
    k = R // 2
    if R % 2:
        denom = binom(m1-1, k-1) * binom(m2-1, k) + binom(m1-1, k) * binom(m2-1, k-1)
        return denom / binom(n, m2)
    else:
        return 2 * binom(m1-1, k-1) * binom(m2-1, k-1) / binom(n, m2)
```

For the case of $n = 25$ and $m_2 = 10$ calculate the distribution

```
n = 25
m2 = 10
m1 = n - m2

df = pd.DataFrame({
    'R': range(0, n+1),
    'p.d.f': [probRuns(m1, m2, R) for R in range(0, n+1)],
})
```

```

})
df['c.d.f'] = np.cumsum(df['p.d.f'])

```

(a) We can determine the median and quantiles from the calculated c.d.f

```

cdf = df['c.d.f'].values

Q1 = np.where(cdf < 0.25)[0][-1]
Me = np.where(cdf < 0.5)[0][-1]
Q3 = np.where(cdf < 0.75)[0][-1]
print(Q1, Me, Q3)

```

```
| 10 12 14
```

$Q_1 = 10$, $M_e = 12$ and $Q_3 = 14$.

(b) From Eq. (9.1.3) and Eq. (9.1.4) we get $\mu_R = 13$ and $\sigma_R = 2.3452$.

```

mu_R = 1 + 2 * m1 * m2 / n
sigma_R = np.sqrt(2*m1*m2*(2*m1*m2 - n) / (n*n*(n-1)))
print(mu_R, sigma_R)

```

```
| 13.0 2.345207879911715
```

(c) $\Pr\{10 \leq R \leq 16\} = 0.8657$.

```

p = df['c.d.f'][16] - df['c.d.f'][9]
print(p)

```

```
| 0.8656750572082381
```

(d) Using the normal approximation $\Pr\{10 \leq R \leq 16\} \approx 0.8644$. Note that we set the limits to [9.5, 16.5].

```

print(stats.norm.cdf((16.5-mu_R)/sigma_R) - stats.norm.cdf((9.5-mu_R)/sigma_R))

```

```
| 0.8644069987336978
```

Exercise 3.2 Use Python to perform a run test on the simulated cycle times from the pistons, which are in data file **CYCLT.csv**. Is the number of runs above the mean cycle time significantly different than its expected value?

Solution 3.2 Load the data and determine number of runs information:

```

data = mistat.load_data('CYCLT')
# convert to up (1) or down (0) information relative to mean
mean_ct = np.mean(data)
runs = [1 if ct > mean_ct else 0 for ct in data]

# determine number of runs
obs_Runs = 0
current = None
for r in runs:
    if r != current:
        obs_Runs += 1

```

```

        current = r
    print(f'Observed number of runs: {obs_Runs}')

    # calculate expected number of runs
    m1 = sum(data > mean_ct)
    m2 = sum(data <= mean_ct)
    n = m1 + m2
    mu_R = 1 + 2 * m1 * m2 / n
    print(f'Expected number of runs {mu_R:.2f}')

    # determine if difference is significant
    mistat.runsTest(data, cutoff=np.mean(data))

```

```

| Observed number of runs: 26
| Expected number of runs 25.64

```

```

| Result(statistic=0.1044134517056721, pval=0.9168412481142088,
| method='Runs Test', alternative='two.sided')

```

The test shows no significant deviations from randomness.

Exercise 3.3

(exc:cyclt-data-R)

- (i) What is the expected number of runs up or down, in a sample of size 50?
- (ii) Compute the number of runs up or down in the cycle time data (CYCLT.csv).
- (iii) Is this number significantly different than expected?
- (iv) What is the probability that a random sample of size 50 will have at least one run of size greater or equal to 5?

Solution 3.3 (i) For $n = 50$, $E\{R^*\} = \frac{2n-1}{3} = 33$.

```

n = 50
mu_Rstar = (2*n-1)/3
print(mu_Rstar)

```

```

| 33.0

```

(ii) Using the cycle time data, $R^* = 34$.

```

# determine direction of change up (1) or down (-1)
y = [1 if xi < xipl else -1 for xi, xipl in zip(data[:-1], data[1:])]

# count number of up and down runs
up = 0
down = 0
current = None
for yi in y:
    if yi == current: # no change of direction
        continue
    if yi < 0:
        down += 1
    else:
        up += 1
    current = yi
Rstar = up + down
print(Rstar, up, down)

```

```

| 34 17 17

```

(iii) We have $\sigma^* = 2.9269$, $\alpha_u^* = 1 - \Phi\left(\frac{1}{2.9269}\right) = 0.3663$. The deviation is not significant.

```
n = 50
mu_Rstar = (2*n-1)/3
sigma_Rstar = np.sqrt((16*n-29)/90)
print(sigma_Rstar)
alpha_L = stats.norm.cdf((Rstar-mu_Rstar)/sigma_Rstar)
alpha_U = 1 - alpha_L
print(alpha_U, alpha_L)
```

```
2.9268868558020253
0.3663034098961011 0.6336965901038989
```

(iv) The probability that a sample of size 50 will have at least one run of length 5 or longer is 0.102.

```
def expected_R_k(n, k):
    return 2 * (n*(k+1) - k*k - k + 1)/factorial(k+2)

print(expected_R_k(50, 5))
# probability to have run greater than 5
print(1 - np.exp(-expected_R_k(50, 5)))
```

```
0.10753968253968255
0.10195911479934461
```

The function `mistat.runStatistics` calculates a variety of statistics for runs.

```
mistat.runStatistics(data)
```

```
{'count': {'mu_R': 25.64, 'sigma_R': 3.4478316167038754, 'observed':
26},
'direction': {'mu_Rstar': 33.0,
'sigma_Rstar': 2.9268868558020253,
'up': 17,
'down': 17,
'Rstar': 34,
'alpha': [0.3663034098961011, 0.6336965901038989]}}
```

Exercise 3.4 Analyze the observations in `YARNSTRG.csv` for runs.

Solution 3.4 We make use of the `mistat.runStatistics` function.

```
data = mistat.load_data('YARNSTRG')
mistat.runStatistics(data)
```

```
{'count': {'mu_R': 50.92, 'sigma_R': 4.966642668235696, 'observed':
49},
'direction': {'mu_Rstar': 66.33333333333333,
'sigma_Rstar': 4.1779846284489315,
'up': 32,
'down': 32,
'Rstar': 64,
'alpha': [0.28825730949895256, 0.7117426905010474]}}
```

```
mistat.runsTest(data, cutoff=data.mean())
```

```
Result (statistic=-0.3865790491189181, pval=0.6990678707195341,
method='Runs Test', alternative='two.sided')
```

{exc:Xbar-S-piston-simulator-top}

Exercise 3.5 Run the piston simulator at the upper level of the seven control parameters and generate 50 samples of size 5. Analyze the output for runs in both \bar{X} and S charts.

Solution 3.5 Define the parameter settings and create 50 simulations with 5 replications each

```
parameter = pd.DataFrame({
    'm': [60],
    's': [0.02],
    'k': [5_000],
    't': [296],
    'p0': [110_000],
    'v0': [0.01],
    't0': [360],
})
simulator = mistat.PistonSimulator(parameter=parameter, n_simulation=50, n_replicate=5, seed=1236)
# simulator = mistat.PistonSimulator(n_simulation=50, n_replicate=5, seed=1)
Ps = simulator.simulate()

# get grouped cycle times
cycleTime = mistat.qcc_groups(Ps['seconds'], Ps['group'])
mistat.runsTest(np.mean(cycleTime, axis=1), np.mean(cycleTime))
```

```
Result (statistic=-0.5610330087428953, pval=0.5747750350233831,
method='Runs Test', alternative='two.sided')
```

The test shows no significant runs. The resulting \bar{X} and S charts are shown in Fig. 3.1.

{fig:Xbar-S-piston-simulator-top}

```
qcc = mistat.QualityControlChart(cycleTime, qcc_type='xbar')
ax = qcc.plot(title='for piston simulation at upper level')
plt.show()
qcc = mistat.QualityControlChart(cycleTime, qcc_type='S')
ax = qcc.plot(title='for piston simulation at upper level')
plt.show()
```

Exercise 3.6

{exc:Xbar-S-piston-simulator-T-change}

- (i) Run the piston simulator at the upper level of the seven control parameters and generate 50 samples of size 5 (both \bar{X} and S charts).
- (ii) Repeat the exercise allowing T to change over time (provide a list of T which specify the changing ambient temperature over time)
- (iii) Compare the results in (i) and (ii). with those of Exercise 3.3.

{exc:cyclt-data-R}

Solution 3.6 (i) See solution of Exercise 3.5 for this part of the exercise. Fig. 3.1 shows the resulting control charts.

{fig:Xbar-S-piston-simulator-top}

(ii) In simulating the piston cycle times, the ambient temperature around the piston is increased 10% per cycle after the shift point, which is after the 16th sample.




Fig. 3.1 \bar{X} and S chart of simulated Piston data (Exercises 6.5 and 6.6)

{fig:Xbar-S-piston-simulator-top}

```

parameter = pd.DataFrame({
    'm': [60]*50,
    's': [0.02]*50,
    'k': [5_000]*50,
    't': [296] * 16 + [296 * 1.1**i for i in range(1, 35)],
    'p0': [110_000]*50,
    'v0': [0.01]*50,
    't0': [360]*50,
})

simulator = mistat.PistonSimulator(parameter=parameter, n_simulation=50,
                                   n_replicate=5, seed=1, check=False)
Ps = simulator.simulate()
cycleTimeTshift = mistat.qcc_groups(Ps['seconds'], Ps['group'])

```

```

qcc = mistat.QualityControlChart(cycleTimeTshift, qcc_type='xbar')
ax = qcc.plot(title='for contact lens data')
plt.show()
qcc = mistat.QualityControlChart(cycleTimeTshift, qcc_type='S')
ax = qcc.plot(title='for contact lens data')
plt.show()

```

The control charts for this simulation are shown in Fig. 3.2. Both show clearly unrandom behavior.

{fig:Xbar-S-piston-simulator-T-change}

(iii) We calculate runs statistics for both simulations.

```

mistat.runStatistics(np.mean(cycleTime, axis=1))

```

```

{'count': {'mu_R': 25.96, 'sigma_R': 3.493555583105112, 'observed':
24},
 'direction': {'mu_Rstar': 33.0,
 'sigma_Rstar': 2.9268868558020253,
 'up': 16,
 'down': 16,
 'Rstar': 32,
 'alpha': [0.36630340989610105, 0.6336965901038989]}}

```

```

mistat.runStatistics(np.mean(cycleTimeTshift, axis=1))

```

```

{'count': {'mu_R': 25.96, 'sigma_R': 3.493555583105112, 'observed':
8},
 'direction': {'mu_Rstar': 33.0,
 'sigma_Rstar': 2.9268868558020253,
 'up': 15,
 'down': 14,
 'Rstar': 29,
 'alpha': [0.08586912144778469, 0.9141308785522153]}}

```

The number of runs in the first simulation is within the expected range. For the second simulation we observe a much smaller number of runs. The statistics for the direction of runs, shows no strong deviation from expected values.

```

mistat.runStatistics(np.mean(cycleTime, axis=1))

```

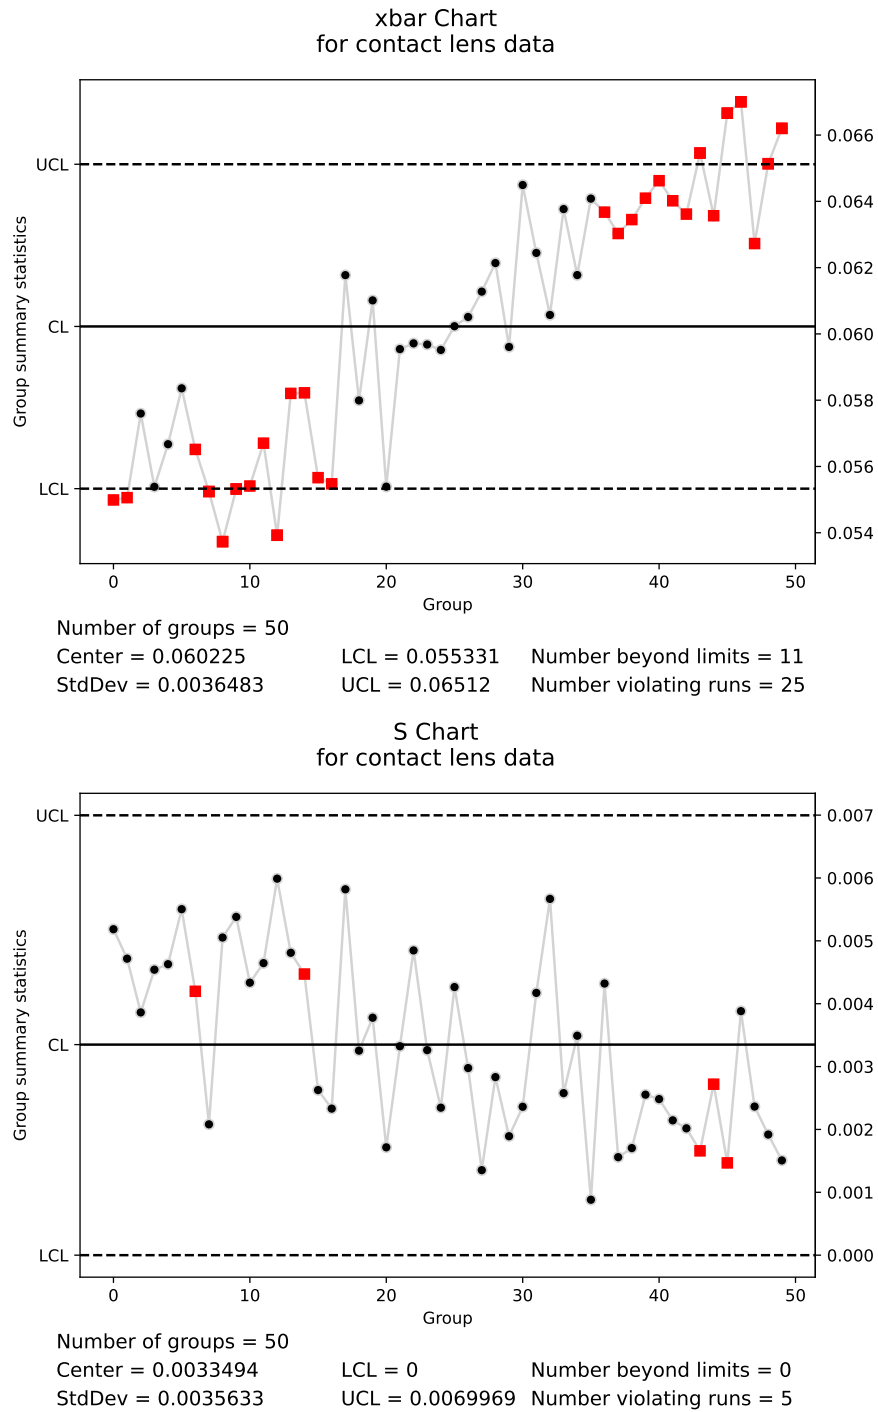


Fig. 3.2 \bar{X} and S chart of simulated Piston data (Exercise 3.6)

fig-Xbar-S-piston-simulator-T-change

{fig:Xbar-S-piston-simulator-T-change}

```
{'count': {'mu_R': 25.96, 'sigma_R': 3.493555583105112, 'observed':
24},
'direction': {'mu_Rstar': 33.0,
'sigma_Rstar': 2.9268868558020253,
'up': 16,
'down': 16,
'Rstar': 32,
'alpha': [0.36630340989610105, 0.6336965901038989]}}
```

```
mistat.runStatistics(np.mean(cycleTimeTshift, axis=1))
```

```
{'count': {'mu_R': 25.96, 'sigma_R': 3.493555583105112, 'observed':
8},
'direction': {'mu_Rstar': 33.0,
'sigma_Rstar': 2.9268868558020253,
'up': 15,
'down': 14,
'Rstar': 29,
'alpha': [0.08586912144778469, 0.9141308785522153]}}
```

We can also perform runs tests for both simulations. For the first simulation we get:

```
print('mean:', mistat.runsTest(np.mean(cycleTime, axis=1), np.mean(cycleTime)).pval)
STD = np.std(cycleTime, axis=1)
print('std:', mistat.runsTest(STD, np.mean(STD)).pval)
```

```
mean: 0.5747750350233831
std: 0.2529990614746843
```

This confirms our observation that the simulation shows no non-random behavior. For the second simulation we get:

```
print('mean:', mistat.runsTest(np.mean(cycleTimeTshift, axis=1), np.mean(cycleTimeTshift)).pval)
STD = np.std(cycleTimeTshift, axis=1)
print('std:', mistat.runsTest(STD, np.mean(STD)).pval)
```

```
mean: 2.7343384749349e-07
std: 0.09297787599834818
```

The run tests reflect the nonrandomness of the sequence, after the change point.

Exercise 3.7 Construct a p -chart for the fraction of defective substrates received at a particular point in the production line. One thousand ($n = 1000$) substrates are sampled each week. Remove data for any week for which the process is not in control. Be sure to check for runs as well as points outside the control limits. Construct the revised p -chart and be sure to check for runs again. The data are in Table 3.1.

(exc:p-def-subst)

Solution 3.7 Construct the np -chart of the data using `mistat.QualityControlChart`. It is shown in Fig. 3.3.

(tbl:data-exc-p-def-subst)

(fig:exc-np-chart-defects)


```
data = pd.Series([18, 14, 9, 25, 27, 18, 21, 16, 18, 24, 20, 19, 22, 22, 20,
38, 29, 35, 24, 20, 23, 17, 20, 19, 17, 16, 10, 8, 10, 9])

qcc = mistat.QualityControlChart(data, qcc_type='np', sizes=1000)
ax = qcc.plot(title='for defects')
plt.show()
```

Table 3.1 Data set for Exercise 3.7

Week	No. Def.	Week	No. Def.
1	18	16	38
2	14	17	29
3	9	18	35
4	25	19	24
5	27	20	20
6	18	21	23
7	21	22	17
8	16	23	20
9	18	24	19
10	24	25	17
11	20	26	16
12	19	27	10
13	22	28	8
14	22	29	10
15	20	30	9

(tbl:data-exc-p-def-subst)

**Fig. 3.3** np-Chart for defective substrates

(fig:exc-np-chart-defects)




Fig. 3.4 Revised *np*-chart for defective substrates

{fig:exc-np-chart-defects-revised}

As we see, the data points of week 16 and 18 fall outside the upper control limit. Also, the last 8 weeks show a significant trend down (improvement).

After removing the data points for week 16 and 18 (note that Python arrays are 0-indexed), we get a revised *np*-chart (Fig. 3.4). The points are now all in control, but the pattern of the last 8 weeks remains.

{fig:exc-np-chart-defects-revised}

```
revised = data[data < 32.751]

qcc = mistat.QualityControlChart(revised, qcc_type='np', sizes=1000)
ax = qcc.plot(title='for number of defects')
plt.show()
```

{exc:two-prod-lines}

Exercise 3.8 Substrates were inspected for defects on a weekly basis, on two different production lines. The weekly sample sizes and the number of defectives are indicated below in the data set in Table 3.2. Plot the data and indicate which of the lines is not in a state of statistical control. On what basis do you make your decision? Use Python to construct control charts for the two production lines.

{tbl:data-two-prod-lines}

Note: When the sample size is not the same for each sampling period, we use **variable** control limits. If $X(i)$ and $n(i)$ represent the number of defects and sample size, respectively, for sampling period i , then the upper and lower control limits for the i -th period are

$$UCL_i = \bar{p} + 3(\bar{p}(1 - \bar{p})/n_i)^{1/2}$$

and

$$LCL_i = \bar{p} - 3(\bar{p}(1 - \bar{p})/n_i)^{1/2}$$

where

$$\bar{p} = \sum X(i) / \sum n(i)$$

is the centerline for the control chart.

(tbl:data-two-prod-lines)

Table 3.2 Data set for Exercise 3.8 ^{exc:two-prod-lines}

	Line 1		Line 2	
Week	X_i	n_i	X_i	n_i
1	45	7920	135	2640
2	72	6660	142	2160
3	25	6480	16	240
4	25	4500	5	120
5	33	5840	150	2760
6	35	7020	156	2640
7	42	6840	140	2760
8	35	8460	160	2980
9	50	7020	195	2880
10	55	9900	132	2160
11	26	9180	76	1560
12	22	7200	85	1680

Solution 3.8 The average proportion of defectives for line 1 is $\hat{p}_1 = 0.005343$, while that for line 2 is $\hat{p}_2 = 0.056631$. The difference is very significant. Indeed,

$$Z = \frac{\hat{p}_2 - \hat{p}_1}{\sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{N_1} + \frac{\hat{p}_2(1-\hat{p}_2)}{N_2}}} = 34.31,$$

where $N_1 = \sum_{i=1}^{12} n_1(i)$ and $N_2 = \sum_{i=1}^{12} n_2(i)$.

(fig:exc-production-lines)

In Fig. 3.5 we present the p -charts for these two production lines. We see that the chart for line 1 reveals that the process was at the beginning (point 2), out of control.

```
data = pd.DataFrame([
    [1, 45, 7920, 135, 2640],
    [2, 72, 6660, 142, 2160],
    [3, 25, 6480, 16, 240],
    [4, 25, 4500, 5, 120],
    [5, 33, 5840, 150, 2760],
    [6, 35, 7020, 156, 2640],
    [7, 42, 6840, 140, 2760],
    [8, 35, 8460, 160, 2980],
    [9, 50, 7020, 195, 2880],
    [10, 55, 9900, 132, 2160],
    [11, 26, 9180, 76, 1560],
    [12, 22, 7200, 85, 1680]],
    columns=['Week', 'Line 1 X', 'Line 1 n', 'Line 2 X', 'Line 2 n']
)

qcc = mistat.QualityControlChart(data['Line 1 X'], qcc_type='p', sizes=data['Line 1 n'])
ax = qcc.plot(title='for line 1 defects')
plt.show()
```

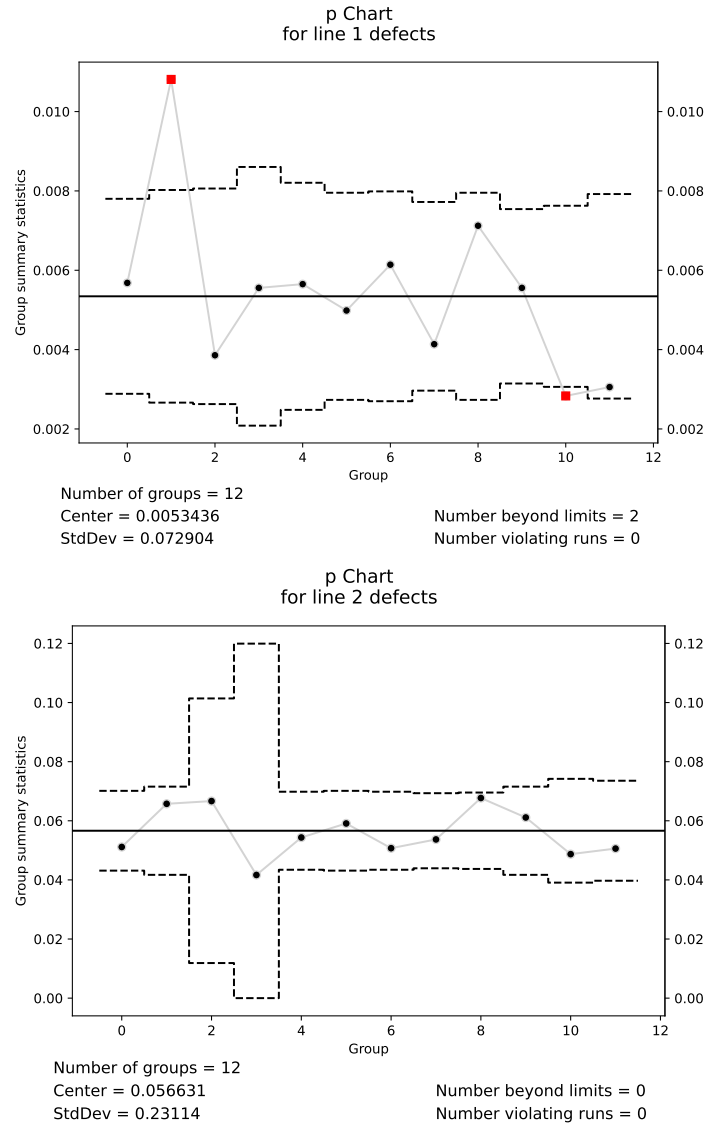


Fig. 3.5 *p*-Charts for production line defects (Exercise 6.8)

```
qcc = mistat.QualityControlChart(data['Line 2 X'], qcc_type='p', sizes=data['Line 2 n'])
ax = qcc.plot(title='for line 2 defects')
plt.show()
```

Exercise 3.9 In designing a control chart for the fraction defectives p , a random sample of size n is drawn from the productions of each day (very large lot). How

large should n be so that the probability of detecting a shift from $p_0 = .01$ to $p_t = .05$, within a 5 day period, will not be smaller than .8?

Solution 3.9 Let θ be the probability of not detecting the shift in a given day. Solving $\theta^5 = 0.2$ we get $\theta = 0.72478$. We find the smallest n for which

$$\theta \geq \Phi\left(\frac{0.01 + 3\sqrt{\frac{0.01 \times 0.99}{n}} - 0.05}{\sqrt{\frac{0.05 \times 0.95}{n}}}\right) - \Phi\left(\frac{0.01 - 3\sqrt{\frac{0.01 \times 0.99}{n}} - 0.05}{\sqrt{\frac{0.05 \times 0.95}{n}}}\right).$$

The solution is $n = 16$. Using Eq. (3.3.7) we get $n \approx 18$. This solution and the one shown above both use the normal approximation to the binomial. Another approach, which may be preferable for small sample sizes, is to use binomial distribution directly. We find the smallest n for which

$$B(n \times 0.01 + 3\sqrt{n \times 0.01 \times 0.99}; n, 0.05) - B(n \times 0.01 - 3\sqrt{n \times 0.01 \times 0.99}; n, 0.05) \leq \theta.$$

In this case the solution is $n = 7$.

{tbl:dock-stock-cycle} **Table 3.3** Dock to Stock Cycle Times

Day	Times				
1	27	43	49	32	36
2	34	29	34	31	41
3	36	32	48	35	33
4	31	41	51	51	34
5	43	35	30	32	31
6	28	42	35	40	37
7	38	37	41	34	44
8	28	44	44	34	50
9	44	36	38	44	35
10	30	43	37	29	32
11	36	40	50	37	43
12	35	36	44	34	32
13	48	49	44	27	32
14	45	46	40	35	33
15	38	36	43	38	34
16	42	37	40	42	42
17	44	31	36	42	39
18	32	28	42	39	27
19	41	41	35	41	44
20	44	34	39	30	37
21	51	43	36	50	54
22	52	50	50	44	49
23	52	34	38	41	37
24	40	41	40	23	30
25	34	38	39	35	33

[exc:dock-stock-cycle]

Exercise 3.10 The data in Table 3.3 represent dock-to-stock cycle times for a certain type of shipment (class *D*). Incoming shipments are classified according to their ‘type’, which is determined by the size of the item and the shipment, the type of handling required, and the destination of the shipment. Samples of five shipments per day are tracked from their initial arrival to their final destination, and the time it takes for this cycle to be complete is noted. The samples are selected as follows: at five preselected times during the day, the next class *D* shipment to arrive is tagged and the arrival time and identity of the shipment are recorded. When the shipment reaches its final destination, the time is again recorded. The difference between these times is the cycle time. The cycle time is always recorded for the day of arrival.

[tbl:dock-stock-cycle]

- (i) Construct \bar{X} and *S*-charts from the data. Are any points out of control? Are there any trends in the data? If there are points beyond the control limits, assume that we can determine special causes for the points, and recalculate the control limits, excluding those points which are outside the control limits.
- (ii) Use a *t*-test to decide whether the mean cycle time for days 21 and 22 was significantly greater than 45.
- (iii) Make some conjectures about possible causes of unusually long cycle times. Can you think of other appropriate data that might have been collected, such as the times at which the shipments reached intermediate points in the cycle? Why would such data be useful?

Solution 3.10 (i) In Fig. 3.6 we see the \bar{X} and *S* control charts. Points for days 21 and 22 are outside the UCL for \bar{X} . Excluding these points, the recalculated control chart limits for \bar{X} are LCL = 29.731 and UCL = 45.781. Also, $\bar{\bar{X}} = 37.757$.

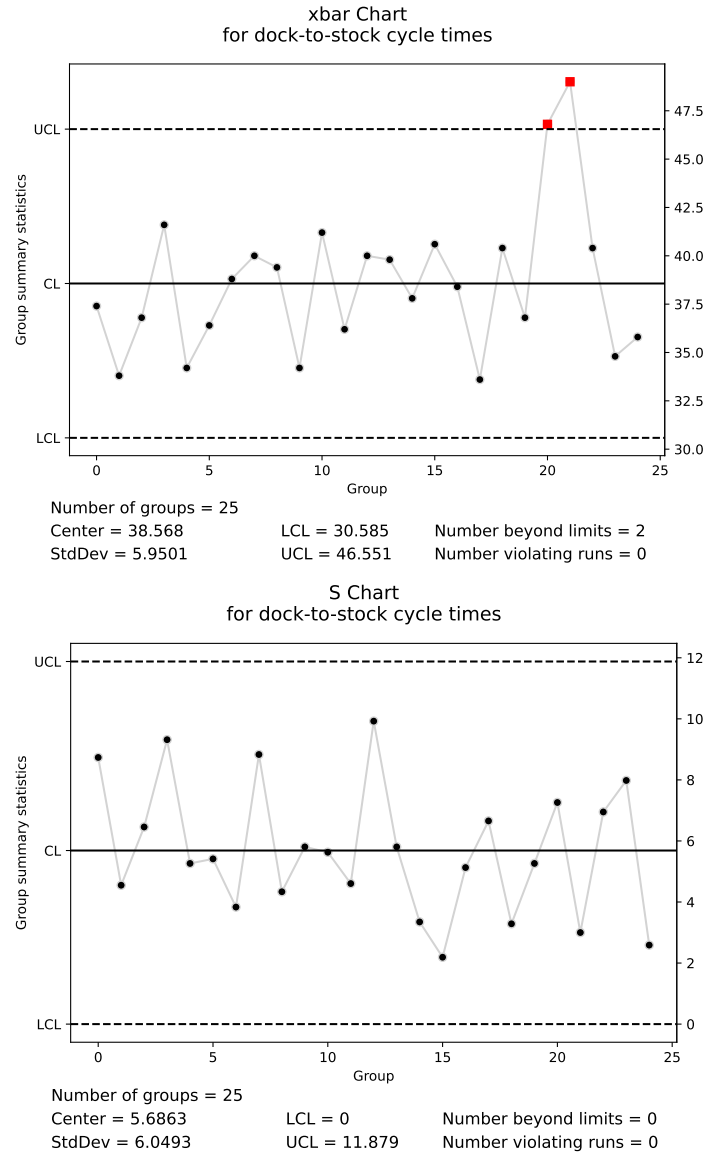
[fig:dock-stock-cycle]

```
cycleTime = pd.DataFrame([
    [1, 27, 43, 49, 32, 36],      [2, 34, 29, 34, 31, 41],
    [3, 36, 32, 48, 35, 33],      [4, 31, 41, 51, 51, 34],
    [5, 43, 35, 30, 32, 31],      [6, 28, 42, 35, 40, 37],
    [7, 38, 37, 41, 34, 44],      [8, 28, 44, 44, 34, 50],
    [9, 44, 36, 38, 44, 35],      [10, 30, 43, 37, 29, 32],
    [11, 36, 40, 50, 37, 43],      [12, 35, 36, 44, 34, 32],
    [13, 48, 49, 44, 27, 32],      [14, 45, 46, 40, 35, 33],
    [15, 38, 36, 43, 38, 34],      [16, 42, 37, 40, 42, 42],
    [17, 44, 31, 36, 42, 39],      [18, 32, 28, 42, 39, 27],
    [19, 41, 41, 35, 41, 44],      [20, 44, 34, 39, 30, 37],
    [21, 51, 43, 36, 50, 54],      [22, 52, 50, 50, 44, 49],
    [23, 52, 34, 38, 41, 37],      [24, 40, 41, 40, 23, 30],
    [25, 34, 38, 39, 35, 33]],
    columns=['Week', 'S1', 'S2', 'S3', 'S4', 'S5'])
cycleTime = cycleTime.set_index('Week')

qcc = mistat.QualityControlChart(cycleTime, qcc_type='xbar')
ax = qcc.plot(title='for dock-to-stock cycle times')
plt.show()

qcc = mistat.QualityControlChart(cycleTime, qcc_type='s')
ax = qcc.plot(title='for dock-to-stock cycle times')
plt.show()

# exclude points for week 21 and 22
qcc = mistat.QualityControlChart(cycleTime.drop(labels=[21, 22]), qcc_type='xbar')
qcc.center, qcc.limits
```



(fig:dock-stock-cycle) **Fig. 3.6** Xbar and S chart of dock to stock cycle times (Exercise 3.10) *fig:dock-stock-cycle*

```
(37.756521739130434,  
      LCL      UCL  
0 29.731454 45.781589)
```

(ii) We can calculate the significance of the day 21 and 22 cycle times relative to the UCL of the recalculated control chart.

```

statistic, pvalue = stats.ttest_1samp(cycleTime.loc[21,], 45.781, alternative='greater')
print(pvalue)
statistic, pvalue = stats.ttest_1samp(cycleTime.loc[22,], 45.781, alternative='greater')
print(pvalue)

```

```

0.3846525285951686
0.03720636409199298

```

We find that only \bar{X}_{22} is significantly larger than 45.781.

(iii) Unusual long cycle times can be due to

1. Missing or misplaced information in accompanying paperwork
2. Missing or misplaced marks on package
3. Defective package
4. Non standard package
5. Wrong information on package destination
6. Overloaded stock room so that packages cannot be accepted
7. Misplaced packages.

Additional data that can be collected to explain long cycle times:

1. Package destination
2. Stock room
3. Package size
4. Package weight
5. Package origin (country, supplier)
6. Package courier.

Exercise 3.11 Consider the modified Shewhart control chart for sample means, with $a = 3$, $w = 2$ and $r = 4$. What is the ARL of this procedure when $\delta = 0, 1, 2$ and the sample size is $n = 10$?

Solution 3.11¹ The ARL of the modified Shewhart control chart with $a = 3$, $w = 2$ and $r = 4$ when $n = 10$ and $\delta = 0$ is 370.3. When $\delta = 1$, the ARL is 1.75 and when $\delta = 2$ the ARL is 1.0.

Exercise 3.12 Repeat the previous exercise for $a = 3$, $w = 1$, $r = 15$; when $n = 5$ and $\delta = 0.5$.

Solution 3.12² **9.2.2** For $a = 3$, $w = 1$ and $r = 15$ when $n = 5$ the ARL is 33.4.

Exercise 3.13 Suppose that a shift in the mean is occurring at random, according to an exponential distribution with mean of 1 hour. The hourly cost is \$100 per shift of size $\delta = \frac{\mu_1 - \mu_0}{\sigma}$. The cost of sampling and testing is $d = \$10$ per item. How often should samples of size $n = 5$ be taken, when lifts³ of size $\delta \geq 1.5$ should be detected?

¹ **TODO:** require clarification on how to calculate this

² **TODO:** require clarification on how to calculate this

³ **TODO:** shifts?

Solution 3.13 ⁴ Samples should be taken every $h^0 = 34 \text{ [min]} = 0.57 \text{ [hr]}$.

Exercise 3.14 Compute the $OC(p)$ function, for a Shewhart 3-sigma control chart for p , based on samples of size $n = 20$, when $p_0 = 0.10$. [Use the formula for exact computations.]

Solution 3.14 With $n = 20$ and $p_0 = 0.10$,

$$OC(p) = B(np_0 + 3\sqrt{np_0(1-p_0)}; n, p) - B(np_0 - 3\sqrt{np_0(1-p_0)}; n, p).$$

we get the following table of values of $OC(p)$, for $p = 0.05, 0.5 (0.05)$.

p	$OC(p)$
0.05	0.999966
0.10	0.997614
0.15	0.978065
0.20	0.913307
0.25	0.785782
0.30	0.608010
0.35	0.416625
0.40	0.250011
0.45	0.129934
0.50	0.057659

Exercise 3.15 How large should the sample size n be, for a 3-sigma control chart for p , if we wish that the probability of detecting a shift from $p_0 = 0.01$ to $p_t = 0.05$ be $1 - \beta = 0.90$?

Solution 3.15 The sample size n_0 is the smallest n for which

$$B(n \times 0.01 + 0.2985\sqrt{n}; n, 0.05) - B(n \times 0.01 - 0.2985\sqrt{n}; n, 0.05) \leq 0.1.$$

From this we obtain $n_0 = 184$.

```
n = 1
while OC_p_chart(0.05, n, 0.01) > 0.1:
    n += 1
n
```

| 184

Using Eq. (3.3.5), which is based on the normal approximation to the binomial, gives $n_0 \approx 209$.

```
# normal approximation
def OC_p_chart_normal(p, n, p0):
    loc = n * p
```

⁴ **TODO:** require clarification on how to calculate this

```

scale = np.sqrt(n*p*(1-p))
c = n * p0
delta = 3 * np.sqrt(n * p0 * (1 - p0))
return (stats.norm(loc, scale).cdf(c + delta) -
        stats.norm(loc, scale).cdf(c - delta))

# alternative implementation
def OC_p_chart_normal_2(p, n, p0):
    delta = 3 * np.sqrt(p0 * (1 - p0) / n)
    UCL = p0 + delta
    LCL = p0 - delta
    denom = np.sqrt(p * (1 - p) / n)
    return (stats.norm().cdf((UCL - p)/denom) -
            stats.norm().cdf((LCL - p)/denom))

n = 1
while OC_p_chart_normal(0.05, n, 0.01) > 0.1:
    n += 1
print(n, OC_p_chart_normal(0.05, n, 0.01))

```

| 209 0.09959381959080052

Exercise 3.16 Suppose that a measurement X , of hardness of brackets after heat treatment, has a normal distribution. Every hour a sample of n units is drawn and a \bar{X} chart with control limits $\mu_0 \pm 3\sigma/\sqrt{n}$ is used. Here μ_0 and σ are the assumed process mean and standard deviation. The OC function is

$$OC(\delta) = \Phi(3 - \delta\sqrt{n}) + \Phi(3 + \delta\sqrt{n}) - 1$$

where $\delta = (\mu - \mu_0)/\sigma$ is the standardized deviation of the true process mean from the assumed one.


- (i) How many hours, on the average, would it take to detect a shift in the process mean of size $\delta = 1$, when $n = 5$?
- (ii) What should be the smallest sample size, n , so that a shift in the mean of size $\delta = 1$ would be on the average detected in less than 3 hours?
- (iii) One has two options: to sample $n_1 = 5$ elements every hour or to sample $n_2 = 10$ elements every two hours. Which one would you choose? State your criterion for choosing between the two options and make the necessary computations.

Solution 3.16 (i) 4.5 [hr]; (ii) 7; (iii) For a shift of size $\delta = 1$, option 1 detects it on the average after 4.5 [hr]. Option 2 detects it on the average after $2 \times 1.77 = 3.5$ [hr]. Option 2 is preferred from the point of view of detection speed.

Exercise 3.17 Electric circuits are designed to have an output of 220 (volts, DC). If the mean output is above 222 (volts DC) you wish to detect such a shift as soon as possible. Examine the sample of data set **OEELECT.csv** for such a shift. For this purpose construct a CUSUM upward scheme with K^+ and h^+ properly designed (consider for h^+ the value $\alpha = .001$). Each observation is of sample of size $n = 1$. Is there an indication of a shift in the mean?

Solution 3.17⁵ **9.4.1** Here $K^+ = 221$ and $h^+ = 55.3724$. Figure 9.4.1 shows the “up” and “down” CUSUM charts. We see that the upper or lower limits are not crossed. There is no signal of change.

```
data = mistat.load_data('OELECT')
analysis = mistat.Cusum(data, center=220)
analysis.plot()
plt.show()
```



Number of groups = 99	Decision interval (std. err.) = 5
Center = 220	Shift detection (std. err.) 1
StdDev = 3.5914	No. of points beyond limits = 8

Exercise 3.18 Estimate the probability of false alarm and the conditional expected delay in the Poisson case, with a CUSUM scheme. The parameters are $\lambda_0 = 15$ and $\lambda_1 = 25$. Use $\alpha = .001$, $\tau = 30$.

Solution 3.18⁶ Using $K^+ = 19.576$, $h^+ = 13.523$, $K^- = 10.497$ and $h^- = -9.064$ we obtained $PFA = 0.014$ and $CED = 3.01 \pm 0.235$. Here the value of λ changed from 15 to 25.

```
lambda0 = 15
lambda1_p = 25
lambda1_m = 7
```

⁵ **TODO:** the cusum charts use z-transformed data. How do I get UCL and LCL from K^+ and h^+ and how are these actually determined?

⁶ **TODO:** I understand how K^+ is derived in this case. The values for K^- and h^- are obtained using a λ_1^- of 7. This should go into the exercise description. Please check if h^- (hm) is calculated correctly.

```

alpha = 0.001
tau = 30
kp = (lambdal_p - lambda0) / np.log(lambdal_p/lambda0)
hp = - np.log(alpha) / np.log(lambdal_p/lambda0)

km = (lambdal_m - lambda0) / np.log(lambdal_m/lambda0)
hm = - np.log(alpha) / np.log(lambdal_m/lambda0)

arl = mistat.cusumPfaCed(randFunc1=stats.poisson(mu=15),
                        randFunc2=stats.poisson(mu=25),
                        tau=tau,
                        kp=kp, km=km,
                        hp=hp, hm=hm,
                        N=4000, limit=1000, seed=1)
result = arl['statistic']

```

```

| PFA 0.01075 CED 2.1663 Std. Error 0.51072

```

Exercise 3.19 A CUSUM control scheme is based on sample means.

- (i) Determine the control parameters K^+ , h^+ , K^- , h^- , when $\mu_0 = 100$, $\mu_1 = 110$, $\sigma = 20$, $n = 5$, $\alpha = 0.001$.
- (ii) Estimate the PFA and CED, when the change point is at $\tau = 10, 20, 30$.
- (iii) How would the properties of the CUSUM change if each sample size is increased from 5 to 20.

Solution 3.19 ⁷ (i) Determine the control parameters:

```

mu0 = 100
mulp = 110
mulm = 90
sigma = 20
n = 5
alpha = 0.001

Kp = (mu0+mulp)/2
hp = -(sigma**2 / n) * np.log(alpha) / (mulp-mu0)
Km = (mu0+mulm)/2
hm = -(sigma**2 / n) * np.log(alpha) / (mulm-mu0)

pd.Series({'Kp': Kp, 'hp': hp, 'Km': Km, 'hm': hm})

```

```

| Kp      105.000000
| hp      55.262042
| Km      95.000000
| hm     -55.262042
| dtype: float64

```

(ii) Estimate PFA and CED for $\tau = 10, 20, 30$.

```

results = []
for tau in [10, 20, 30]:
    arl = mistat.cusumPfaCed(randFunc1=stats.norm(loc=mu0, scale=sigma/np.sqrt(5)),
                            randFunc2=stats.norm(loc=mulp, scale=sigma/np.sqrt(5)),
                            tau=tau, kp=Kp, km=Km, hp=hp, hm=hm,
                            N=300, limit=1000, seed=1, verbose=False)

```

⁷ **TODO:** I assume that mulm is symmetric here. This should go into the exercise description. Please check if h^- (hm) is calculated correctly.

```

        results.append({
            'tau': tau,
            **arl['statistic'], # copy all results from arl['statistic']
        })
    pd.DataFrame(results)

```

	tau	PFA	CED	Std. Error
0	10	0.000000	10.090000	1.050516
1	20	0.003333	10.076923	1.671984
2	30	0.006667	10.449664	2.289643

(iii) Increase n to 20 and estimate PFA and CED for $\tau = 10, 20, 30$.

```

n = 20
hp = -(sigma**2 / n) * np.log(alpha) / (mulp-mu0)
hm = -(sigma**2 / n) * np.log(alpha) / (mulm-mu0)

results = []
for tau in [10, 20, 30]:
    arl = mistat.cusumPfaCed(randFunc1=stats.norm(loc=mu0, scale=sigma/np.sqrt(5)),
                             randFunc2=stats.norm(loc=mulp, scale=sigma/np.sqrt(5)),
                             tau=tau, kp=Kp, km=Km, hp=hp, hm=hm,
                             N=300, limit=1000, seed=1, verbose=False)

    results.append({
        'tau': tau,
        **arl['statistic'], # copy all results from arl['statistic']
    })
pd.DataFrame(results)

```

	tau	PFA	CED	Std. Error
0	10	0.586667	1.709677	1.056345
1	20	0.806667	2.172414	2.910490
2	30	0.920000	2.250000	6.583663

The probability of a false alarm increases considerably.

Exercise 3.20 Show that the Shyryaev-Roberts statistic W_n , for detecting a shift in a Poisson distribution from a mean λ_0 to a mean $\lambda_1 = \lambda_0 + \delta$ is

$$W_m = (1 + W_{m-1})R_m$$

where $W_0 \equiv 0$, $R_m = \exp\{-\delta + x_m \log(\rho)\}$, and $\rho = \lambda_1/\lambda_0$.

Solution 3.20 According to Eq. (3.5.6), in the Poisson case,

$$\begin{aligned}
W_m &= \sum_{i=1}^{m-1} \prod_{j=i+1}^m R_j \\
&= \sum_{i=1}^{m-1} \exp \left\{ -(m-i)\delta + \sum_{j=i+1}^m X_j \log(\rho) \right\} \\
&= e^{-\delta + X_m \log \rho} + \sum_{i=1}^{m-2} \exp \left\{ -(m-i)\delta + \sum_{j=i+1}^m X_j \log(\rho) \right\} \\
&= e^{-\delta + X_m \log \rho} + e^{-\delta + X_m \log \rho} \sum_{i=1}^{m-2} \exp \left\{ -(m-1-i)\delta + \sum_{j=i+1}^{m-1} X_j \log(\rho) \right\} \\
&= (1 + W_{m-1})e^{-\delta + X_m \log \rho} \\
&= (1 + W_{m-1})R_m.
\end{aligned}$$

Exercise 3.21 Analyze the data in data **OELECT**, with an EWMA control chart with $\lambda = .2$.

Solution 3.21 The mean of the data is $\bar{X} = 219.25$, and its standard deviation is $S = 4.004$.

```

data = mistat.load_data('OELECT')
ewma = mistat.EWMA(data, center=data.mean(), std_dev=data.std(),
                   sizes=1, smooth=0.2, nsigmas=3)
ewma.plot()
plt.show()

```

Fig. 3.7 shows that there is no significant shift in the data.

{fig:exc-ewma-oelect}

Exercise 3.22 Analyze the variable diameters in the data set **ALMPIN** with an EWMA control chart with $\lambda = .2$. Explain how you would apply the automatic process control technique described at the end of Section 3.7.

{sec:automatic-proc-control}

Solution 3.22 In Fig. 3.8 we see the EWMA control chart for Diameter 1. The mean and standard deviation are $\bar{X} = 9.993$ and $S = 0.0164$. The target value for Diameter 1 is 10 mm. Notice in the chart that after a drop below 9.98 mm the machine corrects itself automatically, and there is a significant run upwards towards the target value. Using the EWMA chart an automatic control can be based on Eq. **9.8.11**.

{fig:exc-ewma-almphin}

```

data = mistat.load_data('ALMPIN')
data = data['diam1']
ewma = mistat.EWMA(data, center=data.mean(), std_dev=data.std(),
                   sizes=1, smooth=0.2, nsigmas=3)
ax = ewma.plot()
ax.set_ylim(9.96, 10.02)
plt.show()

```

Figure 9.6.2. EWMA chart for Diameter1

Exercise 3.23 Construct the Kalman Filter for the Dow-Jones daily index, which is given in the data set **DOW1941**.




Fig. 3.7 EWMA chart for the **OELECT** dataset

{fig:exc-ewma-oelect}




Fig. 3.8 EWMA chart for the **Diameter 1**

{fig:exc-ewma-almipin}

Solution 3.23 Using the first 50 values of **DOW1941** perform the regression method outlined in the text,

```
dow1941 = mistat.load_data('DOW1941')
# solve the regression equation
m = 50
sqrt_t = np.sqrt(range(1, m + 1))
df = pd.DataFrame({
    'Ut': dow1941[:m]/sqrt_t,
    'x1t': 1 / sqrt_t,
    'x2t': sqrt_t,
})
model = smf.ols(formula='Ut ~ x1t + x2t - 1', data=df).fit()
mu0, delta = model.params
var_eta = np.var(model.resid, ddof=2)
pd.Series({'mu0': mu0, 'delta': delta, 'Var(eta)': var_eta})
```

```
mu0      132.808555
delta    -0.255630
Var(eta)  0.297616
dtype: float64
```

The least squares estimates of the initial parameter values are $\hat{\mu}_0 = 132.809$, $\hat{\delta} = -0.2556$ and $\hat{\sigma}_\epsilon^2 + \hat{\sigma}_2^2 = 0.297616$.

For the Kalman filter, we choose $\hat{\sigma}_\epsilon^2 = 0.15$ and for w_0^2 the value 0.0015. The Python commands used to obtain the Kalman filter estimates of the **DOW1941** data are as follows:

```
# choose sig2e and w20
sig2e = 0.15
w20 = 0.0015
# apply the filter
results = []
mu_tm1 = mu0
w2_tm1 = w20
y_tm1 = mu0
for i in range(0, len(dow1941)):
    y_t = dow1941[i]
    B_t = sig2e / (var_eta + w2_tm1)
    mu_t = B_t * (mu_tm1 + delta) + (1 - B_t) * y_t # X
    results.append({
        't': i + 1,
        'y_t': y_t,
        'mu_t': mu_t,
        'B_t': B_t,
        'W2_t': w2_tm1,
    })
    w2_tm1 = B_t * (var_eta - sig2e + w2_tm1)
    mu_tm1 = mu_t
    y_tm1 = y_t
results = pd.DataFrame(results)
```


The Table 3.4 shows the first 25 values of the DOW 1941 data as well as the values of B_t , w_t^2 , and $\hat{\mu}_t$. Starting with $w_0^2 = 0.0015$ we see that w_t^2 converges fast to 0.0923. The values of $\hat{\mu}_t$ are very close to the data values. The result is shown graphically in Table 3.4 with actual values of Y_t in grey and estimated $\hat{\mu}_t$ in black.

(tab:exc-kalman-dow-jones)

(tab:exc-kalman-dow-jones)

Table 3.4 Results for Kalman filter applied to the Dow-Jones daily 1941 dataset*{tab:exc-kalman-dow-jones}*

	Dow 1941	$\hat{\mu}_t$	B_t	w_t^2
0	131.1300	131.8436	0.5015	0.0015
1	130.5700	130.9800	0.4028	0.0748
2	132.0100	131.5120	0.3874	0.0896
3	132.4000	131.9596	0.3851	0.0919
4	132.8300	132.3967	0.3848	0.0922
5	133.0200	132.6819	0.3847	0.0923
6	133.0200	132.7916	0.3847	0.0923
7	133.3900	133.0614	0.3847	0.0923
8	133.5900	133.2883	0.3847	0.0923
9	133.4900	133.3141	0.3847	0.0923
10	133.2500	133.1763	0.3847	0.0923
11	132.4400	132.6249	0.3847	0.0923
12	131.5100	131.8406	0.3847	0.0923
13	129.9300	130.5667	0.3847	0.0923
14	129.5400	129.8366	0.3847	0.0923
15	129.7500	129.6850	0.3847	0.0923
16	129.2400	129.3128	0.3847	0.0923
17	128.2000	128.5298	0.3847	0.0923
18	128.6500	128.5054	0.3847	0.0923
19	128.3400	128.3053	0.3847	0.0923
20	128.5200	128.3391	0.3847	0.0923
21	128.9600	128.6228	0.3847	0.0923
22	129.0300	128.7750	0.3847	0.0923
23	128.6000	128.5690	0.3847	0.0923
24	126.0000	126.8900	0.3847	0.0923



Chapter 4

Multivariate Statistical Process Control

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats
import matplotlib.pyplot as plt
import seaborn as sns

import mistat
```

Exercise 4.1 In data set **TSQ** we find 368 T^2 values corresponding to the vectors (x, y, θ) in the **PLACE** file. The first $n = 48$ vectors in **PLACE** file were used as a base sample, to compute the vector of means \mathbf{m} and the covariance matrix S . The T^2 values are for the other individual vectors ($m = 1$). Plot the T^2 values in the file **TSQ.csv**. Compute the UCL and describe from the plot what might have happened in the placement process generating the (x, y, θ) values.

Solution 4.1 The values of T^2 were computed, with \mathbf{m} being the mean of the first 48 vectors. Fig. ?? shows the T^2 values for all boards with the calculated control limit of UCL = 17.1953.

{fig:<exc-T2-control-limit}

```
# use eqn 4.1.6 to calculate the upper control limit for monitoring
n = 48; p = 3
UCL = (n-1)*(n+1)*p/(n*(n-p)) * stats.f(p, n-p).ppf(0.997)

tsq = mistat.load_data('TSQ')
ax = tsq.plot()
ax.axhline(UCL, color='black')
ax.axhline(0, color='black')
```

| <matplotlib.lines.Line2D at 0x7f8ab60386a0>

The side by side boxplots in Fig. 4.2 aggregate the information of the T^2 values by board number.

{fig:exc-T2-boxplots}

```
df = pd.DataFrame({
    'T2': tsq,
```




Fig. 4.1 Plot of T^2 values with control limit

(fig:exc-T2-control-limit)




Fig. 4.2 Plot of T^2 values with control limit

(fig:exc-T2-boxplots)

```
'board': mistat.load_data('PLACE')['cncBrd'][48:],
))
ax = df.groupby('board').boxplot(column='T2', subplots=False, rot=90, grid=False)
ax.axhline(UCL, color='grey')
plt.show()
```

We see that even on the first 9 cards there are a few outliers, that is, points whose T^2 is outside the control limits. All points from card 13 on and a majority of points from boards 10, 11 and 12 have T^2 values greater than UCL.

Exercise 4.2 Prove that if \mathbf{X} has a multivariate normal distribution, $(N_v(\boldsymbol{\mu}, \boldsymbol{\sigma}))$, then $(\mathbf{X} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{X} - \boldsymbol{\mu})$ has a χ^2 distribution with v degrees of freedom where $R = \chi^2_{1-p}[v]$ is the corresponding $(1 - p)$ quantile of the χ^2 distribution with v degrees of freedom.

Solution 4.2 As $(N_v(\boldsymbol{\mu}, \boldsymbol{\sigma}))$ is a multivariate normal distribution, $\boldsymbol{\Sigma}$ is positive definite. Therefore there exists a nonsingular matrix \mathbf{P} so that $\boldsymbol{\Sigma} = \mathbf{P}\mathbf{P}'$.

Using this we get,

$$\begin{aligned} (\mathbf{X} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{X} - \boldsymbol{\mu}) &= (\mathbf{X} - \boldsymbol{\mu})' (\mathbf{P}')^{-1} \mathbf{P}^{-1} (\mathbf{X} - \boldsymbol{\mu}) \\ &= (\mathbf{X} - \boldsymbol{\mu})' (\mathbf{P}^{-1})' \mathbf{P}^{-1} (\mathbf{X} - \boldsymbol{\mu}) \\ &= \left(\mathbf{P}^{-1} (\mathbf{X} - \boldsymbol{\mu}) \right)' \mathbf{P}^{-1} (\mathbf{X} - \boldsymbol{\mu}) \\ &= \mathbf{Y}' \mathbf{Y} \end{aligned}$$

with $\mathbf{Y} = \mathbf{P}^{-1} (\mathbf{X} - \boldsymbol{\mu})$.

¹

(exc:car-cylinder-T2)

Exercise 4.3 Sort the dataset **CAR** by variable **cyl**, indicating the number of cylinders in a car, and run a T^2 chart with internally derived targets for the variables **turn**, **hp**, **mpg**, with separate computations for cars with 4, 6 and 8 cylinders. How is the number of cylinders affecting the overall performance of the cars?

Solution 4.3 Create T^2 chart using separate computations by cylinders. The chart is shown in Fig. 4.3.

(fig:exc:car-cylinder-T2)

```
car = mistat.load_data('CAR')
car = car.sort_values('cyl')
columns = ['turn', 'hp', 'mpg']

fig, axes = plt.subplots(nrows=3, figsize=[8, 8])
for cyl, ax in zip([4, 6, 8], axes):
    base = car.loc[car['cyl'] == cyl, columns]
    newdata = car.loc[car['cyl'] != cyl, columns]
    mqcc = mistat.MultivariateQualityControlChart(base, qcc_type='T2single',
                                                  confidence_level=0.99, newdata=newdata)
    mqcc.plot(ax=ax, show_legend=False)
    ax.set_ylabel(f'{cyl} cylinders')
    plt.tight_layout()
plt.show()
```

The T^2 charts for the internally derived targets for 4 or 6 cylinders, show now relevant differences towards the remaining data. For 8 cylinders on the other hand, many of other cars show strong differences.

¹ **TODO:** proof required

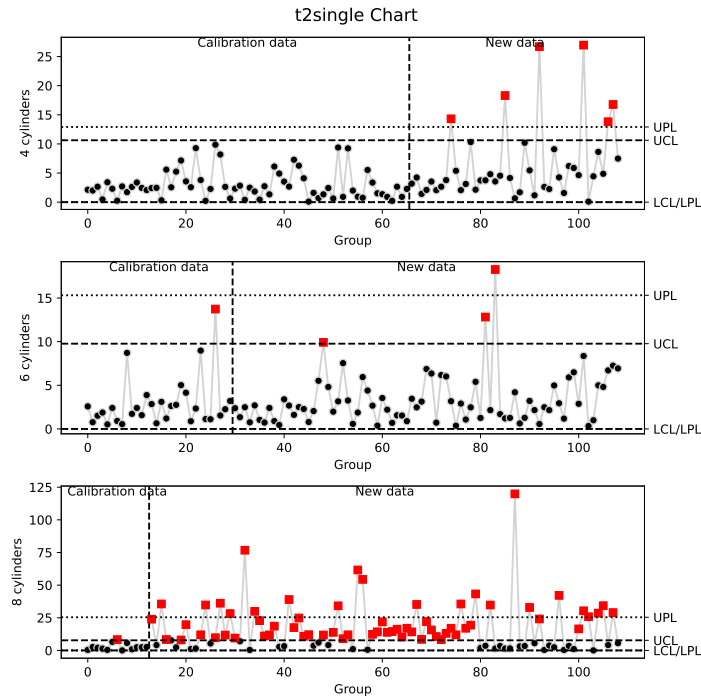


Fig. 4.3 T^2 charts for the CAR dataset using subsets based on number of cylinders to derive internal targets (Exercise 4.3)

(fig:exc:car-cylinder-T2)

Exercise 4.4 Sort the dataset CAR.csv by variable **origin**, indicating the country of origin, and run a T^2 chart with internally derived targets for the variables **turn**, **hp**, **mpg**, with separate computations for cars from 1 = US; 2 = Europe; 3 = Asia. How is the country of origin affecting the overall performance of the cars?

(exc:car-origin-T2)

Solution 4.4 Create T^2 chart using separate computations based on origin. The chart is shown in Fig. 4.4.

(fig:exc:car-origin-T2)

```
car = mistat.load_data('CAR')
car = car.sort_values('cyl')
columns = ['turn', 'hp', 'mpg']
origins = [None, 'US', 'Europe', 'Asia']

fig, axes = plt.subplots(nrows=3, figsize=[8, 8])
for origin, ax in zip([1, 2, 3], axes):
    base = car.loc[car['origin'] == origin, columns]
    newdata = car.loc[car['origin'] != origin, columns]
    mqcc = mistat.MultivariateQualityControlChart(base, qcc_type='T2single',
        confidence_level=0.99, newdata=newdata)
    mqcc.plot(ax=ax, show_legend=False)
    ax.set_ylabel(f'Origin {origins[origin]}')
plt.tight_layout()
plt.show()
```




Fig. 4.4 T^2 charts for the CAR dataset using subsets based on number of cylinders to derive internal targets (Exercise 4.4)

{fig:exc:car-origin-T2}

Using the European cars to derive internal targets, we can see that many of the other cars are above the UCL. This is less frequent in the other two plots.

The differences are also obvious when the data are visualized in a scatterplot matrix. See Fig. 4.5

{fig:exc:car-scatterplot-matrix}

```
sns.pairplot(car[*columns, 'cyl'], hue='cyl', height=1.5)
sns.pairplot(car[*columns, 'origin'], hue='origin', height=1.5)
plt.show()
```

Exercise 4.5 Load the dataset GASOL.csv and compute a T^2 chart for x_1 , x_2 , **astm**, **endPt**, **yield**. Design the chart with an external assigned target based on observations 12-24. Compare the charts. Explain the differences.

{exc:gasol-T2-1}

Solution 4.5 Create T^2 chart for full GASOL dataset.

```
gasol = mistat.load_data('GASOL')
mqcc = mistat.MultivariateQualityControlChart(gasol, qcc_type='T2single',
                                              confidence_level=0.99)
ax = mqcc.plot()
ax.plot((11, 23), (1, 1), color='red')
plt.show()
```

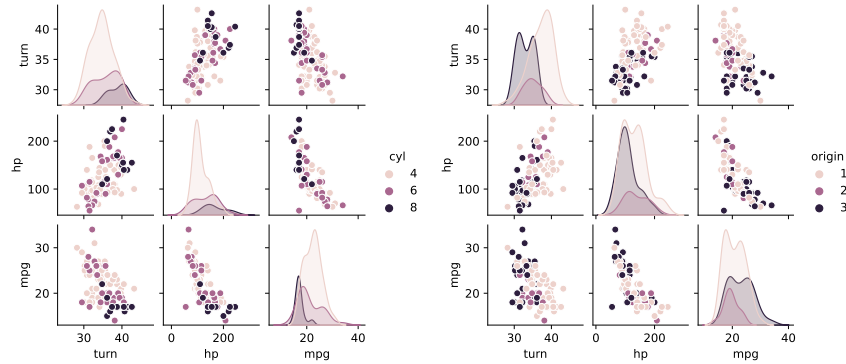


Fig. 4.5 Scatterplot matrix for the CAR dataset colored by cylinders (left) and origin (right).

{fig:exc:car-scatterplot-matrix}




Fig. 4.6 T^2 charts for the GASOL dataset. The horizontal line highlights observations 12 to 24 are highlighted

{fig:exc:gasol-T2}

Using observations 12 to 24 as an external assigned target, we get the T^2 chart in Fig. 4.7. All other data points have now T^2 value greater than UCL. The scatterplot matrix shows that the observations 12 to 24 have very distinct values especially for x_1 .

{fig:exc:gasol-T2-subset1}

```
gasol = mistat.load_data('GASOL')
base = gasol.iloc[11:24]
mqcc_base = mistat.MultivariateQualityControlChart(base, qcc_type='T2single',
                                                    confidence_level=0.99)
mqcc = mistat.MultivariateQualityControlChart(gasol, qcc_type='T2single',
```

```

        center=mqcc_base.stats.center, cov=mqcc_base.stats.cov,
        confidence_level=0.99)
ax = mqcc.plot()
plt.show()
gasol['color'] = ['red' if 11 <= i < 24 else 'black'
                  for i in range(len(gasol))]
sns.pairplot(gasol, hue='color', height=1.75)
plt.show()

```

Exercise 4.6 Repeat Exercise 4.5, but this time design the chart with an externally assigned target based on observations 25-32. Explain the computational difficulty.

{exc:gasol-T2-1}

Solution 4.6 The T^2 chart and a scatterplot matrix are shown in Fig. 4.8. Compared to Exercise 4.5, the effect is even stronger, as the variation of x_1 of the subset is even tighter and closer to 0. This results in very large values of T^2 .

{fig:exc:gasol-T2-subset2}

{exc:gasol-T2-1}

```

gasol = mistat.load_data('GASOL')
base = gasol.iloc[24:32]
mqcc_base = mistat.MultivariateQualityControlChart(base, qcc_type='T2single',
        confidence_level=0.99)
mqcc = mistat.MultivariateQualityControlChart(gasol, qcc_type='T2single',
        center=mqcc_base.stats.center, cov=mqcc_base.stats.cov,
        confidence_level=0.99)
ax = mqcc.plot()
plt.show()
gasol['color'] = ['red' if 24 <= i < 32 else 'black'
                  for i in range(len(gasol))]
sns.pairplot(gasol, hue='color', height=1.75)
plt.show()

```

Exercise 4.7 Calculate control limits for grouped data with 20 subgroups of size 5 and 6 dimensions, with internally derived targets (Eq. (4.4.2)). How will the control limits change if you start monitoring a process with similar data?

Solution 4.7 ²

Exercise 4.8 Let $X_1 = (x_{11}, x_{12}, \dots, x_{1p})$ and $X_2 = (x_{21}, x_{22}, \dots, x_{2p})$ represent the mean dissolution values of tablets at p time instances of a reference product and a batch under test, respectively. The Mahalanobis distance T^2 , between X_1 and X_2 , is defined here as $D_M = \sqrt{(X_2 - X_1)' S_{\text{pooled}}^{-1} (X_2 - X_1)}$, where $S_{\text{pooled}} = (S_{\text{reference}} + S_{\text{test}})/2$, is the pooled covariance matrix of the reference and test samples. The confidence region, CR, of the difference between batch and reference consists of all vectors Y satisfying: $K[(Y - (X_2 - X_1))' S_{\text{pooled}}^{-1} (Y - (X_2 - X_1))] \leq F_{0.90}[p, 2n - p - 1]$ where $F_{0.90}[p, 2n - p - 1]$ is the 90th quantile of the F -distribution with degrees of freedom p and $(2n - p - 1)$. Prove that for measurements conducted at one time instance ($p = 1$) these formulae correspond to the confidence intervals presented in Chapter 3 (Kenett et al., 2022).

Solution 4.8 ³

² TODO

³ TODO




Fig. 4.7 T^2 charts for the GASOL dataset using observations 12 to 24 (note the 0-indexing) as an external assigned target. The bottom scatterplot matrix highlights the subset values.

{fig:exc:gasol-T2-subset1}

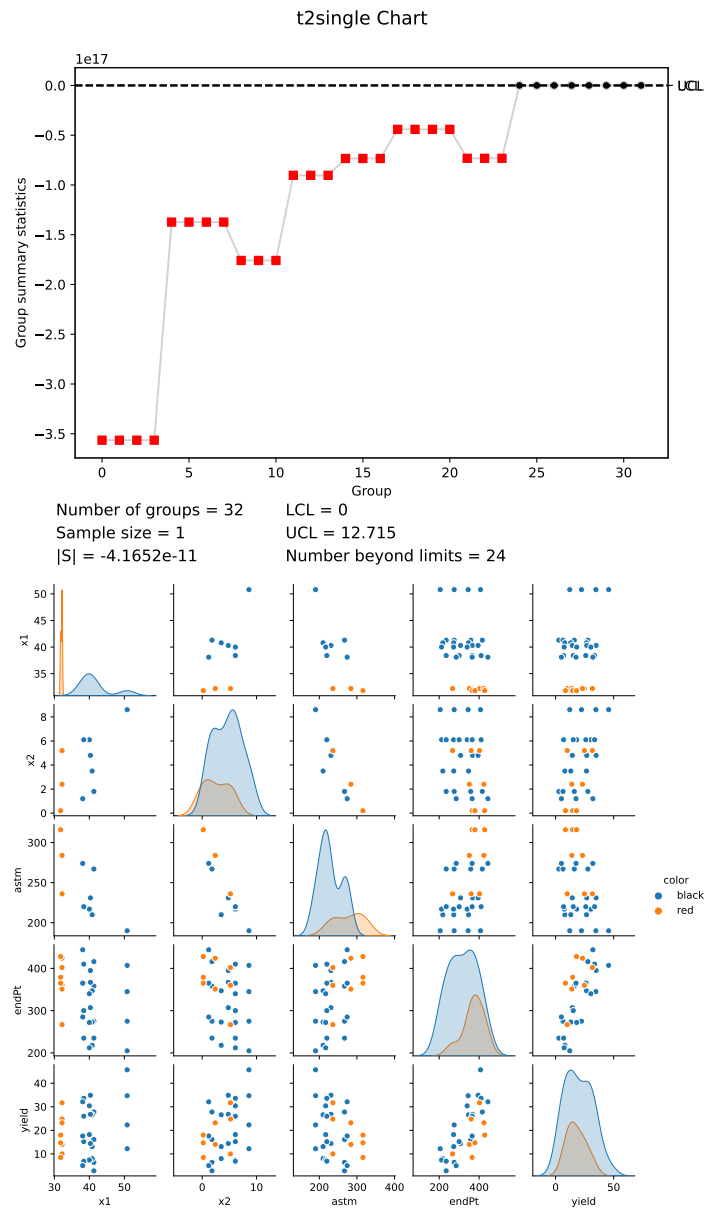


Fig. 4.8 T^2 charts for the GASOL dataset using observations 25 to 32 (note the 0-indexing) as an external assigned target. The bottom scatterplot matrix highlights the subset values.

(fig:exc:gasol-T2-subset2)

Chapter 5

Classical Design and Analysis of Experiments

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats
import statsmodels.formula.api as smf
from statsmodels.stats import anova
import matplotlib.pyplot as plt

import mistat
```

Exercise 5.1 Describe a production process familiar to you, like baking of cakes, or manufacturing concrete. List the pertinent variables. What is (are) the response variable(s)? Classify the variables which effect the response to noise variables and control variables. How many levels would you consider for each variable?

Solution 5.1 Preparing pancakes is a production process readily available for everyone to experiment with. The steps are as follows:

1. Open readymix box.
2. Measure prespecified amount of readymix into measuring cup.
3. Pour prespecified amount of readymix into container.
4. Measure prespecified amount of water in measuring cup.
5. Add prespecified amount of water to container.
6. Mix material in container with one or two eggs.
7. Warm cooking pan.
8. Add oil to pan in liquid, solid or spray form.
9. Pour material into cooking pan for one or more pancakes.
10. Wait for darkening signs on the pancakes rim.
11. Turn over pancakes.
12. Wait again.
13. Remove pancakes from cooking pan.

Examples of response variables include subjective taste testing by household members with classification on a 1 to 5 scale, drop test for measuring pancake body

composition (drop pancake from the table to the floor and count the number of pieces it decomposes into), quantitative tests performed in the laboratory to determine pancake chemical composition and texture.

Noise variables include environmental temperature and humidity, variability in amounts of readymix, water, oil and cooking time, stove heating capacity and differences in quality of raw materials.

Examples of control variables include:

- Amounts of readymix—less and more than present value
- Preheating time of cooking pan—less and more than present value
- Pouring speed of water into readymix container—slower and faster
- Waiting time before turn over—not just ready and overdone.

Exercise 5.2 Different types of adhesive are used in a lamination process, in manufacturing a computer card. The card is tested for bond strength. In addition to the type of adhesive, a factor which might influence the bond strength is the curing pressure (currently at 200 psi). Follow the basic steps of experimental design to set a possible experiment for testing the effects of adhesives and curing pressure on the bond strength.

Solution 5.2 Response variable: Bond strength test. Controllable factors: Adhesive type, Curing pressure. Factor levels: Adhesive type - A, B, C, D, Curing pressure - low, nominal, high. Experimental layout: 4x3 full factorial experiment with 4 replications.

Experimental Run	Adhesive type	Curing pressure	Replication			
1	A	low	1	2	3	4
2	A	nominal	1	2	3	4
3	A	high	1	2	3	4
4	B	low	1	2	3	4
5	B	nominal	1	2	3	4
...

Experiment protocol:

1. Prepare 48 (3x4x4) computer cards.
2. Randomly split computer cards into 4 groups, 12 cards per group.
3. Randomly assign adhesive type to each group, 1 type per group.
4. Randomly assign curing pressure level to group of 12 cards, 1 level for 4 cards.
5. Attach to each computer card a sticker indicating experimental run, replication number, adhesive type and curing pressure.
6. Randomize the order of the 48 cards.
7. Run experiment according to randomized order and factor levels indicated on sticker by keeping factors not participating in experiment fixed (e.g. amount of water).
8. Perform bond strength test and record data.
9. Analyze data with statistical model including main effects and interaction terms.

Exercise 5.3 Provide an example where blocking can reduce the variability of a product.

Solution 5.3 Blocking reduces the variability of the product relative to factors being studied. Examples include:

- Track and field athletes on a college team decided to investigate the effect of sleeping hours on athletic performance. The experiment consisted of sleeping a controlled amount of time prior to competitions. Individual athletes are natural experimental blocks.
- A natural extension of the example in section 5.2 on testing shoe sole materials is an experiment designed to test car tires. A natural block consists of the four wheels of a car, with an additional blocking variable determined by the position of the tires (front or rear).
- Experiments performed on plants are known to be sensitive to environmental conditions. Blocking variables in such experiments consist of neighboring plots of land where soil, humidity and temperature conditions are similar.

Exercise 5.4 Three factors A , B , C are tested in a given experiment, designed to assess their effects on the response variable. Each factor is tested at 3 levels. List all the main effects and interactions.

Solution 5.4 Main effects: τ_i^A , τ_j^B , τ_k^C each at 3 levels. First order interactions: τ_{ij}^{AB} , τ_{ik}^{AC} , τ_{jk}^{BC} each at 9 levels. Second order interaction: τ_{ijk}^{ABC} at 27 levels.

Exercise 5.5 Let x_1 , x_2 be two quantitative factors and Y a response variable. A regression model $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 x_2 + e$ is fitted to the data. Explain why β_{12} can be used as an interaction parameter.

Solution 5.5 $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 x_2 + e$. If $\beta_{12} = 0$ then the model is additive. The regression of Y on x_1 (x_2) are parallel for different values of x_2 (x_1). On the other hand, if $\beta_{12} \neq 0$ then the regression of Y on x_1 , for two values of x_2 , say $x_2^{(1)}$ and $x_2^{(2)}$, are $Y = \beta_0 + (\beta_1 + \beta_{12} x_2^{(1)}) x_1 + \beta_2 x_2^{(1)} + e$, and $Y = \beta_0 + (\beta_1 + \beta_{12} x_2^{(2)}) x_1 + \beta_2 x_2^{(2)} + e$. These are regression lines with different slopes. This means that the model is nonadditive and the extent of the interaction depends on β_{12} .

Exercise 5.6 Consider the ISC values for times t_1 , t_2 and t_3 in data file **SOCELL.csv**. Make a paired comparison for testing whether the mean ISC in time t_2 is different from that in time t_1 , by using a t -test.

Solution 5.6 ¹ Use the function `ttest_rel` from `scipy` to run a paired t -test.

```
ISC = mistat.load_data('SOCELL')
stats.ttest_rel(ISC['t2'], ISC['t1'])
```

```
| Ttest_relResult(statistic=11.302222594952937,
| pvalue=9.758913823226797e-09)
```

¹ **TODO:** is this using a paired t -test or a normal t -test?

*{tbl:exc-penicillin}***Table 5.1** Result of four treatments in pencillin manufacturing

blends	Treatments			
	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	89	88	97	94
2	84	77	92	79
3	81	87	87	85
4	87	92	89	84
5	79	81	80	88

The difference between the two means is significant

Exercise 5.7 Use `permutation_test` from `scipy` to perform a randomization test for the differences in the ISC values of the solar cells in times t_2 and t_3 (data file **SOCELL.csv**).

Solution 5.7 ² The permutation test requires the definition of the target statistic.

```
def statistic(x, y):
    return np.mean(x-y) # - np.mean(y)
```

The permutation test is then run as follows. The keyword argument `permutation_type='samples'` corresponds to the running a paired t-test.

```
def statistic(x, y):
    return np.mean(x) - np.mean(y)

res = stats.permutation_test((ISC['t2'], ISC['t3']), statistic,
                             permutation_type='samples', n_resamples=1000)
res.pvalue.round(5)
```

```
| 0.13786
```

A standard t-test gives;

```
stats.ttest_rel(ISC['t2'], ISC['t3']).pvalue.round(5)
```


```
| 0.11567
```

Both results are comparable. We can also visualize the distribution of the statistic; see Fig. 5.1.

{fig:ex-randomization-test}

```
fig, ax = plt.subplots()
ax.hist(res.null_distribution, bins=20, color='lightgrey')
ax.axvline(statistic(ISC['t2'], ISC['t3']), color='black')
plt.show()
```

² **TODO:** same question - paired t-test or normal t-test? Is the code for RPCOMP.MTB available



{fig:ex-randomization-test}

Fig. 5.1 Distribution resulting from the randomization paired comparison

Exercise 5.8 Box et al. (2005) give the results of four treatments *A*, *B*, *C*, *D* in penicillin manufacturing in five different blends (blocks) shown in Table 5.1.

{tbl:ex-penicillin}

Perform an ANOVA to test whether there are significant differences between the treatments or between the blends.

Solution 5.8 First create the dataset:

```
df = pd.DataFrame([
    ['B1', 'A', 89], ['B1', 'B', 88], ['B1', 'C', 97], ['B1', 'D', 94],
    ['B2', 'A', 84], ['B2', 'B', 77], ['B2', 'C', 92], ['B2', 'D', 79],
    ['B3', 'A', 81], ['B3', 'B', 87], ['B3', 'C', 87], ['B3', 'D', 85],
    ['B4', 'A', 87], ['B4', 'B', 92], ['B4', 'C', 89], ['B4', 'D', 84],
    ['B5', 'A', 79], ['B5', 'B', 81], ['B5', 'C', 80], ['B5', 'D', 88],
], columns=['blend', 'treatment', 'result'])
```

Then use `statsmodels` to perform an ANOVA.

```
model = smf.ols('result ~ C(blend) + C(treatment)', data=df).fit()
anova.anova_lm(model)
```

	df	sum_sq	mean_sq	F	PR(>F)
C(blend)	4.0	264.0	66.000000	3.504425	0.040746
C(treatment)	3.0	70.0	23.333333	1.238938	0.338658
Residual	12.0	226.0	18.833333	NaN	NaN

There are no significant differences between the treatments. There are significant differences between blends, the most extreme difference being between blend 1 and blend 5. See also Fig. 5.2.

{fig:ex-blends-treatments}

```
fig, axes = plt.subplots(ncols=2)
df.groupby('blend').boxplot(column='result', subplots=False,
                             rot=90, grid=False, ax=axes[0])
df.groupby('treatment').boxplot(column='result', subplots=False,
                                 rot=90, grid=False, ax=axes[1])
plt.tight_layout()
plt.show()
```




Fig. 5.2 Boxplot representation shows differences of blend and treatment effect.

{fig:ex-blends-treatments}

Table 5.2 Results of treatments A, B, C, \dots, H for different treatments

{tbl:exc-bibd-8treatment-2perblock}

block	Treatments				block	Treatments			
1	A	38	B	30	15	D	11	G	24
2	C	50	D	27	16	F	37	H	39
3	E	33	F	28	17	A	23	F	40
4	G	62	H	30	18	B	20	D	14
5	A	37	C	25	19	C	18	H	10
6	B	38	H	52	20	E	22	G	52
7	D	89	E	89	21	A	66	G	67
8	F	27	G	75	22	B	23	F	46
9	A	17	D	25	23	C	22	E	28
10	B	47	G	63	24	D	20	H	40
11	C	32	F	39	25	A	27	H	32
12	E	20	H	18	26	B	10	E	40
13	A	5	E	15	27	C	32	G	33
14	B	45	C	38	28	D	18	F	23

{exc:bid-8treatment-2perblock}

Exercise 5.9 Eight treatments A, B, C, \dots, H were tested in a BIBD of 28 blocks, $k = 2$ treatments per block, $r = 7$ and $\lambda = 1$. The results of the experiments are shown in Table 5.2

{tbl:exc-bibd-8treatment-2perblock}

Make an ANOVA to test the significance of the block effects, treatment effects. If the treatment effects are significant, make multiple comparisons of the treatments.

Solution 5.9 First create the dataset:

```
df = pd.DataFrame([
    [1, 'A', 38], [1, 'B', 30], [2, 'C', 50], [2, 'D', 27],
    [3, 'E', 33], [3, 'F', 28], [4, 'G', 62], [4, 'H', 30],
    [5, 'A', 37], [5, 'C', 25], [6, 'B', 38], [6, 'H', 52],
    [7, 'D', 89], [7, 'E', 89], [8, 'F', 27], [8, 'G', 75],
    [9, 'A', 17], [9, 'D', 25], [10, 'B', 47], [10, 'G', 63],
```

```
[11, 'C', 32], [11, 'F', 39], [12, 'E', 20], [12, 'H', 18],
[13, 'A', 5], [13, 'E', 15], [14, 'B', 45], [14, 'C', 38],
[15, 'D', 11], [15, 'G', 24], [16, 'F', 37], [16, 'H', 39],
[17, 'A', 23], [17, 'F', 40], [18, 'B', 20], [18, 'D', 14],
[19, 'C', 18], [19, 'H', 10], [20, 'E', 22], [20, 'G', 52],
[21, 'A', 66], [21, 'G', 67], [22, 'B', 23], [22, 'F', 46],
[23, 'C', 22], [23, 'E', 28], [24, 'D', 20], [24, 'H', 40],
[25, 'A', 27], [25, 'H', 32], [26, 'B', 10], [26, 'E', 40],
[27, 'C', 32], [27, 'G', 33], [28, 'D', 18], [28, 'F', 23],
], columns=['block', 'treatment', 'result'])
```

The ANOVA gives the following result:

```
model = smf.ols('result ~ C(block) + C(treatment)', data=df).fit()
anova.anova_lm(model)
```

	df	sum_sq	mean_sq	F	PR(>F)
C(block)	27.0	15030.482143	556.684524	5.334113	0.000109
C(treatment)	7.0	1901.875000	271.696429	2.603376	0.042207
Residual	21.0	2191.625000	104.363095	NaN	NaN

Both the effects of the treatments and the blocks are significant at the $\alpha = 0.05$ level.

The Scheffé coefficient for $\alpha = 0.05$ is $S_{.05} = (7 \times F_{.95}[7, 21])^{1/2} = 4.173$ and $\hat{c}_p = \sqrt{104.363} = 10.216$. Thus the treatments can be divided into two homogenous groups, $G_1 = \{A, B, C, D, E, F, H\}$ and $G_2 = \{G\}$.³ The 0.95 confidence interval for the difference of the group means is 20.413 ± 17.225 which shows that the group means are significantly different. See Fig. 5.3 for a visualization of the different groups.

(fig:tbl-d-treatment-2perblock)

```
df['group'] = ['G1' if t == 'G' else 'G2' for t in df['treatment']]
fig, axes = plt.subplots(ncols=3)
df.groupby('block').boxplot(column='result', subplots=False,
                             rot=90, grid=False, ax=axes[0])
df.groupby('treatment').boxplot(column='result', subplots=False,
                                 rot=90, grid=False, ax=axes[1])
df.groupby('group').boxplot(column='result', subplots=False,
                             rot=90, grid=False, ax=axes[2])
plt.tight_layout()
plt.show()
```

Exercise 5.10 Four different methods of preparing concrete mixtures A, B, C, D were tested, these methods consisted of two different mixture ratios of cement to water and two blending duration). The four methods (treatments) were blocks in four batches and four days, according to a Latin square design. The concrete was poured to cubes and tested for compressive strength [Kg/cm^2] after 7 days of storage in special rooms with 20°C temperature and 50% relative humidity. The results are in Table 5.3

(tbl:exc-concrete-mixtures)

Are the differences between the strength values of different treatments significant? [Perform the ANOVA.]

³ **TODO** how is Scheffé calculated? why 7,21 and not 7,27? Is the split into groups subjective or somehow derived?




Fig. 5.3 Boxplot representation for Exercise 5.9.

exc:ibid-8treatment-2perblock

{fig:ibid-8treatment-2perblock}

Table 5.3 Compressive strengths of cement for different treatments

{tbl:exc-concrete-mixtures}

Days	batches			
	1	2	3	4
1	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
	312	299	315	290
2	<i>C</i>	<i>A</i>	<i>D</i>	<i>B</i>
	295	317	313	300
3	<i>B</i>	<i>D</i>	<i>A</i>	<i>C</i>
	295	298	312	315
4	<i>D</i>	<i>C</i>	<i>B</i>	<i>A</i>
	313	314	299	300

Solution 5.10 Prepare the dataset.

```
df = pd.DataFrame([
    [1,1,'A',312], [1,2,'B',299], [1,3,'C',315], [1,4,'D',290],
    [2,1,'C',295], [2,2,'A',317], [2,3,'D',313], [2,4,'B',300],
    [3,1,'B',295], [3,2,'D',298], [3,3,'A',312], [3,4,'C',315],
    [4,1,'D',313], [4,2,'C',314], [4,3,'B',299], [4,4,'A',300],
], columns=['day', 'batch', 'mixture', 'result'])
```

Build a model and perform an ANOVA.

```
model = smf.ols('result ~ C(day) + C(batch) + C(mixture)', data=df).fit()
anova.anova_lm(model)
```

Table 5.4 Results of 3^2 factorial experiment with $n = 3$ observations*{tbl:32-factorial-replicates}*

	A_1	A_2	A_3
B_1	18.3	17.9	19.1
	17.9	17.6	19.0
	18.5	16.2	18.9
B_2	20.5	18.2	22.1
	21.1	19.5	23.5
	20.7	18.9	22.9
B_3	21.5	20.1	22.3
	21.7	19.5	23.5
	21.9	18.9	23.3

	df	sum_sq	mean_sq	F	PR(>F)
C(day)	3.0	16.1875	5.395833	0.048079	0.984704
C(batch)	3.0	165.6875	55.229167	0.492111	0.700652
C(mixture)	3.0	388.6875	129.562500	1.154446	0.401167
Residual	6.0	673.3750	112.229167	NaN	NaN

The differences between mixtures, batches or days are not significant.

Exercise 5.11 Repeat the experiments described in Example 5.7 at the low levels of factors B , C , E , F , and G . Perform the ANOVA for the main effects and interaction of spring coefficient and piston weight on the cycle time. Are your results different from those obtained in the example?

{ex:piston-two-factor}

Solution 5.11 ⁴

Exercise 5.12 For the data in Example 5.7 compute the least squares estimates of the main effects on the means and on the standard deviations.

{ex:piston-two-factor}

Solution 5.12 ⁵

Exercise 5.13 A 2^4 factorial experiment gave the following response values, arranged in standard order: 72, 60, 90, 80, 65, 60, 85, 80, 60, 50, 88, 82, 58, 50, 84, 75.

- Estimate all possible main effects.
- Estimate σ^2 under the assumption that all the interaction parameters are zero.
- Determine a confidence interval for σ^2 at level of confidence 0.99.

Solution 5.13 ⁶ (i) The LSE of the main effects are $\hat{A} = -4.0625$, $\hat{B} = 11.8125$, $\hat{C} = -1.5625$ and $\hat{D} = -2.8125$.

(ii) An estimate of σ^2 with 11 d.f. is $\hat{\sigma}^2 = 9.2898$.

(iii) A 0.99 level confidence interval for σ^2 is (3.819, 39.254).

⁴ TODO

⁵ TODO

⁶ **TODO:** in Python - standard order!

Exercise 5.14 A 3^2 factorial experiment, with $n = 3$ replications, gave the observations in Table 5.4.

(exc:32-factorial-replicates)

(tbl:32-factorial-replicates)

Perform an ANOVA to test the main effects and interactions. Break the between treatments sum of squares to one degree of freedom components. Use the Scheffé S_α coefficient to determine which effects are significant.

Solution 5.14 Prepare the dataset:

```
df = pd.DataFrame(
    [['A1', 'B1', v] for v in [18.3, 17.9, 18.5]] +
    [['A2', 'B1', v] for v in [17.9, 17.6, 16.2]] +
    [['A3', 'B1', v] for v in [19.1, 19.0, 18.9]] +
    [['A1', 'B2', v] for v in [20.5, 21.1, 20.7]] +
    [['A2', 'B2', v] for v in [18.2, 19.5, 18.9]] +
    [['A3', 'B2', v] for v in [22.1, 23.5, 22.9]] +
    [['A1', 'B3', v] for v in [21.5, 21.7, 21.9]] +
    [['A2', 'B3', v] for v in [20.1, 19.5, 18.9]] +
    [['A3', 'B3', v] for v in [22.3, 23.5, 23.3]],
    columns=['a', 'b', 'result'])
```

Build the model and perform an ANOVA.

```
model = smf.ols('result ~ C(a) + C(b) + C(a):C(b)', data=df).fit()
anova.anova_lm(model)
```

	df	sum_sq	mean_sq	F	PR(>F)
C(a)	2.0	43.080741	21.540370	70.495758	3.055533e-09
C(b)	2.0	54.169630	27.084815	88.641212	4.802685e-10
C(a):C(b)	4.0	4.345926	1.086481	3.555758	2.633500e-02
Residual	18.0	5.500000	0.305556	NaN	NaN

(fig:32-factorial-replicates)


Visualization of the factor and factor interactions are shown in Fig. 5.4

```
df['a:b'] = [f'{a}:{b}' for a, b in zip(df['a'], df['b'])]
fig, axes = plt.subplots(ncols=3)
df.groupby('a').boxplot(column='result', subplots=False,
                        rot=90, grid=False, ax=axes[0])
df.groupby('b').boxplot(column='result', subplots=False,
                        rot=90, grid=False, ax=axes[1])
df.groupby('a:b').boxplot(column='result', subplots=False,
                        rot=90, grid=False, ax=axes[2])
plt.tight_layout()
plt.show()
```

Exercise 5.15 Construct a 2^{8-2} fractional replication, using the generators $ABCDG$ and $ABEFH$. What is the resolution of this design? Write the aliases to the main effects, and to the first order interactions with the factor A .

Solution 5.15 ⁷

⁷ TODO



{fig:32-factorial-replicates} **Fig. 5.4** Boxplot representation for Exercise 5.14. ex:32-factorial-replicates

Table 5.5 Design matrix and response of 2^2 factorial design

{tbl:22-factorial-response}

X_1	X_2	Y
-1	-1	55.8
-1	-1	54.4
1	-1	60.3
1	-1	60.9
-1	1	63.9
-1	1	64.4
1	1	67.9
1	1	68.5
0	0	61.5
0	0	62.0
0	0	61.9
0	0	62.4

Exercise 5.16 Consider a full factorial experiment of $2^6 = 64$ runs. It is required to partition the runs to 8 blocks of 8. The parameters in the group of defining parameters are confounded with the effects of blocks, and are not estimable. Show which parameters are not estimable if the blocks are generated by ACE , $ABEF$ and $ABCD$.

Solution 5.16 ⁸

⁸ TODO

Exercise 5.17 A 2^2 factorial design is expanded by using 4 observations at 0. The design matrix and the response are in Table 5.5.

(exc:22-factorial-response)

{tbl:22-factorial-response}

- Fit a response function of the form: $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_{12} X_1 X_2 + e$, and plot its contour lines.
- Estimate the variance σ^2 and test the goodness of fit of this model.

Solution 5.17 (i) Prepare the dataset.

```
df = pd.DataFrame([
    [-1, -1, 55.8], [-1, -1, 54.4], [1, -1, 60.3], [1, -1, 60.9],
    [-1, 1, 63.9], [-1, 1, 64.4], [1, 1, 67.9], [1, 1, 68.5],
    [0, 0, 61.5], [0, 0, 62.0], [0, 0, 61.9], [0, 0, 62.4]
], columns=['X1', 'X2', 'Y'])
```

Build the model.

```
formula = ('Y ~ X1 + X2 + X1:X2')
model = smf.ols(formula, data=df).fit()
print(model.summary2())
```

```
Results: Ordinary least squares
=====
Model: OLS Adj. R-squared: 0.986
Dependent Variable: Y AIC: 19.8454
Date: 2022-09-10 20:00 BIC: 21.7851
No. Observations: 12 Log-Likelihood: -5.9227
Df Model: 3 F-statistic: 262.0
Df Residuals: 8 Prob (F-statistic): 2.52e-08
R-squared: 0.990 Scale: 0.23568
=====
              Coef.  Std.Err.  t  P>|t|  [0.025  0.975]
-----
Intercept    61.9917    0.1401  442.3492  0.0000    61.6685    62.3148
X1           2.3875    0.1716   13.9101  0.0000     1.9917     2.7833
X2           4.1625    0.1716   24.2516  0.0000     3.7667     4.5583
X1:X2        -0.3625    0.1716   -2.1120  0.0677    -0.7583     0.0333
=====
Omnibus: 0.321 Durbin-Watson: 2.580
Prob(Omnibus): 0.852 Jarque-Bera (JB): 0.447
Skew: 0.054 Prob(JB): 0.800
Kurtosis: 2.061 Condition No.: 1
=====
```

The response function is $\hat{Y} = 62.0 + 2.39X_1 + 4.16X_2 - 0.363X_1X_2$, with $R^2 = 0.99$. The contour plot is given in Figure 11.1.

```
def plotResponseSurface(model, ncontours=20):
    x1 = np.linspace(-1, 1)
    x2 = np.linspace(-1, 1)
    X1, X2 = np.meshgrid(x1, x2)
    exog = pd.DataFrame({'X1': X1.ravel(), 'X2': X2.ravel()})
    responses = model.predict(exog=exog)
    CS = plt.contour(x1, x2,
        responses.values.reshape(len(x2), len(x1)),
        ncontours, colors='gray')
    ax = plt.gca()
    ax.clabel(CS, inline=True, fontsize=10)
```




Fig. 5.5 Contour plot of equal responses (Exercise 5.17)

fig:22-factorial-response

{fig:22-factorial-response}

```
ax.set_xlabel('X1')
ax.set_ylabel('X2')
return ax

plotResponseSurface(model)
plt.show()
```

(ii)

```
# derive variance around the regression using an ANOVA (mean_sq of residuals)
res = anova.anova_lm(model)
res
```

	df	sum_sq	mean_sq	F	PR(>F)
X1	1.0	45.601250	45.601250	193.490387	6.905851e-07
X2	1.0	138.611250	138.611250	588.140552	8.915958e-09
X1:X2	1.0	1.051250	1.051250	4.460552	6.766219e-02
Residual	8.0	1.885417	0.235677	NaN	NaN

An estimate of the variance is $\hat{\sigma}^2 = 0.13667^9$, 3 d.f. The variance around the regression is $s^2_{y|(x)} = 0.2357$, 8 d.f. Use F-distribution to derive the p -value.

```
# Estimate of variance?
# ['Residual', 'mean_sq'] gives variance around regression
sigma2 = 0.13667
var_residuals = 0.235677
```

⁹ **TODO:** how is this calculated? eqn (5.7.41)? I get a different result with it; about double the value in the solution 0.26785

```
F = var_residuals / sigma2
p = 1 - stats.f(8, 3).cdf(F)
print(f'F-ratio: {F:.4f}; p-value: {p:.2f}')
```

```
| F-ratio: 1.7244; p-value: 0.36
```

The ratio $F = \frac{s_{y|x}^2}{\hat{\sigma}^2} = 1.7246$ gives a p -value of $P = 0.36$. This shows there is no significant differences between the variances.

Table 5.6 Design matrix and the response for a control composite design of Exercise 5.18

X_1	X_2	Y
1.0	0.000	95.6
0.5	0.866	77.9
-0.5	0.866	76.2
-1.0	0	54.5
-0.5	-0.866	63.9
0.5	-0.866	79.1
0	0	96.8
0	0	94.8
0	0	94.4

Exercise 5.18 Table 5.6 represents a design matrix and the response for a control composite design.

- Estimate the response function and its stationary point.
- Plot contours of equal response, in two dimensions.
- Conduct an ANOVA.

Solution 5.18 Prepare the dataset and build a regression model.

```
df = pd.DataFrame([
    [1, 0, 95.6], [0.5, 0.866, 77.9], [-0.5, 0.866, 76.2],
    [-1, 0, 54.5], [-0.5, -0.866, 63.9], [0.5, -0.866, 79.1],
    [0, 0, 96.8], [0, 0, 94.8], [0, 0, 94.4],
], columns=['X1', 'X2', 'Y'])

formula = ('Y ~ X1 + X2 + X1*X2 + I(X1**2) + I(X2**2)')
model = smf.ols(formula, data=df).fit()
print(model.summary2())
```

```
Results: Ordinary least squares
=====
Model: OLS Adj. R-squared: 0.855
Dependent Variable: Y AIC: 59.2942
Date: 2022-09-10 20:00 BIC: 60.4776
No. Observations: 9 Log-Likelihood: -23.647
Df Model: 5 F-statistic: 10.47
Df Residuals: 3 Prob (F-statistic): 0.0408
R-squared: 0.946 Scale: 33.638
=====
Coef. Std.Err. t P>|t| [0.025 0.975]
```

```

-----
Intercept      95.3333    3.3485  28.4703  0.0001   84.6768  105.9898
X1             16.5167    3.3485   4.9325  0.0160   5.8602   27.1732
X2              3.2044    3.3486   0.9569  0.4092  -7.4524   13.8612
X1:X2          -7.7945    6.6972  -1.1638  0.3286  -29.1081   13.5191
I (X1 ** 2)    -20.2833    5.2945  -3.8310  0.0313  -37.1327  -3.4339
I (X2 ** 2)    -21.3179    5.2948  -4.0262  0.0275  -38.1683  -4.4675
-----
Omnibus:                3.213          Durbin-Watson:          3.565
Prob(Omnibus):          0.201          Jarque-Bera (JB):          0.950
Skew:                   0.006          Prob(JB):                  0.622
Kurtosis:               1.408          Condition No.:            4
=====

```

The response function is $Y = 95.3 + 16.5X_1 + 3.20X_2 - 20.3X_1^2 - 21.3X_2^2 - 7.79X_1X_2$.

We can use `ResponseSurfaceMethod` from the `mistat` package to identify the stationary point.

```

rsm = mistat.ResponseSurfaceMethod(model, ['X1', 'X2'])
stationary = rsm.stationary_point()
stationary

```

```

X1    0.407004
X2    0.000751
dtype: float64

```

The stationary point is at (0.4, 0).

(ii) To plot the response surface, we reuse the function defined in Exercise 5.17; see Fig. 5.6.

[\[exc:22-factorial-response\]](#)

[\[fig:design-respose-control-composite\]](#)

```

ax = plotResponseSurface(model)
ax.scatter(*stationary, color='black')
plt.show()

```

(iii) Perform an ANOVA.

```
anova.anova_lm(model)
```

	df	sum_sq	mean_sq	F	PR(>F)
X1	1.0	818.400833	818.400833	24.329813	0.015976
X2	1.0	30.802500	30.802500	0.915712	0.409198
X1:X2	1.0	45.562500	45.562500	1.354504	0.328640
I (X1 ** 2)	1.0	320.800500	320.800500	9.536911	0.053790
I (X2 ** 2)	1.0	545.280333	545.280333	16.210355	0.027533
Residual	3.0	100.913333	33.637778	NaN	NaN

The factor X_1 and the squares X_1 , X_2 are significant. X_2 and the interaction $X_1 : X_2$ are not significant.




Fig. 5.6 Contour plot of equal responses; the stationary point is shown as a black dot (Exercise 5.18)

exc:design-respose-control-composite
(fig:design-respose-control-composite)

Chapter 6

Quality by Design

Import required modules and define required functions

```
import numpy as np
import pandas as pd
import statsmodels.formula.api as smf

import mistat
```

Exercise 6.1 The objective is to find the levels of the factors of the piston, which yield average cycle time of 0.45 [sec]. Execute a `PistonSimulation`, with sample size $n = 100$.

- (i) determine which treatment combination yields the smallest

$$MSE = (\bar{Y} - 0.45)^2 + S^2.$$

- (ii) determine which treatment combination yields the largest SN ratio,

$$\eta = 10 \log_{10} \left(\frac{Y^2}{S^2} - \frac{1}{100} \right).$$

What is the MSE at this treatment combination?

The five factors which are varied are: piston weight; piston surface area; initial gas volume; spring coefficient and ambient temperature. The factors atmospheric pressure and filling gas temperature are kept constant at the midrange level.

Solution 6.1 ¹

Exercise 6.2 Run a `PistonSimulation` with sample size of $n = 100$ and generate the sample means and standard deviation of the $2^7 = 128$ treatment combinations of a full factorial experiment, for the effects on the piston cycle time. Perform regression analysis to find which factors have significant effects on the signal to noise ratio $SN = \log((\bar{X}/S)^2)$.

¹ TODO

Solution 6.2²

Exercise 6.3 Let (X_1, X_2) have joint distribution with means (ξ_1, ξ_2) and covariance matrix

$$V = \begin{pmatrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{12} & \sigma_2^2 \end{pmatrix}.$$

Find approximations to the expected values and variances of:

- (i) $Y = X_1/X_2$;
- (ii) $Y = \log(X_1^2/X_2^2)$;
- (iii) $Y = (X_1^2 + X_2^2)^{1/2}$.

Solution 6.3 The approximations to the expected values and variances are as follows:

(i)

$$E \left\{ \frac{X_1}{X_2} \right\} \approx \frac{\xi_1}{\xi_2} - \sigma_{12} \frac{1}{\xi_2^2} + \sigma_2^2 \frac{\xi_1}{\xi_2^3}$$

$$V \left\{ \frac{X_1}{X_2} \right\} \approx \frac{\sigma_1^2}{\xi_2^2} + \sigma_2^2 \frac{\xi_1^2}{\xi_2^4} - 2\sigma_{12} \frac{\xi_1}{\xi_2^3}$$

(ii)

$$E \left\{ \log \frac{X_1^2}{X_2^2} \right\} \approx \log \left(\frac{\xi_1}{\xi_2} \right)^2 - \frac{\sigma_1^2}{\xi_1^2} + \frac{\sigma_2^2}{\xi_2^2}$$

$$V \left\{ \log \left(\frac{X_1}{X_2} \right)^2 \right\} \approx \frac{4\sigma_1^2}{\xi_1^2} + \frac{4\sigma_2^2}{\xi_2^2} - 8 \frac{\sigma_{12}}{\xi_1 \xi_2}.$$

(iii)

$$E \{ (X_1^2 + X_2^2)^{1/2} \} \approx (\xi_1^2 + \xi_2^2)^{1/2} + \frac{\sigma_1^2}{2(\xi_1^2 + \xi_2^2)^{1/2}} \left(1 - \frac{\xi_1^2}{\xi_1^2 + \xi_2^2} \right)$$

$$+ \frac{\sigma_2^2}{2(\xi_1^2 + \xi_2^2)^{1/2}} \left(1 - \frac{\xi_2^2}{\xi_1^2 + \xi_2^2} \right) - \frac{\sigma_{12} \xi_1 \xi_2}{(\xi_1^2 + \xi_2^2)^{3/2}};$$

$$V \{ (X_1^2 + X_2^2)^{1/2} \} \approx \frac{\sigma_1^2 \xi_1^2}{(\xi_1^2 + \xi_2^2)} + \frac{\sigma_2^2 \xi_2^2}{(\xi_1^2 + \xi_2^2)} + 2 \frac{\sigma_{12} \xi_1 \xi_2}{(\xi_1^2 + \xi_2^2)}.$$

Exercise 6.4 The relationship between the absorption ratio Y of a solid image in a copied paper and the light intensity X is given by the function

$$Y = 0.0782 + \frac{0.90258}{1 + 0.6969X^{-1.4258}}.$$

² TODO

Assuming that X has the gamma distribution $G(1, 1.5)$, approximate the expected value and variance of Y .

Solution 6.4 Approximation formulas yield: $E\{Y\} \approx 0.5159$ and $V\{Y\} \approx 0.06762$. Simulation with 5,000 runs yields the estimates $E\{Y\} \approx 0.6089$, $V\{Y\} \approx 0.05159$. The first approximation of $E\{Y\}$ is significantly lower than the simulation estimate.

Exercise 6.5 Let \bar{X}_n and S_n^2 be the mean and variance of a random sample of size n from a normal distribution $N(\mu, \sigma)$. We know that \bar{X}_n and S_n^2 are independent, $\bar{X}_n \sim N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$ and $S_n^2 \sim \frac{\sigma^2}{n-1} \chi^2[n-1]$. Find an approximation to the expected value and variance of $Y = \log\left(\frac{\bar{X}_n^2}{S_n^2}\right)$.

Solution 6.5 We have that $E\{\bar{X}_n\} = \mu$, $V\{\bar{X}_n\} = \frac{\sigma^2}{n}$, $E\{S_n^2\} = \sigma^2$ and $V\{S_n^2\} = \frac{2\sigma^4}{n-1}$. The first and second order partial derivatives of $f(\mu, \sigma^2) = 2 \log(\mu) - \log(\sigma^2)$ are

$$\begin{aligned} \frac{\partial}{\partial \mu} f &= \frac{2}{\mu}, & \frac{\partial}{\partial \sigma^2} f &= -\frac{1}{\sigma^2}, \\ \frac{\partial^2}{\partial \mu \partial \sigma^2} f &= 0, & \frac{\partial^2}{\partial \mu^2} f &= -\frac{2}{\mu^2} \quad \text{and} \quad \frac{\partial^2}{\partial (\sigma^2)^2} f = \frac{1}{\sigma^4}. \end{aligned}$$

Thus the approximations to the expected value and variance of $Y = \log\left(\frac{\bar{X}_n^2}{S_n^2}\right)$ are

$$\begin{aligned} E\left\{\log\left(\frac{\bar{X}_n^2}{S_n^2}\right)\right\} &\approx \log\left(\frac{\mu^2}{\sigma^2}\right) - \frac{\sigma^2}{n\mu^2} + \frac{1}{n-1} \quad \text{and} \\ V\left\{\log\left(\frac{\bar{X}_n^2}{S_n^2}\right)\right\} &\approx \frac{4\sigma^2}{n\mu^2} + \frac{2}{n-1}. \end{aligned}$$

Exercise 6.6 An experiment based on an L_{18} orthogonal array involving eight factors, gave the results listed in Table 6.1 (see Phadke et al., 1983). Each run had $n = 5$ replications. (tbl:phadke-data)

Analyze the effects of the factors of the SN ratio $\eta = \log(\bar{X}/S)$.

Solution 6.6 First, create the data frame.

```
df = pd.DataFrame([
    [1, 1, 1, 1, 1, 1, 1, 1, 2.5, 0.0827],
    [1, 1, 2, 2, 2, 2, 2, 2, 2.684, 0.1196],
    [1, 1, 3, 3, 3, 3, 3, 3, 2.66, 0.1722],
    [1, 2, 1, 1, 2, 2, 3, 3, 1.962, 0.1696],
    [1, 2, 2, 2, 3, 3, 1, 1, 1.87, 0.1168],
    [1, 2, 3, 3, 1, 1, 2, 2, 2.584, 0.1106],
    [1, 3, 1, 2, 1, 3, 2, 3, 2.032, 0.0718],
    [1, 3, 2, 3, 2, 1, 3, 1, 3.267, 0.2101],
    [1, 3, 3, 1, 3, 2, 1, 2, 2.829, 0.1516],
    [2, 1, 1, 3, 3, 2, 2, 1, 2.66, 0.1912],
    [2, 1, 2, 1, 1, 3, 3, 2, 3.166, 0.0674],
    [2, 1, 3, 2, 2, 1, 1, 3, 3.323, 0.1274],
```

Table 6.1 Results of experiment based on L_{18} orthogonal array

(tbl-phadke-data)

run	Factors								\bar{X}	S
	1	2	3	4	5	6	7	8		
1	1	1	1	1	1	1	1	1	2.500	0.0827
2	1	1	2	2	2	2	2	2	2.684	0.1196
3	1	1	3	3	3	3	3	3	2.660	0.1722
4	1	2	1	1	2	2	3	3	1.962	0.1696
5	1	2	2	2	3	3	1	1	1.870	0.1168
6	1	2	3	3	1	1	2	2	2.584	0.1106
7	1	3	1	2	1	3	2	3	2.032	0.0718
8	1	3	2	3	2	1	3	1	3.267	0.2101
9	1	3	3	1	3	2	1	2	2.829	0.1516
10	2	1	1	3	3	2	2	1	2.660	0.1912
11	2	1	2	1	1	3	3	2	3.166	0.0674
12	2	1	3	2	2	1	1	3	3.323	0.1274
13	2	2	1	2	3	1	3	2	2.576	0.0850
14	2	2	2	3	1	2	1	3	2.308	0.0964
15	2	2	3	1	2	3	2	1	2.464	0.0385
16	2	3	1	3	2	3	1	2	2.667	0.0706
17	2	3	2	1	3	1	2	3	3.156	0.1569
18	2	3	3	2	1	2	3	1	3.494	0.0473

```
[2, 2, 1, 2, 3, 1, 3, 2, 2.576, 0.085],
[2, 2, 2, 3, 1, 2, 1, 3, 2.308, 0.0964],
[2, 2, 3, 1, 2, 3, 2, 1, 2.464, 0.0385],
[2, 3, 1, 3, 2, 3, 1, 2, 2.667, 0.0706],
[2, 3, 2, 1, 3, 1, 2, 3, 3.156, 0.1569],
[2, 3, 3, 2, 1, 2, 3, 1, 3.494, 0.0473],
], columns=['F1', 'F2', 'F3', 'F4', 'F5',
            'F6', 'F7', 'F8', 'Xbar', 'S'])
```

Transform the factors to [-1, 1] for factor 1 and to [-1, 0, 1] for the remaining factors. Calculate the ratio.

```
df['F1'] = (df['F1']-1)*2-1
for column in ['F2', 'F3', 'F4', 'F5', 'F6', 'F7', 'F8']:
    df[column] = df[column] - 2
df['SNR'] = np.log(df['Xbar'] / df['S'])
```

Perform a regression analysis taking only the linear effects into account.

```
model = smf.ols('SNR ~ F1 + F2 + F3 + F4 + F5 + F6 + F7 + F8', data=df).fit()
```

Regression analysis yields the following:

```
print(model.summary2().tables[1].round(5))
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
Intercept	3.22601	0.07518	42.90997	0.00000	3.05594	3.39608
F1	0.26565	0.07518	3.53348	0.00638	0.09558	0.43572

F2	0.07900	0.09208	0.85797	0.41317	-0.12929	0.28729
F3	0.13848	0.09208	1.50392	0.16686	-0.06982	0.34677
F4	-0.14339	0.09208	-1.55733	0.15382	-0.35169	0.06490
F5	-0.31231	0.09208	-3.39185	0.00798	-0.52061	-0.10402
F6	0.12630	0.09208	1.37164	0.20340	-0.08200	0.33459
F7	0.02632	0.09208	0.28587	0.78145	-0.18197	0.23462
F8	-0.17109	0.09208	-1.85810	0.09610	-0.37938	0.03720

The linear effects of F1, F5, and F8 are significant. There might be significant interactions or quadratic effects.

Exercise 6.7 Using `PistonSimulation` perform a full factorial (2^7), a $1/8$ (2^{7-3}), $1/4$ (2^{7-2}) and $1/2$ (2^{7-1}) fractional replications of the cycle time experiment. Estimate the main-effects of the seven factors with respect to $SN = \log(\bar{X}/S)$ and compare the results obtained from these experiments.

{exc:piston-simulation-full-factorial}

Solution 6.7 ³

Exercise 6.8 To see the effect of the variances of the random variables on the expected response, in non-linear cases, execute `PistonSimulation`, with $n = 20$, and compare the output means to the values in Exercise 6.7.

{exc:piston-simulation-full-factorial}

Solution 6.8 ⁴

Exercise 6.9 Run `PowerCircuitSimulation` with 1% and 2% tolerances, and compare the results to those of Table 6.13.

{tbl:ex-perf-char-tol-des-exp}

Solution 6.9 ⁵

³ TODO

⁴ TODO

⁵ TODO

Chapter 7

Computer Experiments

Import required modules and define required functions

```
import random
import numpy as np
import pandas as pd
from scipy import stats
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
from pyKriging.krige import kriging
import mistat
```

Exercise 7.1 The birthday problem states that if there are more than 22 people at a birthday party, the probability that at least two people have the same birthday is greater than 0.5. Write a Python program to simulate this problem. Show that if there are more than 22 people in the party, the probability is greater than 1/2 that at least 2 will have birthdays on the same day.

Solution 7.1 In Python:

```
nrepeat = 10000
days = 365
for size in (22, 23): # range(1, 365+1):
    same_birthday = 0
    for _ in range(nrepeat):
        birthdays = stats.randint.rvs(1, 365+1, size=size)
        if len(birthdays) != len(set(birthdays)):
            same_birthday += 1
    print(f'size of party: {size}, ',
          f'p(same birthday) {same_birthday / nrepeat:.2f}')
```

```
| size of party: 22,  p(same birthday) 0.47
| size of party: 23,  p(same birthday) 0.51
```

Exercise 7.2 The Deming funnel experiment was designed to show that an inappropriate reaction to common cause variation will make matters worse. Common cause

and special causes affecting processes over time have been discussed in Chapters 2 to 4.

{ch:basic-proc-control}

{ch:multivar-proc-control}

In the actual demonstration, a funnel is placed above a circular target. The objective is to drop a marble through the funnel as close to the target as possible. A pen or pencil is used to mark the spot where the marble actually hits. Usually, 20 or more drops are performed in order to establish the pattern and extent of variation about the target. The funnel represents common causes affecting a system. Despite the operator's best efforts, the marble will not land exactly on the target each time. The operator can react to this variability in one of four ways:

1. do not move the funnel,
2. measure the distance the hit is from the target and move the funnel an equal distance, but in the opposite direction (error relative to the previous position),
3. measure the distance the hit is from the target and move the funnel this distance in the opposite direction, starting at the target (error relative to the target) and
4. move the funnel to be exactly over the location of the last hit.

Use Python to compare these four strategies using simulation data.

Solution 7.2 We first define a function that simulates the experiment. The marble will drop close to the drop position with a deviation that has a bivariate normal distribution. As the scale is not defined, we set the standard deviation to 1. The target is positioned at (0, 0).

```
def funnel_drop(position):
    ''' based on funnel position, returns marble drop position'''
    x_new = position[0] + stats.norm().rvs()
    y_new = position[1] + stats.norm().rvs()
    return np.array([x_new, y_new])

def result(strategy, dropped):
    ''' returns result information '''
    return {
        'strategy': strategy,
        'x': dropped[0],
        'y': dropped[1],
        'distance': np.sqrt(dropped[0]**2 + dropped[1]**2),
    }
```

{fig:funnel-drop-simulation}

Next we simulate the funnel drop experiment for different strategies. See Fig. 7.1 for a visualization of the simulation results.

```
np.random.seed(1)
results = []
nrepeat = 100

# strategy 1
# funnel fixed
position = (0, 0)
for _ in range(nrepeat):
    dropped = funnel_drop(position)
    results.append(result(1, dropped))

# strategy 2
# position funnel to compensate for error relative to funnel position
```

```

position = np.array([0, 0])
for _ in range(nrepeat):
    dropped = funnel_drop(position)
    results.append(result(2, dropped))
    position = position - dropped

# strategy 3
# position funnel to compensate for error relative to target
position = np.array([0, 0])
for _ in range(nrepeat):
    dropped = funnel_drop(position)
    results.append(result(3, dropped))
    position = - dropped

# strategy 4
# position funnel to compensate for error relative to target
position = np.array([0, 0])
for _ in range(nrepeat):
    dropped = funnel_drop(position)
    results.append(result(4, dropped))
    position = dropped

results = pd.DataFrame(results)
g = sns.FacetGrid(results, col='strategy', col_wrap=2)
g.map(sns.scatterplot, 'x', 'y')
plt.show()
sns.boxplot(x='distance', y='strategy', data=results, orient='h')
#g = sns.FacetGrid(results, col='strategy')
#g.map(sns.boxplot, 'distance', order='strategy', orient='v')
plt.show()

```

The result of the simulation shows that strategy 1, no interference, leads to results closest to the target. Strategy 2 still creates results close to the target, but with a larger deviation. Strategies 3 and 4 lead to results far away from the target.

Exercise 7.3 Design a 50 runs experimental array for running the piston simulator using different options available in the doepy package.

(exc:exp-array-options)

- Latin hypercube (simple): `doepy.build.lhs`
- Latin hypercube (space-filling): `doepy.build.space_filling_lhs`
- Random k-means cluster: `doepy.build.random_k_means`
- Maximin reconstruction: `doepy.build.maximin`
- Halton sequence based: `doepy.build.halton`
- Uniform random matrix: `doepy.build.uniform_random`

Compare the results.

Solution 7.3 We first create a series of design using the different methods

```

from doepy import build

np.random.seed(1) # set random seed for reproducibility
Factors = {
    'm': [30, 60],
    's': [0.005, 0.02],
    'v0': [0.002, 0.01],
    'k': [1_000, 5_000],
    'p0': [90_000, 110_000],
    't': [290, 296],
    't0': [340, 360],

```

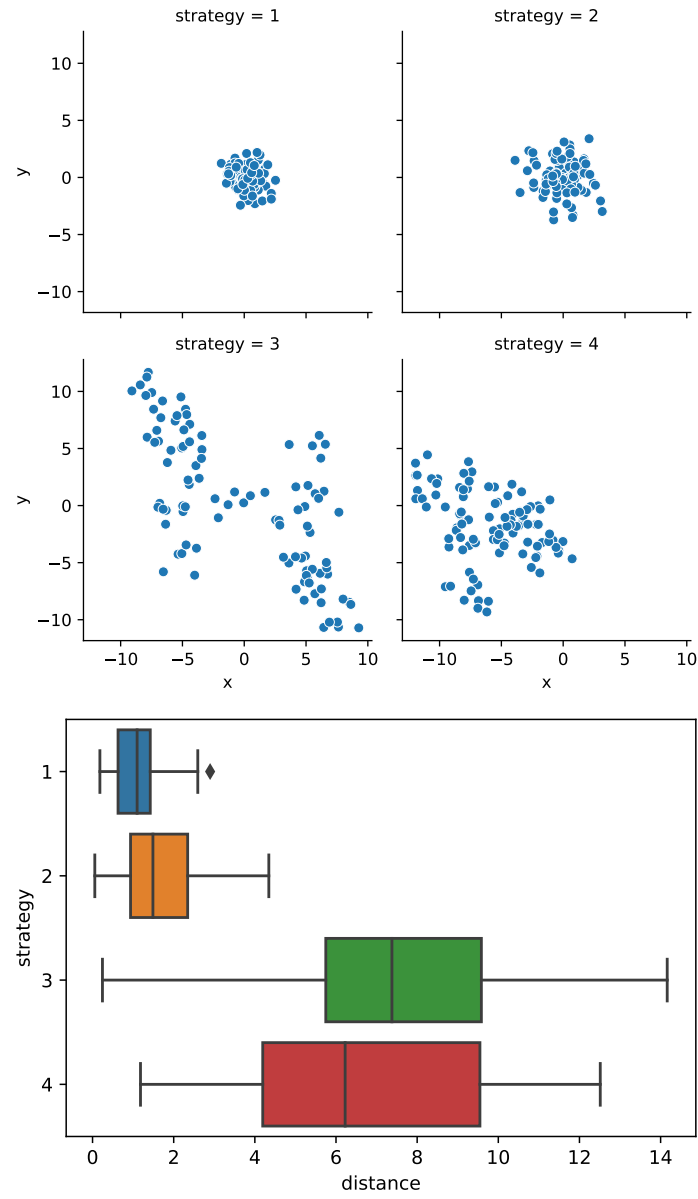


Fig. 7.1 Result of funnel drop experiment for the different strategies

(fig:funnel-drop-simulation)

```

}
Designs = {
    'Latin hypercube': build.lhs(Factors, num_samples=50),
    'Latin hypercube (space-filling)': build.space_filling_lhs(Factors, num_samples=50),
    'Random k-means cluster': build.random_k_means(Factors, num_samples=50),
    'Maximin reconstruction': build.maximin(Factors, num_samples=50),
    'Halton sequence based': build.halton(Factors, num_samples=50),
    'Uniform random matrix': build.uniform_random(Factors, num_samples=50),
}

```

Next we can visualize each design using the pandas scatterplot matrix method. We hide axis ticks for clarity.

```

for method, design in Designs.items():
    sm = pd.plotting.scatter_matrix(design, figsize=[4, 4])
    # hide all axis labels in visualization
    for subaxis in sm:
        for ax in subaxis:
            ax.xaxis.set_ticks([])
            ax.yaxis.set_ticks([])
            ax.set_ylabel("")
    plt.suptitle(method)
    plt.show()

```

The result is shown in Fig. 7.2. The methods lead to clearly distinct distributions of the design. The distributions for both latin hypercube designs, the maximin reconstruction design, and the uniform random matrix design lead to random coverage of the design space. Looking at the distribution for each factor, both latin hypercube designs show the most uniform distributions, while the maximin reconstruction and uniform random matrix methods lead to a design where factors are not uniformly explored.

{fig:design-visualization}

The random k -means cluster and the Halton sequence based designs are clearly distinct. The random k -means cluster method leads to a design that favors the extreme values for each factor. The Halton sequence based design shows patterns in the pair-wise scatter plots.

Exercise 7.4 Fit a Gaussian Process model to data generated by the six designs listed in Exercise 7.3 and compare the MSE of the model fits.

{exc:exp-array-options}

Solution 7.4 We run a 5-fold cross-validation using the kriging model from the pyKriging package. The model training will take some time.

```

random.seed(1)
np.random.seed(1)

outcome = 'seconds'
predictors = ['m', 's', 'v0', 'k', 'p0', 't', 't0']

performance = []
for method, Design in Designs.items():
    # randomize the design and run the piston simulator
    DesignRandomized = Design.sample(frac=1, random_state=1).reset_index(drop=True)
    simulator = mistat.PistonSimulator(parameter=DesignRandomized)
    result = simulator.simulate()

    # convert to X and y matrices to use with the

```

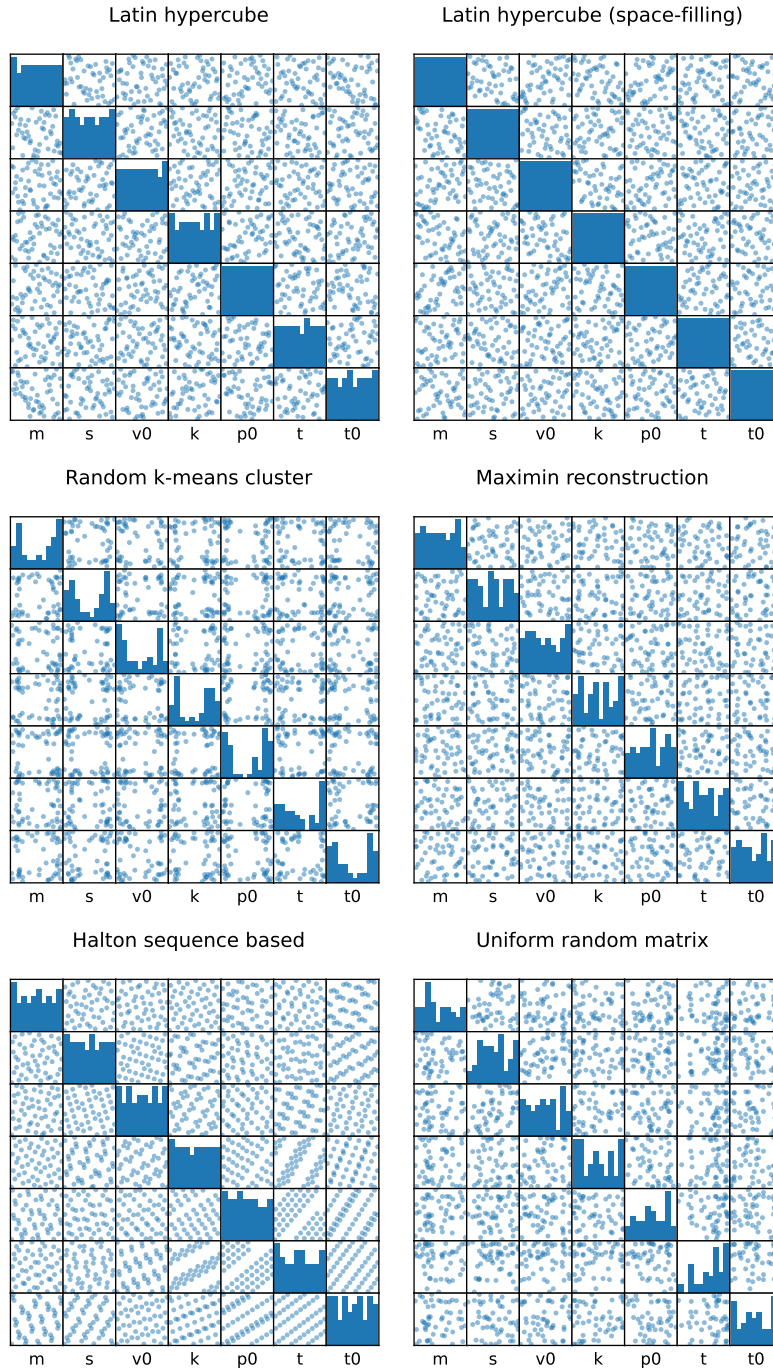


Fig. 7.2 Visualizations of different designs

{fig:design-visualization}

```

X = result[predictors].values
y = result[outcome].values

# run 5-fold cross-validation to determine MSE values
# we collect the out-of-fold predictions to determine
# the MSE
predicted = []
actual = []
for train_idx, test_idx in KFold(n_splits=5).split(result):
    model = kriging(X[train_idx], y[train_idx])
    model.train()

    # predict using the test set
    ypred = [model.predict(row) for row in X[test_idx]]
    # keep prediction results for performance metric calculation
    actual.extend(y[test_idx])
    predicted.extend(ypred)

mse = mean_squared_error(actual, predicted)
performance.append({
    'design': method,
    'mse': mse,
    'rmse': np.sqrt(mse),
})
print(performance)
# skip the remaining designs
break

```

The performance results are:

design	mse	rmse
Latin hypercube	0.00048	0.02183
Latin hypercube (space-filling)	0.00033	0.01820
Random k-means cluster	0.00015	0.01232
Maximin reconstruction	0.00029	0.01700
Halton sequence based	0.00022	0.01478
Uniform random matrix	0.00019	0.01360

In all cases, we observe good predictive performance. Due to randomness in the calculations, the relative order can be different between runs.

Exercise 7.5 Using a Uniform random Design, generate a Stochastic Emulator for the piston simulator in order to get 0.02 seconds cycle time with minimal variability.

{exc:uniform-piston}

Solution 7.5 We first generate the uniform random design

```

np.random.seed(1)

num_samples = 1000
Design = build.uniform_random(Factors, num_samples=num_samples)

```

Run the piston simulator using replicates of each parameter set. As the next exercise requires the same process, we wrap the required steps into a custom function

```

def evaluateDesign(Design, nrepeat=20):
    #Design = Design.loc[np.repeat(Design.index.values, nrepeat), :]
    #settings = {c: list(Design[c]) for c in Design.columns}
    #simulator = mistat.PistonSimulator(seed=1, **settings)

```

```

simulator = mistat.PistonSimulator(parameter=Design, seed=1, n_replicate=nrepeat)
result = simulator.simulate()
result = mistat.simulationGroup(result, nrepeat)

# next we aggregate the replicates and determine mean and
# standard deviation of each group.
options = {p: 'mean' for p in ['m', 's', 'v0', 'k', 'p0', 't', 't0']}
options['seconds'] = ['mean', 'std']
result = result.groupby('group').aggregate(options)
# convert multi-index to single index
result.columns = [' '.join(col) if col[0] == 'seconds' else col[0]
                  for col in result.columns]
return result

result = evaluateDesign(Design, nrepeat=30)

```

In order to find the parameter set that gives a cycle time around 0.02 with minimal variability, we filter the results by the mean values and sort by the standard deviation.

```

# determine rows with seconds around 0.02 and
# identify row with smallest standard deviation
rows = [0.0195 < v < 0.0205 for v in result['seconds mean'].values]
target = result.loc[rows, :]
target = target.sort_values('seconds std')
target.head(3)

```

	m	s	v0	k	p0
t \ group					
871	59.694778	0.018177	0.004087	1152.373331	109480.515301
293.342719					
331	46.278492	0.018261	0.004200	1432.269941	93825.441035
292.322277					
947	37.548930	0.012126	0.003148	1280.072660	105859.307614
291.076845					

	t0	seconds mean	seconds std
group			
871	350.794200	0.020356	0.002469
331	357.081196	0.020009	0.002751
947	342.524145	0.019582	0.003092


We can see that all three possible settings lead to a cycle time close to 0.02 seconds with low variability. The visualization of the results in Fig. 7.3 shows that the settings with a cycle time around 0.02 can be achieved with different values of the piston surface area s and the initial gas volume v_0 . The two variables show correlation. Only a small number of the design parameters match our criteria. It would be useful to repeat the simulation with a larger number of samples.

```

ax = target.plot.scatter(x='v0', y='s')
target.head(3).plot.scatter(x='v0', y='s', ax=ax, color='red')
ax = target.plot.scatter(x='seconds mean', y='seconds std')
target.head(3).plot.scatter(x='seconds mean', y='seconds std', ax=ax, color='red')
plt.show()

```

Exercise 7.6 Using a Latin Hypercube Design, generate a Stochastic Emulator for the piston simulator in order to achieve 0.02 seconds cycle time with minimal variability. Compare your results to what you got in Exercise 7.5



{fig:random-uniform-simulation}

Fig. 7.3 Visualization of results from stochastic emulator based on random uniform design.

Solution 7.6 We can make use of the function defined in the previous exercise to evaluate a latin hyper cube design

```
np.random.seed(1)

num_samples = 1000
Design = build.lhs(Factors, num_samples=num_samples)
result = evaluateDesign(Design, nrepeat=30)

# determine rows with seconds around 0.02 and
# identify row with smallest standard deviation
rows = [0.0195 < v < 0.0205 for v in result['seconds mean'].values]
target = result.loc[rows, :]
target = target.sort_values('seconds std')
target.head(3)
```

	m	s	v0	k	p0
t \ group					
368	33.166208	0.016760	0.003876	2998.808445	96800.530567
291.237264					
940	39.401808	0.019256	0.004388	3024.658005	97630.295509
295.393884					
150	33.434638	0.019884	0.004685	1401.296554	92707.973181
292.463943					

	t0	seconds mean	seconds std
group			
368	341.684282	0.020127	0.001956
940	351.093709	0.020001	0.002313
150	353.536655	0.020030	0.002467

The results derived using this design are similar to what was obtained in Exercise 7.5. We can see more points matching our criteria. This could be due to a better coverage of the design space. However, as already mentioned in Exercise 7.5, we could improve the analysis by increasing the number of design points.

{exc:uniform-piston}

{exc:uniform-piston}

```
ax = target.plot.scatter(x='v0', y='s')
target.head(3).plot.scatter(x='v0', y='s', ax=ax, color='red')
ax = target.plot.scatter(x='seconds mean', y='seconds std')
target.head(3).plot.scatter(x='seconds mean', y='seconds std', ax=ax, color='red')
plt.show()
```




Fig. 7.4 Visualization of results from stochastic emulator based on a latin hypersquare design.

(fig:lhs-simulation)

Chapter 8

Cybermanufacturing and digital twins

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats
import statsmodels.formula.api as smf
import lifelines
import pingouin as pg
import seaborn as sns
import matplotlib.pyplot as plt
import mistat
```

(exc:pensim-simulator)

Exercise 8.1 The PENSIM simulation software modelling penicillin production (Birolet al., 2002) is a fed-batch fermentor (<http://www.industrialpenicillinsimulation.com>). It simulates a fed-batch penicillin production process and includes variables such as pH, temperature, aeration rate, agitation power, feed flow rate of the substrate and a Raman probe.

The **PENSIM_100** dataset consists of 100 observations derived from the simulator. These are observational data collected under the same process set up. Variability in responses is induced by the varying process variables

A. Process set up:

1. **S0**: initial sugar concentration (15 g/L)
2. **X0**: initial biomass concentration (0.1 g/L)
3. **pH**: pH set point (5)
4. **T**: temperature set point (298 °K)
5. **air**: aeration (8.6 L/min)
6. **stirring**: agitation rate (29.9 W)
7. **time**: culture time (350 h)
8. **feed**: sugar feed rate (0.0426 L/h)

B. Process outputs:

1. **P**: Final penicillin concentration
2. **X**: Final biomass concentration

C. Process variables:

1. **Fg**: aeration rate
2. **RPM**: agitation rate
3. **Fs**: subst. feed
4. **Ts**: subst. temp.
5. **S**: substrate
6. **DO**: dissolved oxygen
7. **Uvis**: viscosity
8. **CO2**: off-gas CO2
9. **Hi**: heat inflow
10. **Ti**: temperature inflow
11. **Ho**: heat outflow
12. **Fw**: water for injection

Predict the process outputs P and X from the 12 process variables using two different models. Compare and contrast the models.

Some options are multivariate least square regressions (Chapter 4, Vol 1), regression trees, random forests and neural networks (Chapter 7, Vol 1), Bayesian networks (Chapter 8, Vol 1), response surfaces (Chapter 5, Vol 2) and Kriging (Gaussian) models (Chapter 7, Vol 2).

{ch:design-of-exp}
{ch:comp-exp}

Solution 8.1 ¹²

{exc:pensim-simulator}
{sec:expl-resp-surface}

Exercise 8.2 The PENSIM simulator introduced in Exercise 8.1 has been used to design and analyse a central composite design experiment (see Section 5.9). The file with the data is available as PENSIM_CCD.csv.

{sec:evalDesignExp}
{sec:expl-resp-surface}

1. Evaluate the experimental design set up (see Section 5.10)
2. Fit a second order response surface model to both X and P (see Section 5.9)
3. Compose a qualitative description of the models

Solution 8.2 ³⁴

{ex:railway-vehicle-vibrations}

Exercise 8.3 Example 4.5 provides an example of time series tracking engine vibrations of railway vehicle suspension systems. These suspensions can be affected by wheel flats with significant impact on system performance and safety. The file Order_PSD.csv includes three series where time is in units of revolution order. This angular resampling transformation is eliminating the variability in revolution time. The three time series correspond to vibrations in healthy suspensions and with wheel flats of 10 mm and 20 mm. In the analysis use the log transformed data.

1. Fit an ARIMA model to the healthy suspension vibration data (see Chapter 6, Vol 1)

¹ TODO

² TODO: add files to mistat

³ TODO

⁴ TODO: add files to mistat

2. Fit the same model to the 10 mm and 20 mm wheel flat
3. Compare the model parameters

Solution 8.3 ⁵⁶

Exercise 8.4 A company operates in 8 cities. The company product is temperature sensitive and management is interested in clustering the 8 locations by temperature characteristics. Monthly average daily minimum and maximum temperatures in these 8 cities, from 2000–2012, is available from https://www.stat.auckland.ac.nz/~wild/data/data_from_inZight/TimeSeriesDatasets_130207/TempWorld1.csv.

1. Use a smoother to compare the maximum and minimum monthly average temperatures in the 8 cities.
2. Group the cities using maximum and minimum temperature patterns using hierarchical clusters.
3. Group the cities using maximum and minimum temperature patterns using *K*-means clusters.

Solution 8.4 ⁷⁸

⁵ TODO

⁶ TODO: add files to mistat

⁷ TODO

⁸ TODO: add files to mistat

Chapter 9

Reliability Analysis

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats
import statsmodels.formula.api as smf
import lifelines
import pingouin as pg
import seaborn as sns
import matplotlib.pyplot as plt
import mistat
```

Exercise 9.1 During 600 hours of manufacturing time a machine was up 510 hours. It had 100 failures which required a total of 11 hours of repair time. What is the MTTF of this machine? What is its mean time till repair, MTTR? What is the intrinsic availability?

Solution 9.1 $MTTF = 5.1$ [hr]; $MTTR = 6.6$ [min]; Intrinsic Availability = 0.979.

Exercise 9.2 The frequency distribution of the lifetime in a random sample of $n = 2,000$ solar cells, under accelerated life testing is the following:

$t[10^3 \text{ [hr]}]$	0-1	1-2	2-3	3-4	4-5	5-
prof. freq.	0.15	0.25	0.25	0.10	0.10	0.15

The relationship of the scale parameters of the life distributions, between normal and accelerated conditions is 10:1.

- (i) Estimate the reliability of the solar cells at age $t = 4.0$ [yr].
- (ii) What proportion of solar cells are expected to survive 40,000 [hr] among those which survived 20,000 [hr]?

Solution 9.2 ¹

¹ TODO

- (i) Since 4 yrs = 35,040 hr, the reliability estimate is $R(4) = 0.30$.
(ii) Of the solar cells that survive 20,000 hours, the proportion expected to survive 40,000 hours is $\frac{0.25}{0.60} = 0.42$.

Exercise 9.3 The CDF of the lifetime [months] of an equipment is

$$F(t) = \begin{cases} t^4/20736, & 0 \leq t < 12 \\ 1, & 12 \leq t \end{cases}$$

- (i) What is the failure rate function of this equipment?
(ii) What is the MTTF?
(iii) What is the reliability of the equipment at 4 months?

Solution 9.3 ²

- (i) The hazard function is

$$h(t) = \begin{cases} \frac{4t^3}{20736 \cdot \left(1 - \frac{t^4}{20736}\right)}, & 0 \leq t < 12 \\ \infty, & 12 \leq t. \end{cases}$$

- (ii) $MTTF = \int_0^{12} \left(1 - \frac{t^4}{20736}\right) dt = 9.6$ [Months].
(iii) $R(4) = 1 - F(4) = 0.9877$.

Exercise 9.4 The reliability of a system is

$$R(t) = \exp\{-2t - 3t^2\}, \quad 0 \leq t < \infty.$$

- (i) What is the failure rate of this system at age $t = 3$?
(ii) Given that the system reached the age of $t = 3$, what is its reliability for two additional time units?

Solution 9.4 ³ (i) $h(t) = 2 + 6t$, $h(3) = 20$. (ii) $R(5)/R(3) = 0$.

Exercise 9.5 An aircraft has four engines but can land using only two engines.

- (i) Assuming that the reliability of each engine, for the duration of a mission, is $R = .95$, and that engine failures are independent, compute the mission reliability of the aircraft.
(ii) What is the mission reliability of the aircraft if at least one functioning engine must be on each wing?

Solution 9.5 ⁴ (i) $1 - B(1; 4, 0.95) = 0.9995$ (ii) $(1 - B(0; 2, 0.95))^2 = (1 - 0.05^2)^2 = 0.9950$.

² TODO

³ TODO

⁴ TODO




Fig. 9.1 System schema

{fig:ex-system-schema}

Exercise 9.6 (i) Draw a block diagram of a system having the structure function

$$R_{\text{sys}} = \psi_s(\psi_p(\psi_{M_1}, \psi_{M_2}), R_6), \quad \psi_{M_1} = \psi_p(R_1, R_2 R_3), \quad \psi_{M_2} = \psi_2(R_4, R_5)$$

(ii) Determine R_{sys} if all the components act independently, and have the same reliability $R = .8$.

Solution 9.6 (a) The block diagram of the system is shown in Fig. 9.1.

{fig:ex-system-schema}

(b) $R_{M_1} = 1 - (1 - R_1)(1 - R_2)(1 - R_3)$ and $R_{M_2} = R_4 R_5$. Thus,

$$\begin{aligned}
 R_{\text{sys}} &= [1 - (1 - R_1)(1 - R_2)(1 - R_3)(1 - R_4 R_5)] R_6 \\
 &= R_1 R_6 + R_2 R_6 + R_3 R_6 - R_1 R_2 R_6 + R_4 R_5 R_6 \\
 &\quad - R_1 R_3 R_6 - R_2 R_3 R_6 + R_1 R_2 R_3 R_6 \\
 &\quad - R_1 R_4 R_5 R_6 - R_2 R_4 R_5 R_6 - R_3 R_4 R_5 R_6 \\
 &\quad + R_2 R_3 R_4 R_5 R_6 + R_1 R_3 R_4 R_5 R_6 + R_1 R_2 R_4 R_5 R_6 \\
 &\quad - R_1 R_2 R_3 R_4 R_5 R_6.
 \end{aligned}$$

If all values of $R = 0.8$ then $R_{\text{sys}} = 0.7977$.

Exercise 9.7 Consider a system of n components in a series structure. Let R_1, \dots, R_n be the reliabilities of the components. Show that

$$R_{\text{sys}} \geq 1 - \sum_{i=1}^n (1 - R_i).$$

Solution 9.7 Extending Bonferroni's inequality,

$$\begin{aligned}
\Pr\{C_1 = 1, \dots, C_n = 1\} &= 1 - \Pr\left\{\bigcup_{i=1}^n \{C_i = 0\}\right\} \\
&\geq 1 - \sum_{i=1}^n \Pr\{C_i = 0\} \\
&= 1 - \sum_{i=1}^n (1 - R_i);
\end{aligned}$$

where $\{C_i = 1\}$ is the event that the i -th component functions and $\{C_i = 0\}$ is the event that it fails. Since $R_{sys} = \Pr\{C_1 = 1, \dots, C_n = 1\}$ for a series structure, the inequality follows.

Exercise 9.8 A 4 out of 8 system has identical components whose life lengths T [weeks] are independent and identically distributed like a Weibull $W\left(\frac{1}{2}, 100\right)$. What is the reliability of the system at $t_0 = 5$ weeks?

Solution 9.8 The reliability of a component is $\theta(5) = \Pr\left\{W\left(\frac{1}{2}, 100\right) > 5\right\} = 0.79963$.

Therefore, $R_{sys} = 1 - B(3; 8, 0.7996) = 0.9895$.

Exercise 9.9 A system consists of a main unit and two standby units. The lifetimes of these units are exponential with mean $\beta = 100$ [hr]. The standby units undergo no failure while idle. Switching will take place when required. What is the MTTF of the system? What is the reliability function of this system?

Solution 9.9 The life length of the system is the sum of 3 independent $E(100)$, that is $G(3, 100)$ (gamma distribution). Hence the $MTTF = 300$ hr. The reliability function of the system is

$$R_{sys}(t) = 1 - G(t; 3, 100).$$

Exercise 9.10 Suppose that the TTF in a renewal cycle has a $W(\alpha, \beta)$ distribution and that the TTR has a lognormal distribution $LN(\mu, \sigma)$. Assume further that TTF and TTR are independent. What are the mean and standard deviation of a renewal cycle.

Solution 9.10 Let C be the length of the renewal cycle.

$$\begin{aligned}
E\{C\} &= \beta \Gamma\left(1 + \frac{1}{\alpha}\right) + e^{\mu + \sigma^2/2}. \\
\sigma(C) &= \left[\beta^2 \left(\Gamma\left(1 + \frac{2}{\alpha}\right) - \Gamma^2\left(1 + \frac{1}{\alpha}\right) \right) + e^{2\mu + \sigma^2} (e^{\sigma^2} - 1) \right]^{1/2}.
\end{aligned}$$

Exercise 9.11 Suppose that a renewal cycle has the normal distribution $N(100, 10)$. Determine the p.d.f of $N_R(200)$.

Solution 9.11 Let C_1, C_2, \dots denote the renewal cycle times, $C_i \sim N(100, 10)$. $N_R(200)$ is a random variable representing the number of repairs that occur during the time interval $(0, 200]$.

$$\Pr\{N_R(200) = 0\} = \Pr\{C_1 > 200\} = 1 - \Phi\left(\frac{200 - 100}{10}\right) = 1 - \Phi(10) = 0.$$

For $k = 1, 2, 3, \dots$, $C_1 + \dots + C_k \sim N(100k, 10\sqrt{k})$ and so

$$\begin{aligned} \Pr\{N_R(200) = k\} &= \Pr\{N_R(200) \geq k\} - \Pr\{N_R(200) \geq k+1\} \\ &= \Pr\{C_1 + \dots + C_k \leq 200\} - \Pr\{C_1 + \dots + C_{k+1} \leq 200\} \\ &= \Phi\left(\frac{200 - 100k}{10\sqrt{k}}\right) - \Phi\left(\frac{200 - 100(k+1)}{10\sqrt{k+1}}\right) \\ &= \Phi\left(\frac{20}{\sqrt{k}} - 10\sqrt{k}\right) - \Phi\left(\frac{20}{\sqrt{k+1}} - 10\sqrt{k+1}\right), \end{aligned}$$

where $\Phi(\cdot)$ denotes the c.d.f. of a $N(0, 1)$ distribution. Thus the p.d.f. of $N_R(200)$ is given by the following table

k	$\Pr\{N_R(200) = k\}$
0	0
1	0.5
2	0.5
3	0
4	0
\vdots	\vdots

Exercise 9.12 Let the renewal cycle C be distributed like $N(100, 10)$. Approximate $V(1000)$.

Solution 9.12 $V(1000) = \sum_{n=1}^{\infty} \Phi\left(\frac{100-10n}{\sqrt{n}}\right) \approx 9.501$.

Exercise 9.13 Derive the renewal density $v(t)$ for a renewal process with $C \sim N(100, 10)$.

Solution 9.13 $v(t) = \frac{1}{10\sqrt{n}} \sum_{n=1}^{\infty} \phi\left(\frac{t-100n}{10\sqrt{n}}\right)$, where $\phi(Z)$ is the p.d.f. of $N(0, 1)$.

Exercise 9.14 Two identical components are connected in parallel. The system is not repaired until both components fail. Assuming that the TTF of each component is exponentially distributed, $E(\beta)$, and the total repair time is $G(2, \gamma)$, derive the Laplace transform of the availability function $A(t)$ of the system.

Solution 9.14 Since $TTF \sim \max(E_1(\beta), E_2(\beta))$, $F(t) = (1 - e^{-t/\beta})^2$ and $R(t) = 1 - (1 - e^{-t/\beta})^2$.

Using Laplace transforms we have

$$R^*(s) = \frac{\beta(3 + s\beta)}{(1 + s\beta)(2 + s\beta)}, \quad f^*(s) = \frac{2}{(1 + \beta s)(2 + \beta s)}, \quad g^*(s) = \frac{1}{(1 + s\gamma)^2} \quad \text{and}$$

$$A^*(s) = \frac{\beta(3 + s\beta)(1 + \gamma s)^2}{s[3\beta + 4\gamma + (\beta^2 + 6\beta\gamma + 2\gamma^2)s + (2\beta^2\gamma + 3\beta\gamma^2)s^2 + \beta^2\gamma^2s^3]}.$$

Exercise 9.15 Simulate a sample of 100 TTF of a system comprised of two components connected in parallel, where the life distribution of each component (in hours) is $E(100)$. Similarly, simulate a sample of 100 repair times (in hours), having a $G(2, 1)$ distribution. Estimate the expected value and variance of the number of renewals in 2,000 [hr].

Solution 9.15 For our samples, with $M = 500$ runs, we get the following estimates:

$$E\{N_R(2000)\} = V(2000) = 16.574 \quad \text{and} \quad \text{Var}\{N_R(2000)\} = 7.59.$$

{exc:independent-units}

Exercise 9.16 In a given life test, $n = 15$ units are placed to operate independently. The time till failure of each unit has an exponential distribution with mean 2,000 [hr]. The life test terminates immediately after the 10-th failure. How long is the test expected to last?

Solution 9.16 The expected length of the test is 2069.8 [hr].

Exercise 9.17 If n units are put on test and their TTF are exponentially distributed with mean β , the time elapsed between the r -th and $(r + 1)$ -st failure, i.e., $\Delta_{n,r} = T_{n,r+1} - T_{n,r}$, is exponentially distributed with mean $\beta/(n - r)$, $r = 0, 1, \dots, n - 1$. Also, $\Delta_{n,0}, \Delta_{n,2}, \dots, \Delta_{n,n-1}$ are independent. What is the variance of $T_{n,r}$? Use this result to compute the variance of the test length in the previous exercise.

Solution 9.17 Let $T_{n,0} \equiv 0$. Then $T_{n,r} = \sum_{j=1}^r (T_{n,j} - T_{n,j-1}) = \sum_{j=1}^r \Delta_{n,j-1}$. Thus,

$$V\{T_{n,r}\} = \beta^2 \sum_{j=1}^r \left(\frac{1}{n - j + 1} \right)^2.$$

{exc:independent-units}

In Exercise 9.16, $\beta = 2,000$ hr, $n = 15$ and $r = 10$. This yields

$$V\{T_{15,10}\} = 4 \times 10^6 \times 0.116829 = 467,316 \text{ (hr)}^2.$$

Exercise 9.18 Consider again the previous exercise. How would you estimate unbiasedly the scale parameter β , if the r failure times, $T_{n,1}, T_{n,2}, \dots, T_{n,r}$ are given? What is the variance of this unbiased estimator?

Solution 9.18 Let $S_{n,r} = \sum_{i=1}^r T_{n,i} + (n - r)T_{n,r}$. An unbiased estimator of β is

$$\hat{\beta} = \frac{S_{n,r}}{r}, \quad V\{\hat{\beta}\} = \frac{\beta^2}{r}.$$

Exercise 9.19 Simulate a random sample of 100 failure times, following the Weibull distribution $W(2.5, 10)$. Draw a Weibull Probability plot of the data. Estimate the

parameters of the distribution from the parameters of the linear regression fitted to the Q - Q plot.

Solution 9.19 The Weibull QQ plotting is based on the regression of $\ln T_{(i)}$ on

$$W_i = \ln \left(-\ln \left(1 - \frac{i}{n+1} \right) \right).$$

The linear relationship is $\ln T_{(i)} = \frac{1}{\nu} W_i + \ln \beta$. The slope of the straight line estimates $\frac{1}{\nu}$ and the intercept estimates $\log \beta$.

```
np.random.seed(1)
n = 100
lnTi = np.log(sorted(stats.weibull_min(2.5, scale=10).rvs(n)))
Wi = [np.log(-np.log(1 - i / (n+1))) for i in range(1, n + 1)]
df = pd.DataFrame({'Wi': Wi, 'lnTi': lnTi})
model = smf.ols('lnTi ~ Wi + 1', data=df).fit()
print(model.params)
intercept, slope = model.params
fig, ax = plt.subplots(figsize=(4, 4))
ax.plot((Wi[0], Wi[-1]),
        (slope * Wi[0] + intercept, slope * Wi[-1] + intercept),
        color='grey')
ax.scatter(Wi, lnTi, color='black')
ax.set_xlabel(r'$W_{(i)}$')
ax.set_ylabel(r'$\ln(T_{(i)})$')
plt.show()
```

```
| Intercept    2.294805
| Wi           0.496692
| dtype: float64
```

For our sample of $n = 100$ from $W(2.5, 10)$ the regression of $\ln T_{(i)}$ on W_i is

$$\ln(\text{TTF}) = 2.295 + 0.497W$$

Predictor	Coef	Stdev	t-ratio	p
Constant	2.2948	0.018	130.703	0.000
W	0.4967	0.013	37.629	0.000

R-sq = 0.935 R-sq(adj) = 0.935.

```
beta = np.exp(intercept)
nu = 1 / slope
print(f'beta {beta:.3f}, nu {nu:.3f}')
```

```
| beta 9.922, nu 2.013
```

The graphical estimate of β is $\hat{\beta} = \exp(2.295) = 9.922$. The graphical estimate of ν is $\hat{\nu} = \frac{1}{0.4643} = 2.013$. Both estimates are close to the nominal values. In Fig. 9.2 we see the Weibull probability plot. This graph puts the sorted observations $W_{(i)}$ on the x -axis and $\ln T_i$ on the y -axis where T_i is the calculated probability of occurrence for each observation $W_{(i)}$ assuming a Weibull distribution. As shown above, the $\hat{\beta}$ and $\hat{\nu}$ can be calculated from the slope and intercept of the liner regression line.

(fig:ex-weibull-prob-plot)




Fig. 9.2 Weibull probability plot

{fig:ex-weibull-prob-plot}

Exercise 9.20 The following is a random sample of the compressive strength of 20 concrete cubes [kg/cm²].

94.9, 106.9, 229.7, 275.7, 144.5, 112.8, 159.3, 153.1, 270.6, 322.0,
216.4, 544.6, 266.2, 263.6, 138.5, 79.0, 114.6, 66.1, 131.2, 91.1

Make a lognormal Q - Q plot of these data and estimate the mean and standard deviation of this distribution.

Solution 9.20 The regression of $Y_i = \ln X_{(i)}$, $i = 1, \dots, n$ on the normal scores

$$Z_i = \Phi^{-1} \left(\frac{i-3/8}{n+1/4} \right) \text{ is}$$

```
Xi = [94.9, 106.9, 229.7, 275.7, 144.5, 112.8, 159.3, 153.1,
      270.6, 322.0, 216.4, 544.6, 266.2, 263.6, 138.5, 79.0,
      114.6, 66.1, 131.2, 91.1]
n = len(Xi)
Zi = [stats.norm().ppf((i - 3/8)/(n + 1/4)) for i in range(1, n+1)]

df = pd.DataFrame({'Zi': Zi, 'lnXi': np.log(sorted(Xi))})
model = smf.ols('lnXi ~ Zi + 1', data=df).fit()

intercept, slope = model.params
fig, ax = plt.subplots(figsize=(4, 4))
ax.plot((Zi[0], Zi[-1]),
```




Fig. 9.3 Normal Probability Plot of $\ln X$

(fig:ex-normal-prob-plot-concrete)

```
(slope * Zi[0] + intercept, slope * Zi[-1] + intercept),
color='grey')
ax.scatter(Zi, df['lnXi'], color='black')
ax.set_xlabel(r'$Z_{(i)}$')
ax.set_ylabel(r'$\ln(X_{(i)})$')
plt.show()
```

$$\ln(X) = 5.10 + 0.556Z.$$

Predictor	Coef	Stdev	t-ratio	p
Constant	5.09571	0.02201	232.23	0.000
Z	0.55589	0.02344	23.79	0.000
R-sq = 0.969 R-sq(adj) = 0.967.				

```
mu = intercept
sigma = slope
print(f'mu {mu:.3f}, sigma {sigma:.3f}')
```

```
| mu 5.096, sigma 0.556
```

The intercept $\hat{a} = 5.09571$ is an estimate of μ , and the slope $\hat{b} = 0.5559$ is an estimate of σ . The mean of $\text{LN}(\xi, \sigma^2)$ is $\xi = e^{\mu + \sigma^2/2}$ and its variance is $D^2 = e^{2\mu + \sigma^2}(e^{\sigma^2} - 1)$. Thus, the graphical estimates are $\hat{\xi} = 190.608$ and $\hat{D} = 13154.858$.

The mean \bar{X} and variance S^2 of the given sample are 189.04 and 12948.2. In Fig. 9.3 we see the normal probability plotting of the data.

{fig:ex-normal-prob-plot-concrete}

Exercise 9.21 The following data represent the time till first failure [days] of electrical equipment. The data were censored after 400 days.

13, 157, 172, 176, 249, 303, 350, 400⁺, 400⁺.

(Censored values appear as x^+ .) Make a Weibull $Q-Q$ plot of these data and estimate the median of the distribution.

Solution 9.21⁵ The $Q-Q$ plot is given in Fig. 9.4.

{fig:ex-weibull-prob-plot-electrical}

```
np.random.seed(1)
data = [13, 157, 172, 176, 249, 303, 350, 400, 400]
n = len(data)
lnTi = np.log(data)
Wi = [np.log(-np.log(1 - i / (n+1))) for i in range(1, n + 1)]
# exclude the censored data for the regression analysis
df = pd.DataFrame({'Wi': Wi[:7], 'lnTi': lnTi[:7]})
model = smf.ols('lnTi ~ Wi + 1', data=df).fit()
print(model.params)
intercept, slope = model.params
fig, ax = plt.subplots(figsize=(4, 4))
ax.plot((Wi[0], Wi[-1]),
        (slope * Wi[0] + intercept, slope * Wi[-1] + intercept),
        color='grey')
ax.scatter(Wi[:-2], lnTi[:-2], color='black')
ax.scatter(Wi[-2:], lnTi[-2:], color='lightgrey')
ax.set_xlabel(r'$W_{(i)}$')
ax.set_ylabel(r'$\ln(T_i)$')
plt.show()
```

```
Intercept    5.959374
Wi           1.168853
dtype: float64
```

The regression of $\ln T_{(i)}$ on $W_i = \ln \left(-\ln \left(1 - \frac{i}{10} \right) \right)$ for $i = 1, \dots, 7$ is

$$\ln(T) = 5.96 + 1.17W \quad (9.0.1)$$

Predictor	Coef	Stdev	t-ratio	p
Constant	5.9594	0.3068	19.43	0.000
W	1.1689	0.2705	4.32	0.008


$s = 0.5627$ $R\text{-sq} = 0.789$ $R\text{-sq(adj)} = 0.747$.

According to this regression line, the estimates of β and ν are $\hat{\beta} = \exp(5.9594) = 387.378$ and $\hat{\nu} = 1/1.1689 = 0.855$. The estimate of the median of the distribution is

$$\hat{M}e = \hat{\beta} \left(-\ln \left(\frac{1}{2} \right) \right)^{1/\hat{\nu}} = 252.33.$$

Exercise 9.22 Make a PL (Kaplan-Meir) estimate of the reliability function of an electronic device, based on 50 failure times in file **ELECFAIL.csv**.

⁵ TODO



{fig:ex-weibull-prob-plot-electrical}

Fig. 9.4 Weibull probability plot

Solution 9.22 ⁶ Use the `lifelines` package to fit a Kaplan-Meier estimator. The reliability function is shown in Fig. 9.5.

{fig:kaplan-meier-elecfail}

```
elecfail = mistat.load_data('ELECFAIL.csv')
```

```
kmf = lifelines.KaplanMeierFitter()
kmf.fit(elecfail)
kmf.plot_survival_function()
```

```
| <AxesSubplot:xlabel='timeline'>
```

Exercise 9.23 Assuming that the failure times in file `ELECFAIL.csv` come from an exponential distribution $E(\beta)$, compute the MLE of β and of $R(50; \beta) = \exp\{-50/\beta\}$. [The MLE of a function of a parameter is obtained by substituting the MLE of the parameter in the function.] Determine confidence intervals for β and for $R(50; \beta)$ at level of confidence 0.95.

Solution 9.23 Use the `lifelines` package to fit an exponential model.

⁶ TODO




Fig. 9.5 Kaplan-Meier Estimator of The Reliability of An Electric Device

(fig:kaplan-meier-elecfail)

```
elecfail = mistat.load_data('ELECFAIL.csv')

kmf = lifelines.ExponentialFitter()
kmf.fit(elecfail)
kmf.print_summary()
```

	coef	se(coef)	coef lower 95%	coef upper 95%	z	p	-log2(p)
lambda_	57.07	8.07	41.25	72.89	6.95	0.00	37.97

The MLE of β is $\hat{\beta} = \bar{X} = 57.07$. Since $R(t) = \exp(-t/\beta)$, the MLE of $R(50)$ is $\hat{R}(50) = 0.4164$. Note that $\tilde{X} \sim \frac{\beta}{2n}\chi^2[2n]$. Hence, confidence limits for β , with $1 - \alpha = 0.95$ are (44.048, 76.891). Substituting these limits into $R(t)$ gives (0.3214, 0.5219) as a 0.95 confidence interval for $R(50)$.⁷

Exercise 9.24 The following are values of 20 random variables having an exponential distribution $E(\beta)$. The values are censored at $t^* = 200$.

96.88, 154.24, 67.44, 191.72, 173.36, 200, 140.81, 200, 154.71, 120.73,
24.29, 10.95, 2.36, 186.93, 57.61, 99.13, 32.74, 200, 39.77, 39.52.

⁷ TODO: the Python code prints CI based on normal distribution

Determine the MLE of β . Use β equal to the MLE, to estimate the standard deviation of the MLE, and to obtain confidence interval for β , at level $1 - \alpha = .95$. [This simulation is called an empirical Bootstrap.]

Solution 9.24 Use the `lifelines` package to fit an exponential model.

```
T = [96.88, 154.24, 67.44, 191.72, 173.36, 200, 140.81, 200,
      154.71, 120.73, 24.29, 10.95, 2.36, 186.93, 57.61, 99.13,
      32.74, 200, 39.77, 39.52]
E = [ti < 200 for ti in T]

kmf = lifelines.ExponentialFitter()
kmf.fit(T, E)
kmf.print_summary()
```

	coef	se(coef)	coef lower 95%	coef upper 95%	z	p	-log2(p)
lambda_	129.01	31.29	67.68	190.34	4.09	0.00	14.51

We have $n = 20$, $K_n = 17$ and $\hat{\beta} = 129.011$. With $M = 500$ runs, the estimated STD of $\hat{\beta}$ is 34.912, with confidence interval (78.268, 211.312).⁸

Exercise 9.25 Determine n^0 and r for a frequency censoring test for the exponential distribution, where the cost of a unit is 10 times bigger than the cost per time unit of testing. We wish that $S.E.\{\hat{\beta}_n\} = 0.1\beta$, and the expected cost should be minimized at $\beta = 100$ [hr]. What is the expected cost of this test, at $\beta = 100$, when $c_1 = \$1$ [hr].

Solution 9.25 For $\beta = 100$, $S.E.\{\hat{\beta}_n\} = 0.1\beta = 10 = \frac{100}{\sqrt{r}}$. Hence $r = 100$ and

$$n^0 \approx \frac{100}{2} \left(1 + \left(1 + 4 \frac{100}{1000} \right)^{1/2} \right) = 109.16 \approx 110.$$

$E\{T_{n,r}\} = 235.327$ and the expected cost of testing is $E\{C\} = 10 \times n + 1 \times E\{T_{n,r}\} = \1335.33 .

Exercise 9.26 File **WEIBUL.csv** contains the values of a random sample of size $n = 50$ from a Weibull distribution.

- (i) Obtain MLE of the scale and shape parameters β and ν .
- (ii) Use the MLE estimates

of β and ν , to obtain parametric bootstrap EBD of the distribution of $\hat{\beta}$, $\hat{\nu}$, with $M = 500$ runs. Estimate from this distribution the standard deviations of $\hat{\beta}$ and $\hat{\nu}$. Compare these estimates to the large sample approximations.

Solution 9.26 (a)

```
T = mistat.load_data('WEIBUL.csv')

kmf = lifelines.WeibullFitter()
kmf.fit(T)
kmf.print_summary()
```

⁸ TODO: where in the book is this calculation described?

	coef	se(coef)	coef lower 95%	coef upper 95%	z	p	-log2(p)
lambda_	27.08	2.94	21.31	32.85	8.86	0.00	60.17
rho_	1.37	0.15	1.09	1.66	2.57	0.01	6.63

For the WEIBUL.csv data we obtained $\hat{\beta} = 27.0789$ and $\hat{\nu} = 1.374$.

(b)⁹

```

n_boot=5
idx = list(range(len(T)))
def stat_func(x):
    epf = lifelines.WeibullFitter().fit(T[x])
    return epf.params_['lambda_']

ci, dist = pg.compute_bootci(idx, func=stat_func, n_boot=n_boot, confidence=0.95,
                             method='per', seed=1, return_dist=True)
print(f'std(beta): {np.std(dist)}')

def stat_func(x):
    epf = lifelines.WeibullFitter().fit(T[x])
    return epf.params_['rho_']
ci, dist = pg.compute_bootci(idx, func=stat_func, n_boot=n_boot, confidence=0.95,
                             method='per', seed=1, return_dist=True)
print(f'nu(std): {np.std(dist)}')

```

```

std(beta): 2.8735759059232406
nu(std): 0.07224217322172025

```

The EBD estimates are $SE\{\hat{\beta}_{50}\} = 2.971^{10}$ and $SE\{\hat{\nu}_{50}\} = 0.1822$. The asymptotic estimates are $SE\{\hat{\beta}_{50}\} = 2.94^{11}$ and $SE\{\hat{\nu}_{50}\} = 0.15$. The EBD estimates and the asymptotic estimates of the standard deviations of $\hat{\beta}$ and $\hat{\nu}$ are very similar.

Exercise 9.27 In binomial life testing by a fixed size sample, how large should the sample be in order to discriminate between $R_0 = 0.99$ and $R_1 = 0.90$, with $\alpha = \beta = 0.01$? [α and β denote the probabilities of error of Type I and II.]

Solution 9.27 To discriminate between $R_0 = 0.99$ and $R_1 = 0.90$ with $\alpha = \beta = 0.01$, $n \approx 107$.

Exercise 9.28 Design the Wald SPRT for binomial life testing, in order to discriminate between $R_0 = 0.99$ and $R_1 = 0.90$, with $\alpha = \beta = 0.01$. What is the expected sample size, ASN, if $R = 0.9$?

Solution 9.28 ¹² For $R_0 = 0.99$, $R_1 = 0.90$, and $\alpha = \gamma = 0.01$ we get $s = 0.9603$, $h_1 = 1.9163$, $h_2 = 1.9163$ and $ASN(0.9) = 31.17$.

Exercise 9.29 Design a Wald SPRT for exponential life distribution, to discriminate between $R_0 = 0.99$ and $R_1 = 0.90$, with $\alpha = \beta = 0.01$. What is the expected sample size, ASN, when $R = 0.90$?

⁹ TODO: set n_boot to 500 once chapter is done!!!

¹⁰ TODO: should this be $SE\{\hat{\beta}_{500}\}$

¹¹ TODO: should this be $SE\{\hat{\beta}\}$

¹² TODO

Solution 9.29 Without loss of generality, assume that $t = 1$. Thus, $R_0 = e^{-1/\beta_0} = 0.99$, or $\beta_0 = 99.5$. Also, $R_1 = 0.9$, or $\beta_1 = 9.49$. The parameters of the SPRT are: $h_1 = 48.205$, $h_2 = 48.205$ and $s = 24.652$. $ASN(0.9) \approx 3$.

Exercise 9.30 $n = 20$ computer monitors are put on accelerated life testing. The test is an SPRT for Poisson processes, based on the assumption that the TTF of a monitor, in those conditions, is exponentially distributed. The monitors are considered to be satisfactory if their MTBF is $\beta \geq 2,000$ [hr] and considered to be unsatisfactory if $\beta \leq 1,500$ [hr]. What is the expected length of the test if $\beta = 2,000$ [hr].

Solution 9.30 Using $\alpha = \gamma = 0.05$ we have $n = 20$, $H_0 : \beta = 2,000$, $H_1 : \beta = 1,500$. Thus, $\lambda_0 = 0.0005$ and $\lambda_1 = 0.00067$. The parameters of the SPRT are:

$$h_1 = \frac{\log(\frac{95}{5})}{\log(\frac{\lambda_1}{\lambda_0})} = 10.061, \quad h_2 = 10.061$$

and

$$\begin{aligned} s &= \frac{\lambda_1 - \lambda_0}{\log(\frac{\lambda_1}{\lambda_0})} = 0.000581. \\ E_{\beta_0}\{\tau\} &= \frac{2000}{20} E_{\beta}\{X_{20}(\tau)\} \\ &\approx 100 \left(\frac{10.061 - 0.95 \times 20.122}{1 - 2000 \times 0.000581} \right) \\ &= 5589.44 \text{ [hr]}. \end{aligned}$$

Exercise 9.31 A product has an exponential life time with MTTF $\beta = 1,000$ [hr]. 1% of the products come out of production with MTTF of $\gamma = 500$ [hr]. A burn in of $t^* = 300$ [hr] takes place. What is the expected life of units surviving the burn-in? Is such a long burn-in justified?

Solution 9.31 ¹³ With $\beta = 1,000$, $p = 0.01$, $\gamma = 500$ and $t^* = 300$, we get

$$F^*(t) = 1 - \frac{0.01e^{-t/500} + 0.99e^{-t/1,000}}{0.01e^{-300/500} + 0.99e^{-300/1,000}}, \quad t \geq 300.$$

¹⁴ The expected life of units surviving the burn-in is thus,

```
factor = 0.01*np.exp(-300/500) + 0.99*np.exp(-300/1000)
integral = 0.01*500*np.exp(-300/500) + 0.99*1000*np.exp(-300/1000)
with_burn_in = 300 + integral/factor
no_burn_in = 0.01 * 500 + 0.99 * 1000
with_burn_in - no_burn_in
```

¹³ TODO

¹⁴ TODO: the calculation of expected life of units is missing in the current solution. Please check what I added

| 301.2862836203999

$$\begin{aligned}
 \beta^* &= 300 + \frac{1}{0.7389} \int_{300}^{\infty} \left(0.01 \exp \left\{ -\frac{t}{500} \right\} + 0.99 \exp \left\{ -\frac{t}{1,000} \right\} \right) dt \\
 &= 300 + \frac{5}{0.7389} \exp \left\{ -\frac{300}{500} \right\} + \frac{990}{0.7389} \exp \left\{ -\frac{300}{1,000} \right\} \\
 &= 1,296.29 \text{ [hr]}.
 \end{aligned}$$

Without burn in, the expected life of these units is $0.01 * 500 + 0.99 * 1000 = 995$ [hr].
 The expected life of the units in the field increases by 300 [hr].

Chapter 10

Bayesian Reliability Estimation and Prediction

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats
import statsmodels.formula.api as smf
import lifelines
import pingouin as pg
import seaborn as sns
import matplotlib.pyplot as plt
import mistat
```

Exercise 10.1 Suppose that the TTF of a system is a random variable having exponential distribution, $E(\beta)$. Suppose also that the prior distribution of $\lambda = 1/\beta$ is $G(2.25, .01)$.

- (i) What is the posterior distribution of λ , given $T = 150$ [hr]?
- (ii) What is the Bayes estimator of β , for the squared-error loss?
- (iii) What is the posterior SD of β ?

Solution 10.1 ¹

Exercise 10.2 Let $J(t)$ denote the number of failures of a device in the time interval $(0, t]$. After each failure the device is instantaneously renewed. Let $J(t)$ have a Poisson distribution with mean λt . Suppose that λ has a gamma prior distribution, with parameters $\nu = 2$ and $\tau = .05$.

- (i) What is the predictive distribution of $J(t)$?
- (ii) Given that $J(t)/t = 10$, how many failures are expected in the next time unit?
- (iii) What is the Bayes estimator of λ , for the squared-error loss?
- (iv) What is the posterior SD of λ ?

Solution 10.2 ²

¹ TODO

² TODO

Exercise 10.3 The proportion of defectives, θ , in a production process has a uniform prior distribution on $(0, 1)$. A random sample of $n = 10$ items from this process yields $K_{10} = 3$ defectives.

- (i) What is the posterior distribution of θ ?
- (ii) What is the Bayes estimator of θ for the absolute error loss?

Solution 10.3 ³

Exercise 10.4 Let $X \sim \mathcal{P}(\lambda)$ and suppose that λ has the Jeffrey improper prior $h(\lambda) = \frac{1}{\sqrt{\lambda}}$. Find the Bayes estimator for squared-error loss and its posterior SD.

Solution 10.4 ⁴

Exercise 10.5 Apply formula (10.2.3) to determine the Bayes estimator of the reliability when $n = 50$ and $K_{50} = 49$.

Solution 10.5 ⁵

Exercise 10.6 A system has three modules, M_1, M_2, M_3 . M_1 and M_2 are connected in series and these two are connected in parallel to M_3 , i.e.,

$$R_{sys} = \psi_p(R_3, \psi_s(R_1, R_2)) = R_3 + R_1 R_2 - R_1 R_2 R_3,$$

where R_i is the reliability of module M_i . The TTFs of the three modules are independent random variables having exponential distributions with prior $\text{IG}(\nu_i, \tau_i)$ distributions of their MTTF. Moreover, $\nu_1 = 2.5$, $\nu_2 = 2.75$, $\nu_3 = 3$, $\tau_1 = \tau_2 = \tau_3 = 1/1000$. In separate independent trials of the TTF of each module we obtained the statistics $T_n^{(1)} = 4565$ [hr], $T_n^{(2)} = 5720$ [hr] and $T_n^{(3)} = 7505$ [hr], where in all three experiments $n = r = 10$. Determine the Bayes estimator of R_{sys} , for the squared-error loss.

Solution 10.6 ⁶

Exercise 10.7 $n = 30$ computer monitors were put on test at a temperature of 100°F and relative humidity of 90% for 240 [hr]. The number of monitors which survived this test is $K_{30} = 28$. Determine the Bayes credibility interval for $R(240)$, at level $\gamma = .95$, with respect to a uniform prior on $(0, 1)$.

Solution 10.7 ⁷

(exc:rtf)

Exercise 10.8 Determine a $\gamma = .95$ level credibility interval for $R(t)$ at $t = 25$ [hr] when $\text{TTF} \sim E(\beta)$, $\beta \sim \text{IG}(3, .01)$, $r = 27$, $T_{n,r} = 3500$ [hr].

³ TODO

⁴ TODO

⁵ TODO

⁶ TODO

⁷ TODO

Solution 10.8 ⁸

Exercise 10.9 Under the conditions of Exercise 10.8 determine a Bayes prediction interval for the total life of $s = 2$ devices. {excr:rf}

Solution 10.9 ⁹

Exercise 10.10 A repairable system has exponential TTF and exponential TTR, which are independent of each other. $n = 100$ renewal cycles were observed. The total times till failure were 10,050 [hr] and the total repair times were 500 [min]. Assuming gamma prior distributions for λ and μ with $\nu = \omega = 4$ and $\tau = .0004$ [hr], $\zeta = .01$ [min], find a $\gamma = .95$ level credibility interval for A_∞ .

Solution 10.10 ¹⁰

Exercise 10.11 In reference to Example 10.6, suppose that the data of Table 10.1 were obtained for a Poisson random variable where $\lambda_1, \dots, \lambda_{188}$ have a gamma (ν, τ) prior distribution. {tbl:emp-dist-solder-defects}

- (i) What is the predictive distribution of the number of defects per batch?
- (ii) Find the formulae for the first two moments of the predictive distribution.
- (iii) Find, from the empirical frequency distribution of Table 10.1, the first two sample moments. {tbl:emp-dist-solder-defects}
- (iv) Use the method of moment equations to estimate the prior parameters ν and τ .
- (v) What is the Bayes estimator of λ if $X_{189} = 8$?

Solution 10.11 ¹¹

⁸ TODO

⁹ TODO

¹⁰ TODO

¹¹ TODO

Chapter 11

Sampling Plans for Batch and Sequential Inspection

Import required modules and define required functions

```
import numpy as np
import pandas as pd
from scipy import stats, optimize
import statsmodels.formula.api as smf
import lifelines
import pingouin as pg
import seaborn as sns
import matplotlib.pyplot as plt
import mistat
```

(exc:single-sampling-plan)

Exercise 11.1 Determine single sampling plans for attributes, when the lot is $N = 2500$, $\alpha = \beta = 0.01$, and

- (i) $AQL = 0.005$, $LQL = 0.01$
- (ii) $AQL = 0.01$, $LQL = 0.03$
- (iii) $AQL = 0.01$, $LQL = 0.05$

Solution 11.1 The single-sampling plans for attributes with $N = 2500$, $\alpha = \beta = 0.01$, and AQL and LQL as specified are (i) $n = 1878$, $c = 13$; (ii) $n = 702$, $c = 12$; (iii) $n = 305$, $c = 7$.¹

Exercise 11.2 Investigate how the lot size, N , influences the single sampling plans for attributes, when $\alpha = \beta = 0.05$, $AQL = 0.01$, $LQL = 0.03$, by computing the plans for $N = 100$, $N = 500$, $N = 1,000$, $N = 2,000$.

Solution 11.2 For $\alpha = \beta = 0.05$, $AQL = 0.01$, and $LQL = 0.03$ the single-sampling plans for attributes are

N	n	c
100	87	1
500	254	4
1000	355	6
2000	453	8

¹ TODO: example Python code?

We see that as the lot size, N , increases then the required sample size increases, but n/N decreases from 87% to 22.6%. The acceptance number c increases very slowly.

Exercise 11.3 Compute the $OC(p)$ function for the sampling plan computed in Exercise 11.1(iii). What is the probability of accepting a lot having 2.5% of nonconforming items?

Solution 11.3 For the sampling plan in Exercise 11.1(iii), $OC(p) = H(7; 2500, M_p, 305)$. When $p = 0.025$, we get $M_p = 62$ and $OC(0.025) = 0.5091$.

Exercise 11.4 Compute the large sample approximation to a single sample plan for attributes (n^*, c^*) , with $\alpha = \beta = .05$ and $AQL = .025$, $LQL = .06$. Compare these to the exact results. The lot size is $N = 2,000$.

Solution 11.4 The large sample approximation yields $n^* = 292$, $c^* = 11$. The "exact" plan is $n = 311$, $c = 12$. Notice that the actual risks of the large sample approximation plan (n^*, c^*) are $\alpha^* = 0.0443$ and $\beta^* = 0.0543$. The actual risks of the "exact" plan are $\alpha = 0.037$ and $\beta = 0.0494$.

Exercise 11.5 Repeat the previous Exercise with $N = 3000$, $\alpha = \beta = 0.10$, $AQL = .01$ and $LQL = .06$.

Solution 11.5 The large sample approximation is $n^* = 73$, $c^* = 1$ with actual risks of $\alpha^* = 0.16$ and $\beta^* = 0.06$. The exact plan is $n = 87$, $c = 2$ with actual risks of $\alpha = 0.054$ and $\beta = 0.097$.

Exercise 11.6 Obtain the OC and ASN functions of the double sampling plan, with $n_1 = 200$, $n_2 = 2n_1$ and $c_1 = 5$, $c_2 = c_3 = 15$, when $N = 2,000$.

- (i) What are the attained α and β when $AQL = 0.015$ and $LQL = 0.05$?
- (ii) What is the ASN when $p = AQL$?
- (iii) What is a single sampling plan having the same α and β ? How many observations we expect to save if $p = AQL$? Notice that if $p = LQL$ the present double sampling plan is less efficient than the corresponding single sampling plan.

Solution 11.6 (i) The attained α and β are $\alpha = 0.021$, $\beta = 0.0532$. (ii) $ASN = 228.7$. (iii) The single-sampling plan is $n = 253$, $c = 7$ with $\alpha^* = 0.0283$, $\beta^* = 0.0486$. If $p = AQL$ we expect to save 24 observations.

Exercise 11.7 Compute the OC and ASN values for a double sampling plan with $n_1 = 150$, $n_2 = 200$, $c_1 = 5$, $c_2 = c_3 = 10$, when $N = 2,000$. Notice how high β is when $LQL = 0.05$. The present plan is reasonable if $LQL = 0.06$. Compare this plan to a single sampling one for $\alpha = 0.02$ and $\beta = 0.10$, $AQL = 0.02$, $LQL = 0.06$.

Solution 11.7 The double-sampling plan with $n_1 = 150$, $n_2 = 200$, $c_1 = 5$, $c_2 = c_3 = 10$, and $N = 2000$ yields

p	$OC(p)$	$ASN(p)$	p	$OC(p)$	$ASN(p)$
0	1.000	150	0.06	0.099	247.7
0.01	0.999	150.5	0.07	0.039	218.9
0.02	0.966	164.7	0.08	0.014	191.5
0.03	0.755	206.0	0.09	0.005	171.9
0.04	0.454	248.4	0.10	0.001	160.4
0.05	0.227	262.8			

The single sampling plan for $\alpha = 0.02$, $\beta = 0.10$, $AQL = 0.02$, $LQL = 0.06$ is $n = 210$, $c = 8$ with actual $\alpha = 0.02$, $\beta = 0.10$.

Exercise 11.8 Determine a sequential plan for the case of $AQL = 0.02$, $LQL = 0.06$, $\alpha = \beta = 0.05$. Compute the OC and ASN functions of this plan. What are the ASN values when $p = AQL$, $p = LQL$ and $p = 0.035$?

Solution 11.8 For the sequential plan, $OC(0.02) = 0.95 = 1 - \alpha$, $OC(0.06) = 0.05 = \beta$. $ASN(0.02) = 140$, $ASN(0.06) = 99$ and $ASN(0.035) = 191$.

Exercise 11.9 Compare the single sampling plan and the sequential one when $AQL = 0.01$, $LQL = 0.05$, $\alpha = \beta = 0.01$ and $N = 10,000$. What are the expected savings in sampling cost, if each observation costs \$1, and $p = AQL$?

Solution 11.9 The single-sampling plan for $N = 10,000$, $\alpha = \beta = 0.01$, $AQL = 0.01$ and $LQL = 0.05$ is $n = 341$, $c = 8$. The actual risks are $\alpha = 0.007$, $\beta = 0.0097$. The corresponding sequential plan, for $p = AQL = 0.01$ has $ASN(0.01) = 182$. On the average, the sequential plan saves, under $p = AQL$, 159 observations. This is an average savings of \$159 per inspection.

Exercise 11.10 Use the `mistat` function “simulateOAB” to simulate the expected rewards, for $p = 0.4(0.05)0.8$, when $N = 75$, $\lambda = 0.6$, $k = 15$, $\gamma = 0.95$, $Ns = 1000$.

Solution 11.10 ²

Exercise 11.11 Use the `mistat` function “optimalOAB” to predict the expected reward under the optimal strategy, when $N = 75$, $\lambda = 0.6$.

Solution 11.11 ³

Exercise 11.12 Consider the TAB^4 with $N = 40$ and $K = 10$. Make a table of all the possible predicted rewards.

Solution 11.12 ⁵

² TODO

³ TODO

⁴ **TODO**: write out acronym

⁵ TODO

Exercise 11.13 Determine n and k for a continuous variable size sampling plan, when $(p_0) = AQL = .01$ and $(p_t) = LQL = .05$, $\alpha = \beta = 0.05$.

Solution 11.13 For a continuous variable-size sampling plan when $(p_0) = AQL = 0.01$, $(p_t) = LQL = 0.05$, and $\alpha = \beta = 0.05$, we obtain $n = 70$ and $k = 1986$.

Exercise 11.14 Consider data file **ALMPIN.csv**. An aluminum pin is considered as defective if its cap diameter is smaller than 14.9 [mm]. For the parameters $p_0 = 0.01$, $\alpha = 0.05$, compute k and decide whether to accept or reject the lot, on the basis of the sample of $n = 70$ pins. What is the probability of accepting a lot with proportion defectives of $p = 0.03$?

Solution 11.14 From the data we get $\bar{X} = 14.9846$ and $S = 0.019011$. For $p_0 = 0.01$ and $\alpha = 0.05$ we obtain $k = 2.742375$. Since $\bar{X} - kS = 14.9325 > \zeta = 14.9$, the lot is *accepted*. $OC(0.03) \approx 0.4812$.

Exercise 11.15 Determine the sample size and k for a single sampling plan by a normal variable, with the parameters $AQL = 0.02$, $LQL = 0.04$, $\alpha = \beta = 0.10$.

Solution 11.15 For $AQL = 0.02$, $LQL = 0.04$ and $\alpha = \beta = 0.10$, $n = 201$, $k = 1.9022$.

Exercise 11.16 A single sampling plan for attributes, from a lot of size $N = 500$, is given by $n = 139$ and $c = 3$. Each lot which is not accepted is rectified. Compute the AOQ , when $p = 0.01$, $p = 0.02$, $p = 0.03$ and $p = 0.05$. What are the corresponding ATI values?

Solution 11.16 For a single-sampling plan for attributes where $n = 139$, $c = 3$ and $N = 500$ we obtain $OC(p) = H(3; 500, M_p, 139)$, and $R^* = H(2; 499, M_p - 1, 138)/H(3; 500, M_p, 139)$.

p	AOQ	ATI
0.01	0.0072	147.2
0.02	0.0112	244.2
0.03	0.0091	369.3
0.05	0.0022	482.0

Exercise 11.17 A single sampling plan, under normal inspection has probability $\alpha = 0.05$ of rejection, when $p = AQL$. What is the probability, when $p = AQL$ in 5 consecutive lots, that there will be a switch to tightened inspection? What is the probability of switching to a tightened inspection if p increases so that $OC(p) = .7$?

Solution 11.17 A switch to a tightened plan, when all 5 consecutive lots have $p = AQL$, with $\alpha = 0.05$, is $\sum_{j=2}^5 b(j; 5, 0.05) = 0.0226$. The probability of switching to a tightened plan if $\alpha = 0.3$ is 0.4718.

Exercise 11.18 Compute the probability for qualifying for State 2, in a skip-lot sampling plan, when $n = 100$, $c = 1$. What is the upper bound on S_{10} , in order to qualify for State 2, when $AQL = 0.01$? Compute the probability QP for State 2 qualification.

Solution 11.18 The total sample size from 10 consecutive lots is 1000. Thus, from Table 11.15⁶ S_{10} should be less than 5. The last 2 samples should each have less than 2 defective items. Hence, the probability for qualification is

$$\begin{aligned} QP &= b^2(1; 100, 0.01)B(2; 800, 0.01) \\ &\quad + 2b(1; 100, 0.01)b(0; 100, 0.01)B(3; 800, 0.01) \\ &\quad + b^2(0; 100, 0.01)B(4; 800, 0.01) \\ &= 0.0263. \end{aligned}$$

Exercise 11.19 The **FAILURE_J2** dataset contains the cumulative failure counts of a software project collected over a period of 181 weeks. Fit the following models to the data and visualize the results.

- Goel-Okumoto $f(t) = a[1 - \exp(-bt)]$
- Musa-Okumoto $f(t) = \frac{1}{\varphi} \log(\lambda_0 \varphi t + 1)$
- S-shaped Yamada $f(t) = a(1 - (1 + bt) \exp(-bt))$
- inflected S-shaped Ohba $f(t) = \frac{a(1 - \exp(-bt))}{1 + c \exp(-bt)}$

Which model describes the data best?

Solution 11.19 Define function for the four models

```
def modelGoelOkumoto(t, a, b):
    return a * (1 - np.exp(-b * t))

def modelMusaOkumoto(t, phi, lam):
    return (1/phi) * np.log(lam*phi*t + 1)

def modelYamada(t, a, b):
    return a * (1 - (1+b*t)*np.exp(-b*t))

def modelInflectedSshaped(t, a, b, c):
    return a * (1 - np.exp(-b * t)) / (1 + c * np.exp(-b * t))
```

Fit the four functions to the cumulative failure count data CFC

```
def optimizeModelFit(model, data, x, y):
    fit = optimize.curve_fit(model, data[x], data[y], method='lm')
    popt = fit[0]
    # add the fit curve to the data
    data[model.__name__] = [model(t, *popt) for t in data[x]]
    return popt
data = mistat.load_data('FAILURE_J2')
optimizeModelFit(modelGoelOkumoto, data, 'T', 'CFC')
optimizeModelFit(modelMusaOkumoto, data, 'T', 'CFC')
optimizeModelFit(modelYamada, data, 'T', 'CFC')
optimizeModelFit(modelInflectedSshaped, data, 'T', 'CFC')
```

```
| array([2.26933527e+02, 5.40121092e-02, 5.96246327e+01])
```

The fits are shown in Fig. 11.1. The inflected S-shaped model describes the data best.

⁶ TODO: is this the correct table?

{exc:failure-j2}

{fig:exc-failure-j2}

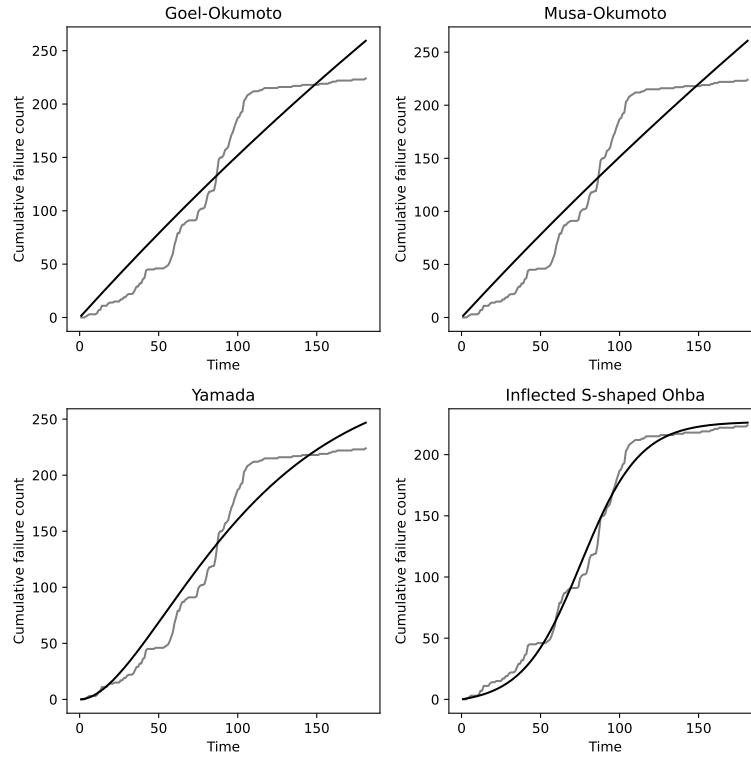


Fig. 11.1 NHPP models fitted to cumulative failure count data **FAILURE_J2**

(fig:exc-failure-j2)

Exercise 11.20 Continuing with Exercise 11.19, simulate cases where you only have data for the first 25, 50, 75, 100, or 125 weeks and fit Goel-Okumoto and inflected S-shaped models to these subsets.

(exc:failure-j2)

Discuss the results with respect to using the models to extrapolate the expected failure count to predict the expected failure count into the future.

Solution 11.20 Modify the `optimizeModelFit` function from the previous exercise to restrict the fit to a subset of the data. The

```
def optimizeModelFit(model, data, x, y, subset):
    # create the subset
    subsetX = data[x][:subset]
    subsetY = data[y][:subset]
    # fit curve to subset - a increase of maxfev is required
    fit = optimize.curve_fit(model, subsetX, subsetY,
                             method='lm', maxfev=2000)

    popt = fit[0]
    data[f'{model.__name__} {subset}'] = [model(t, *popt) for t in data[x]]
    return popt

for subset in [25, 50, 75, 100, 125, 150]:
    optimizeModelFit(modelGoelOkumoto, data, 'T', 'CFC', subset)
    optimizeModelFit(modelInflectedSshaped, data, 'T', 'CFC', subset)
```

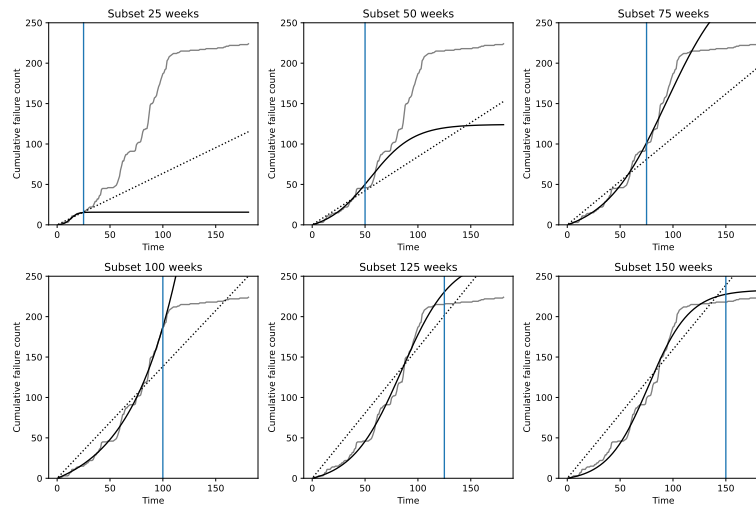


Fig. 11.2 NHPP models fitted to cumulative failure count data **FAILURE_J2** using subsets. The actual data are shown in grey, the Goel-Okumoto fits as dotted black lines, and the inflected S-shaped models as solid black lines. The vertical line shows the cutoff used for subsetting the data for the fits.

{fig:exc-failure-j2-subset}

{fig:exc-failure-j2-subset}

The result of this simulation is shown in Fig. 11.2. As in the previous exercise, the inflected S-shaped model performs best. However, the fit using only 25 weeks of data is insufficient to extrapolate into the future. With 50 or 75 weeks of data, the inflected S-shaped model predicts the following 20-30 weeks well. Shortly after 100 weeks, the cumulative failure count curve flattens. This is not predicted by the model fit. Even with 125 weeks of data, this is not well described. Only at 150 weeks, the inflected S-shaped model starts to show flattening.

Exercise 11.21 The dataset **FAILURE_DS2** contains cumulative failure counts of a software project collected over a period of 18 weeks. Fit an inflected S-shaped models to these data.

You will see that the large initial number of failures leads to an insufficient fit of the data. Create a second fit where the initial number of failures is ignored and discuss the results.

Solution 11.21 Load the data and fit the model

```
data = mistat.load_data('FAILURE_DS2')
fit = optimize.curve_fit(modelInflectedSshaped, data['T'], data['CFC'])
popt = fit[0]

# add the fit curve to the data and visualize
data['Model'] = [modelInflectedSshaped(t, *popt) for t in data['T']]

ax = data.plot(x='T', y='CFC', color='grey')
data.plot(x='T', y='Model', color='black', ax=ax)
plt.show()
```




Fig. 11.3 NHPP models fitted to cumulative failure count data **FAILURE_DS2**. Left: initial model, right: after correction

{fig:exc-failure-ds2}

The result of this simulation is shown in Fig. 11.3. It is clear that the fit could be better if the model ignores the initial failure count. Define a new model to correct for this problem.

{fig:exc-failure-ds2}

```

initial = data['CFC'][0]
def correctedModel(t, a, b, c):
    return modelInflectedSshaped(t, a, b, c) + initial

fit = optimize.curve_fit(correctedModel, data['T'], data['CFC'])
popt = fit[0]
# add the fit curve to the data and visualize
data['Corrected Model'] = [correctedModel(t, *popt) for t in data['T']]

ax = data.plot(x='T', y='CFC', color='grey')
data.plot(x='T', y='Model', color='black', ax=ax, linestyle=':')
data.plot(x='T', y='Corrected Model', color='black', ax=ax)
plt.show()

```

Including the offset in the model, leads to a better fit of the data.